

Learning with Meshes

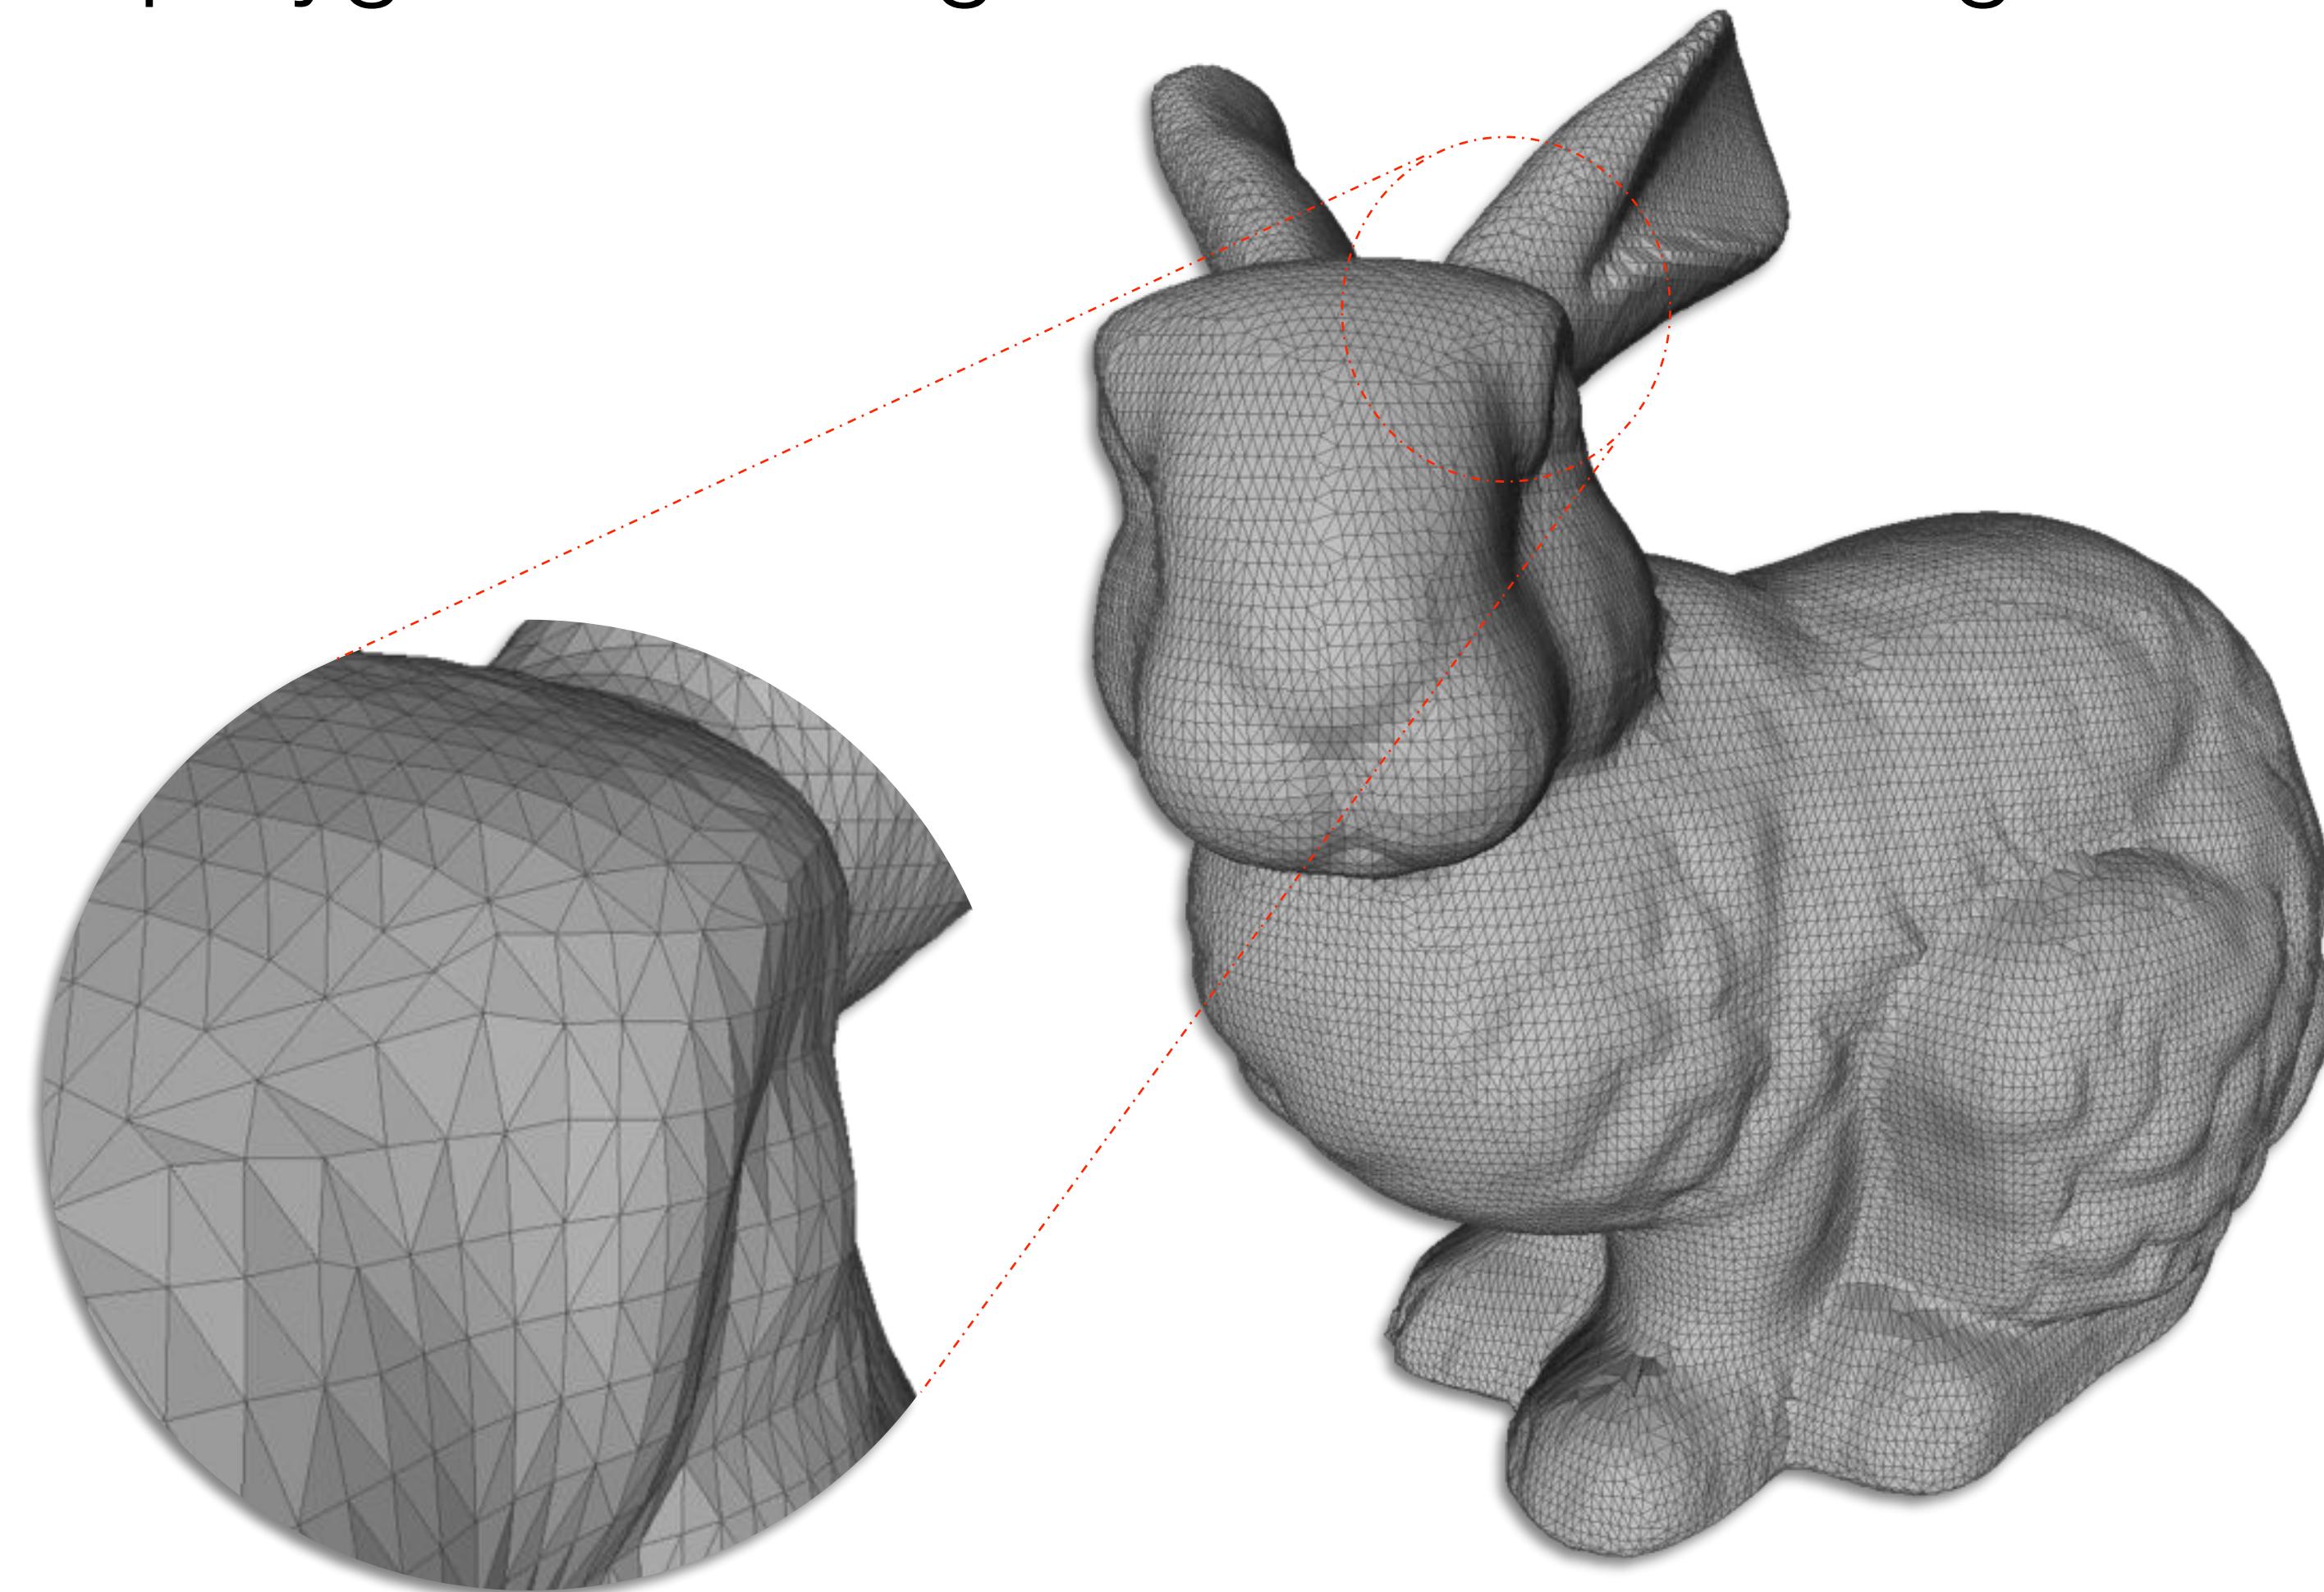
Alexey Artemov

MLSS 2019 tutorial

Acknowledgement: Enrico Puppo, Daniele Panozzo,
Michael Bronstein, and Albert Matveev

What is a mesh?

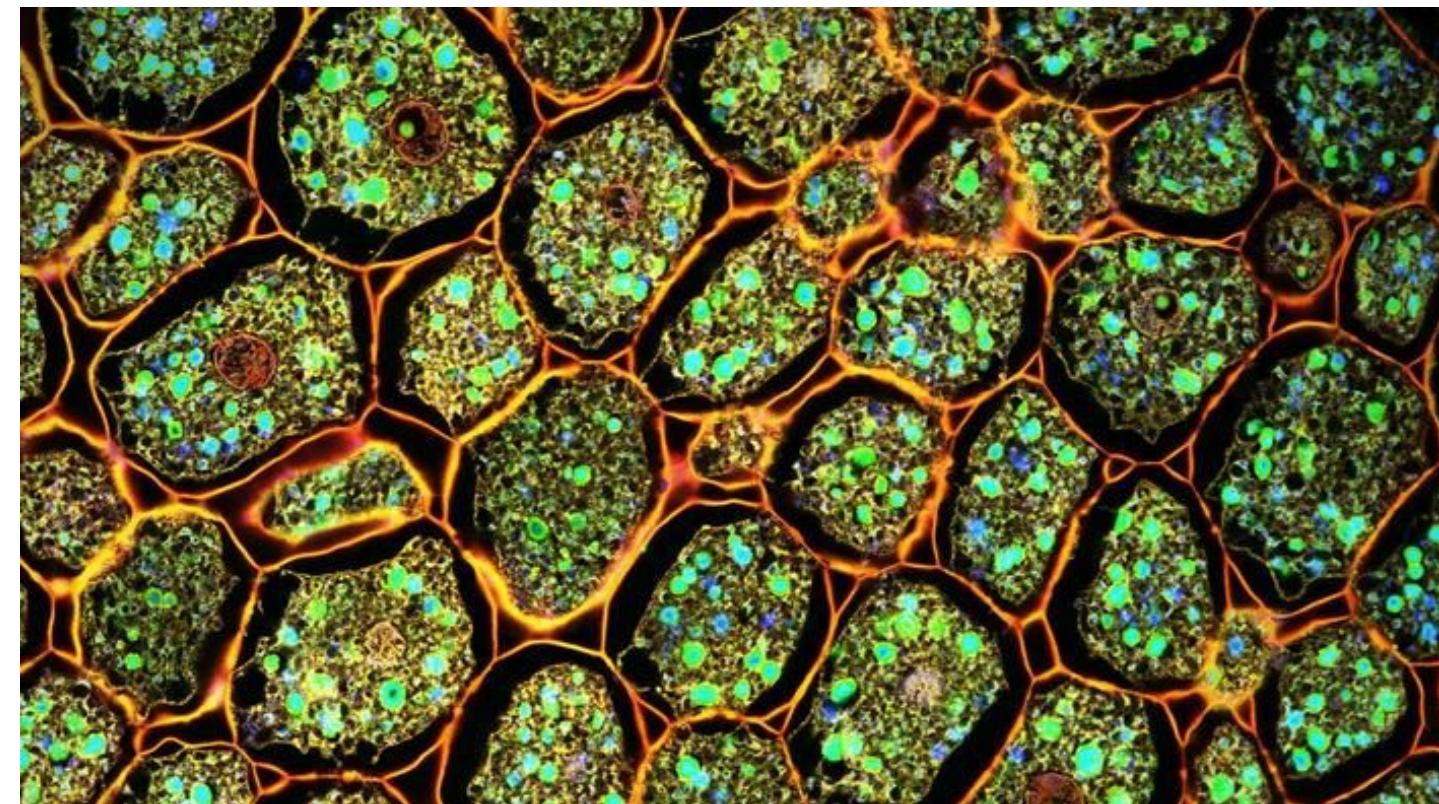
A surface made of polygonal faces glued at common edges



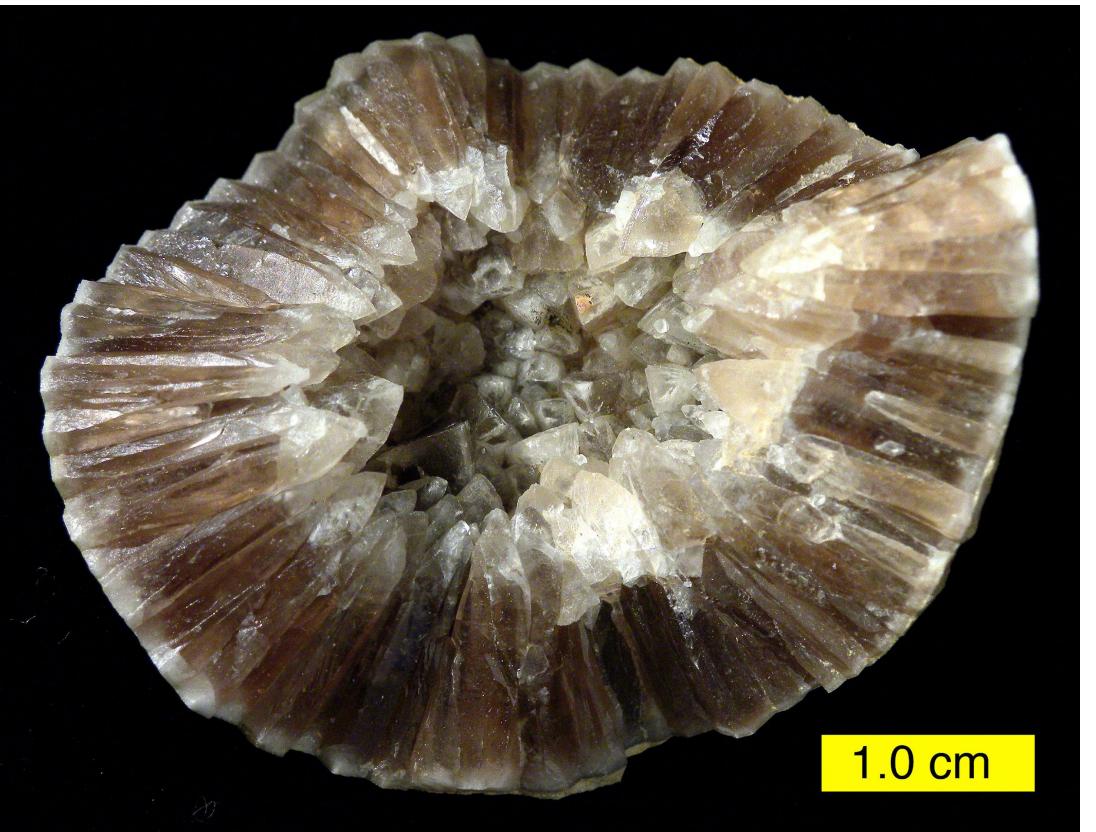
Origin of Meshes

- In nature, meshes arise in a variety of contexts:

- Cells in organic tissues

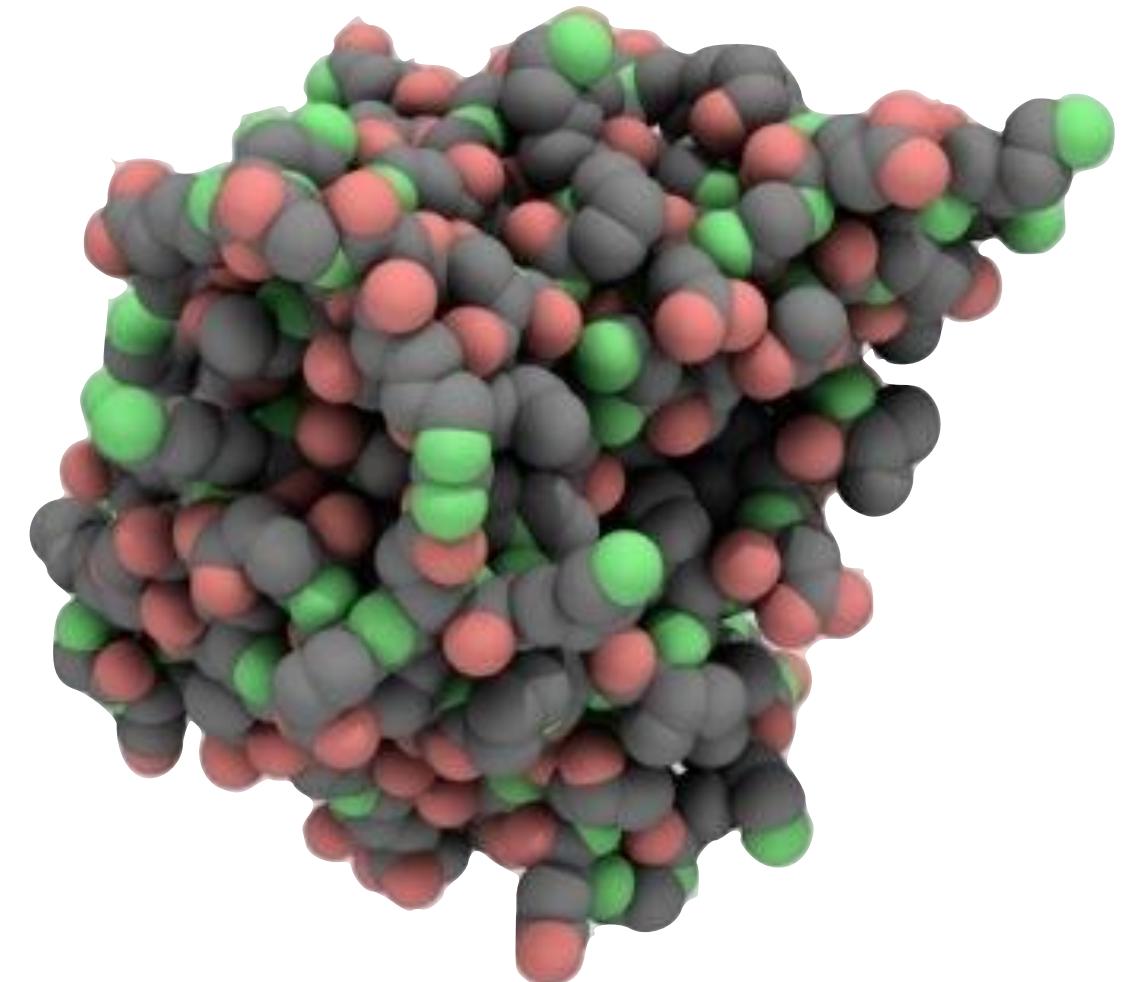


- Crystals



- Molecules

- Mostly *convex* but *irregular* cells



- Common concept: *complex* shapes can be described as *collections* of *simple building blocks*

Credit: Fernan Federici Moment Getty Images

By Mark A. Wilson (Department of Geology, The College of Wooster).[1] -
Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=12666593>

Basic Math of Meshes

- A n -cell is a set homeomorphic to a Euclidean disc of dimension n :

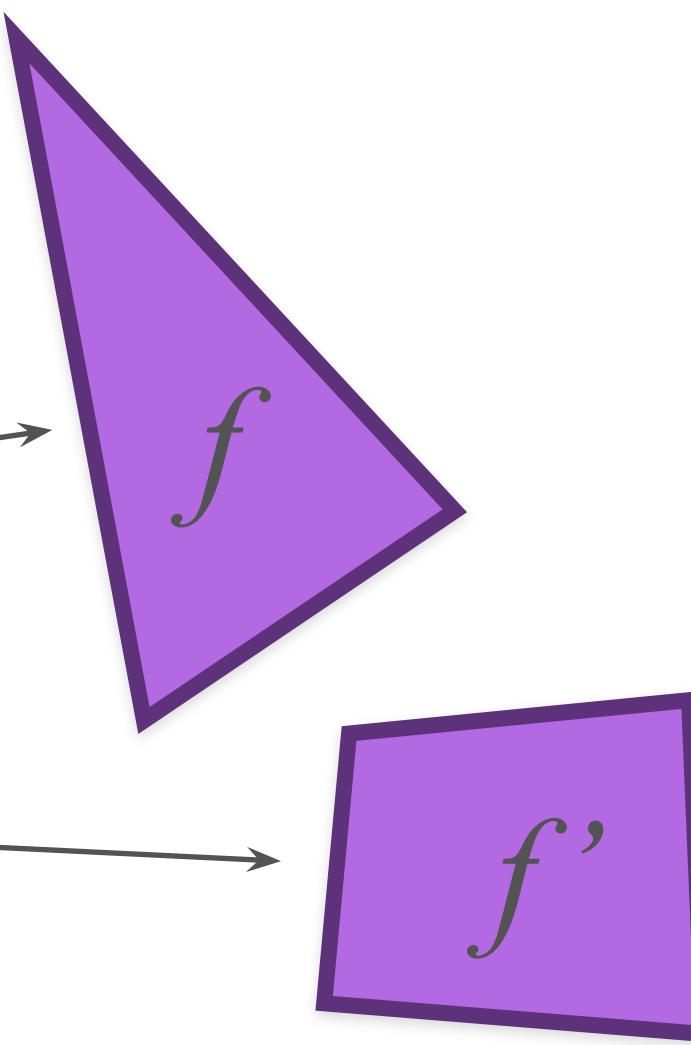
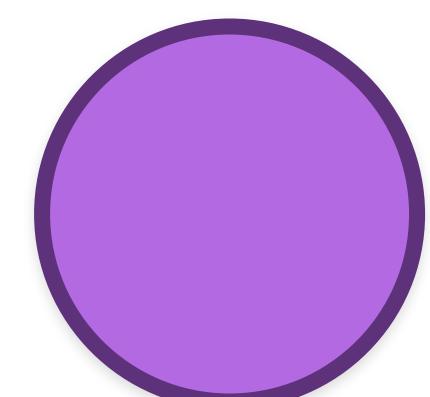
- 0-cell: *vertex*

v_\bullet

- 1-cell: *edge*

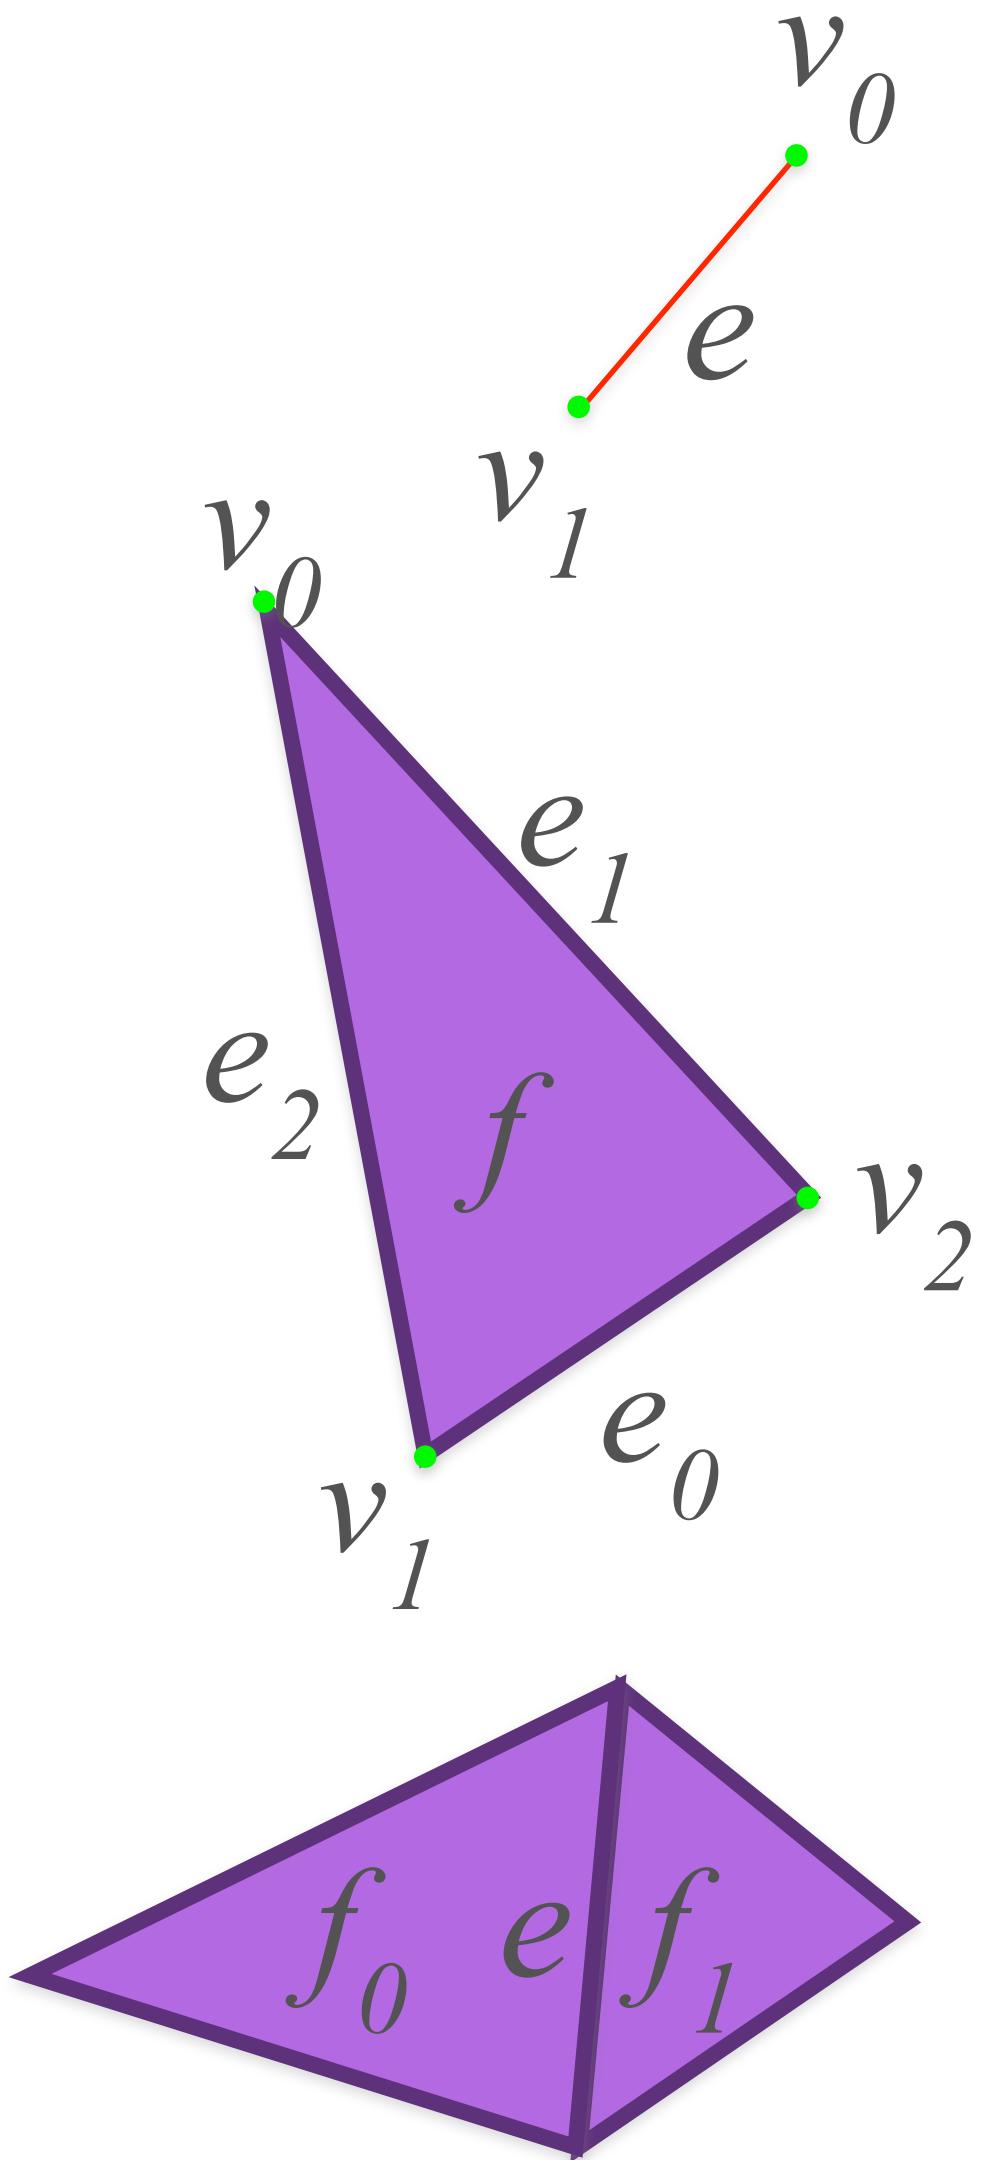
$[0,1] \rightarrow e$

- 2-cell: *face*



Structure

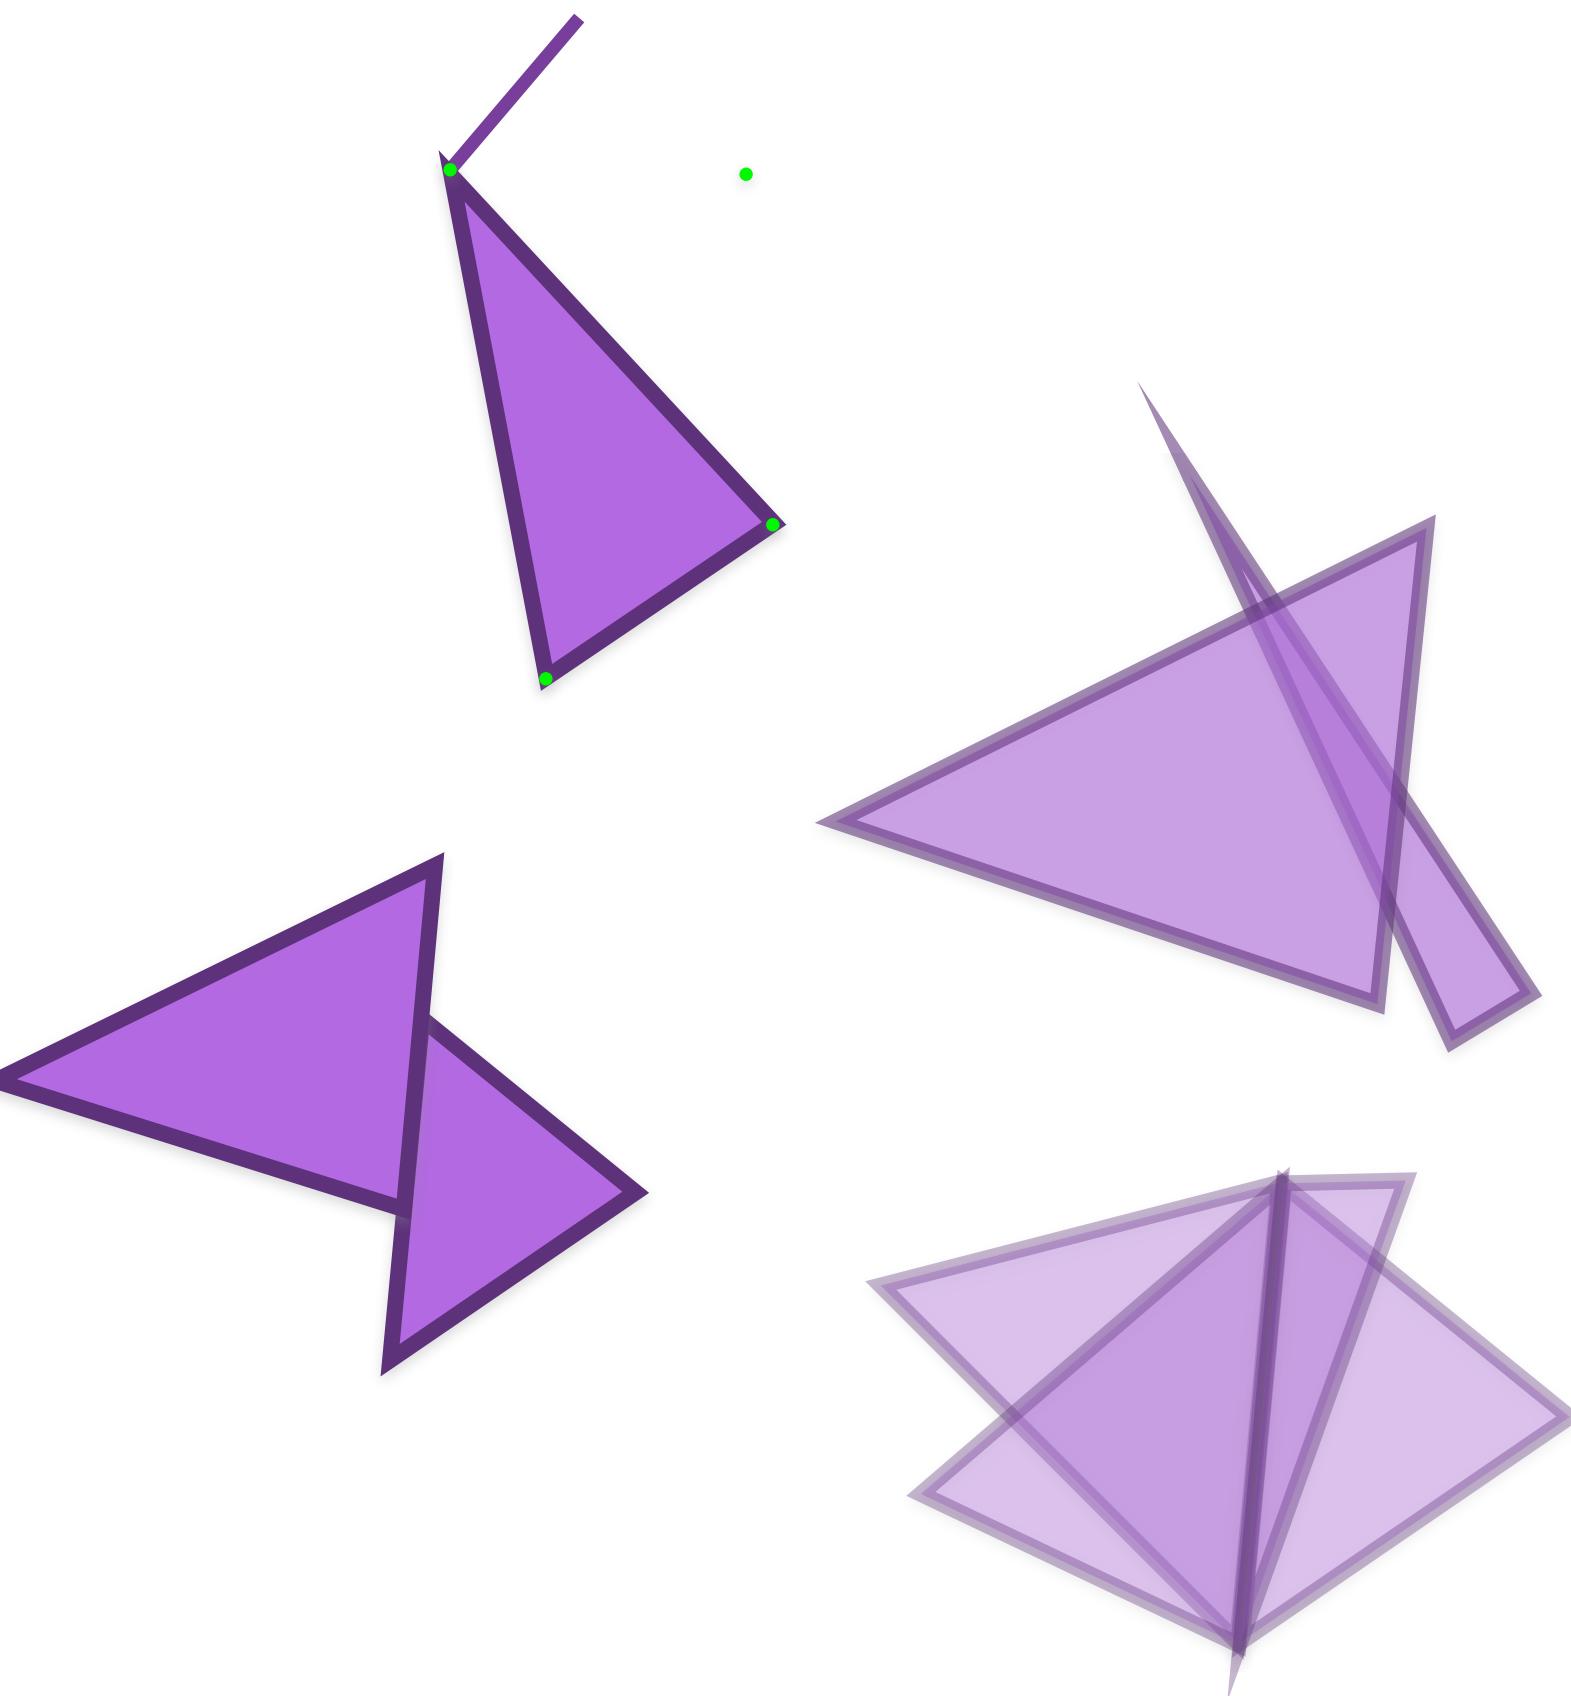
- A **mesh** $M=(V,E,F)$ of dimension 2 is made of a collection of k -cells for $k = 0, 1, 2$:
 - 0-cells of V lie on the boundary of 1-cells of E
 - 1-cells of E lies on the boundary of 2-cells of F
 - **(manifoldness)** each 1-cell of E lies on the boundary of either one or two 2-cells of F
 - the intersection of two distinct 1- / 2-cells is either empty or it coincides with a collection of 0- / 1-cells



Structure

Forbidden configurations:

- Dangling edges and isolated vertices
- Intersecting faces
- Non-conforming adjacency
- Non-manifold edges



Topological Informations

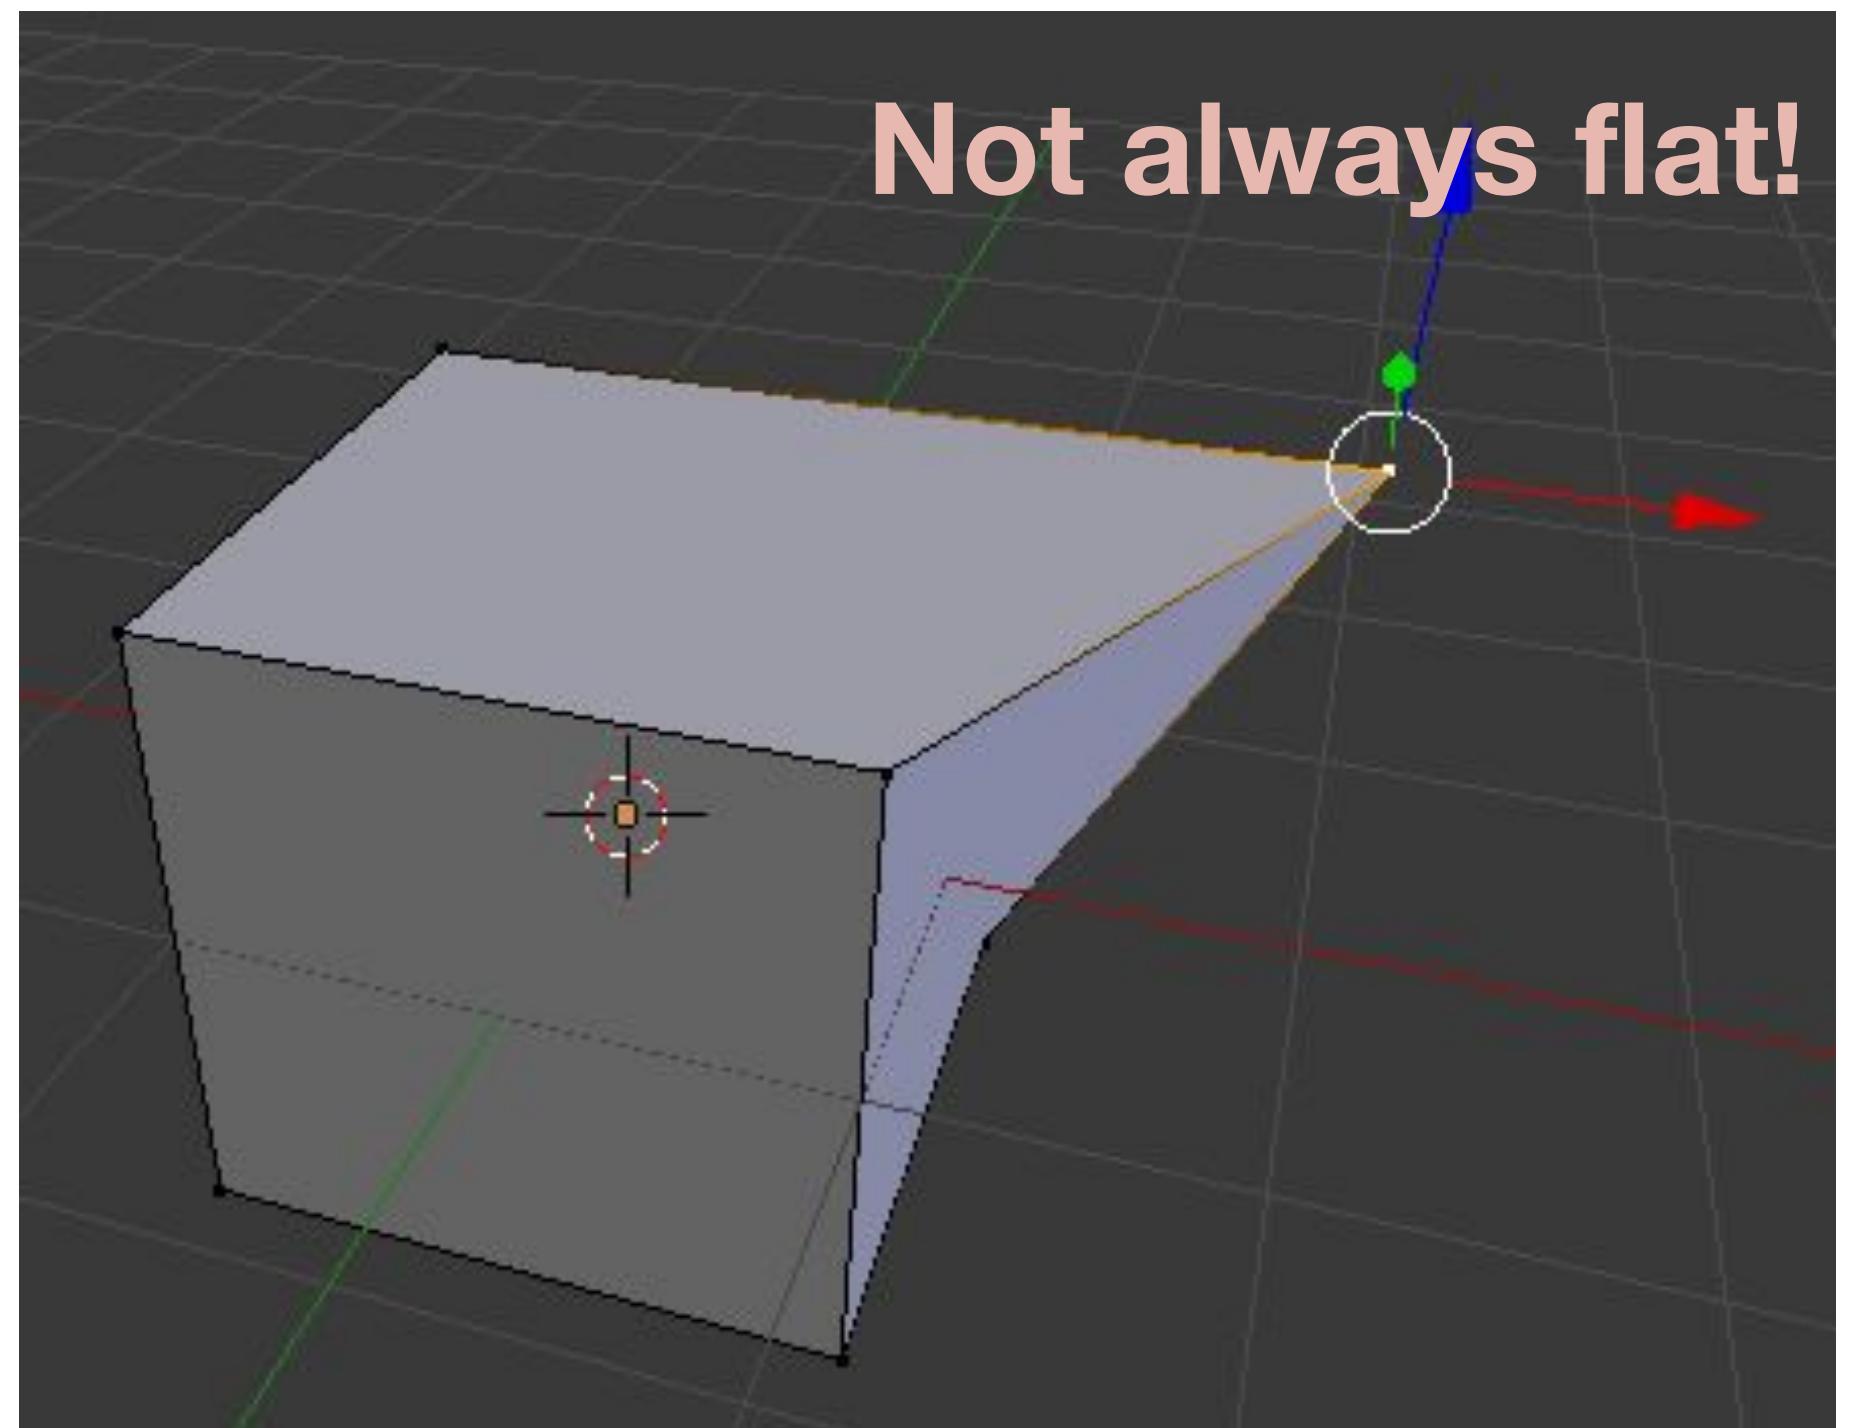
- A mesh can be treated as a purely combinatorial structure

$$M = (V, E, F)$$

- For some applications, geometry of edges and faces it not relevant. Just encode:
 - vertices as singletons (V)
 - how vertices are connected among them (E)
 - how cycles of vertices bound faces (F)

Geometrical Information

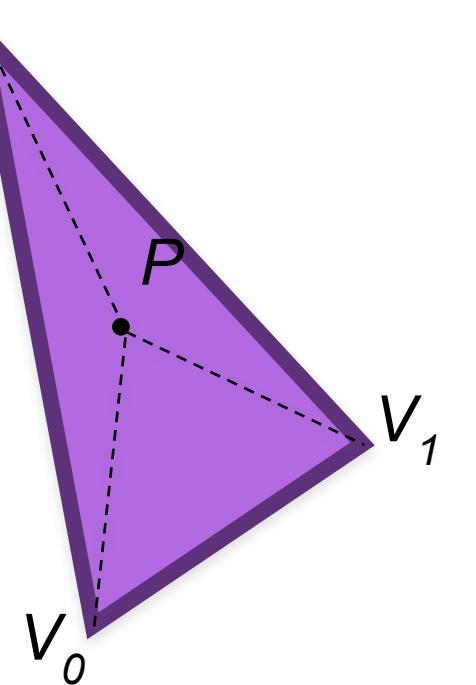
- Geometric embedding:
 - position in space for each 0-cell (vertex - point)
 - geometry for each 1-cell (edge - line) and 2-cell (face - disk-like surface)
- Polygonal meshes are embedded:
 - edges are straight-line segments
 - are faces flat? not always true:
vertices of a face might be *not coplanar*



Triangle Meshes

- A *triangle mesh* is a polygonal mesh with all triangular faces
 - all faces are flat (there exist a unique plane for three points)
- All cells are *simplices*, i.e., they are the convex combinations of their vertices

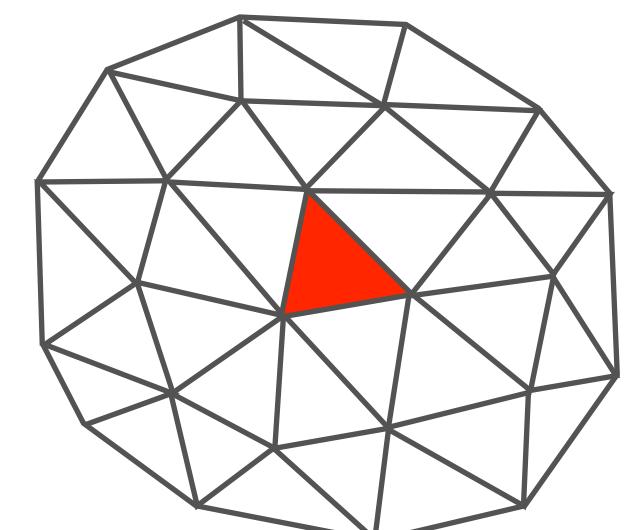
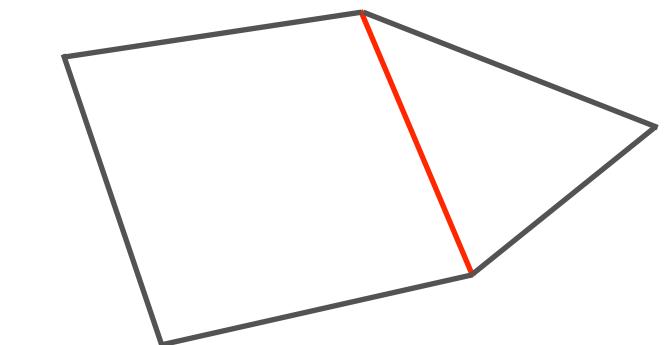
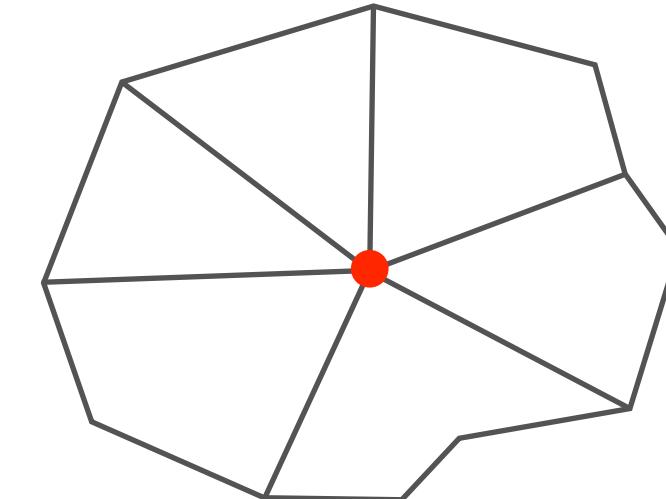
$$P = \lambda_0 V_0 + \lambda_1 V_1 + \lambda_2 V_2 \quad \lambda_i \in [0,1] \quad \lambda_0 + \lambda_1 + \lambda_2 = 1$$



- embedding of vertices + combinatorial structure characterize the embedding of the whole mesh

Stars and Rings

- The star of a vertex v is formed by v plus the set of cells incident at v (edges and faces of its co-boundary)
- The star of an edge e is formed by e plus the set of faces incident at e (faces of its co-boundary)
- The 1-ring of a face f is the formed by the union of the stars of its boundary vertices
- The k -ring of a face f , for $k>1$ is the formed by the union of the 1-rings of faces in its $(k-1)$ -ring



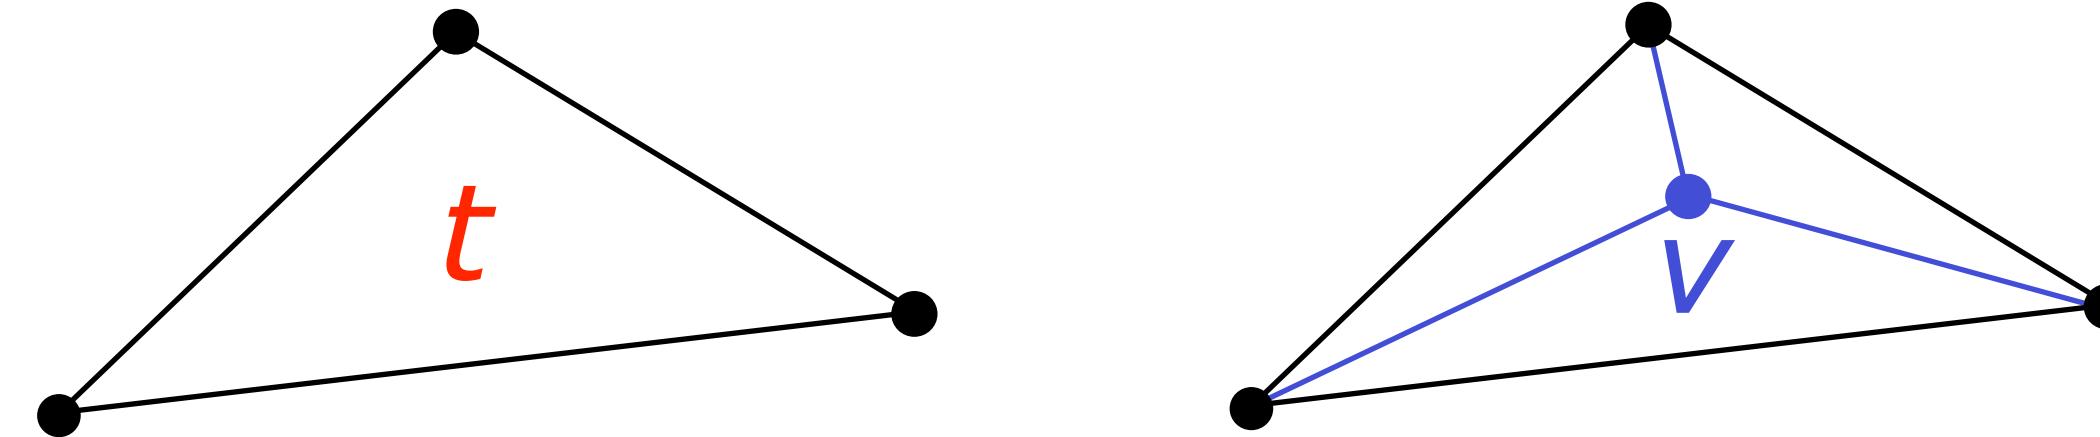
Editing Operators

Operators for Triangle Meshes

- Specific operators for triangle meshes include
- **Refinement operators**: produce a mesh with more vertices/edges/faces
- **Simplification operators**: produce a mesh with less vertices/edges/faces
- Can be implemented on any topological data structure for triangle meshes

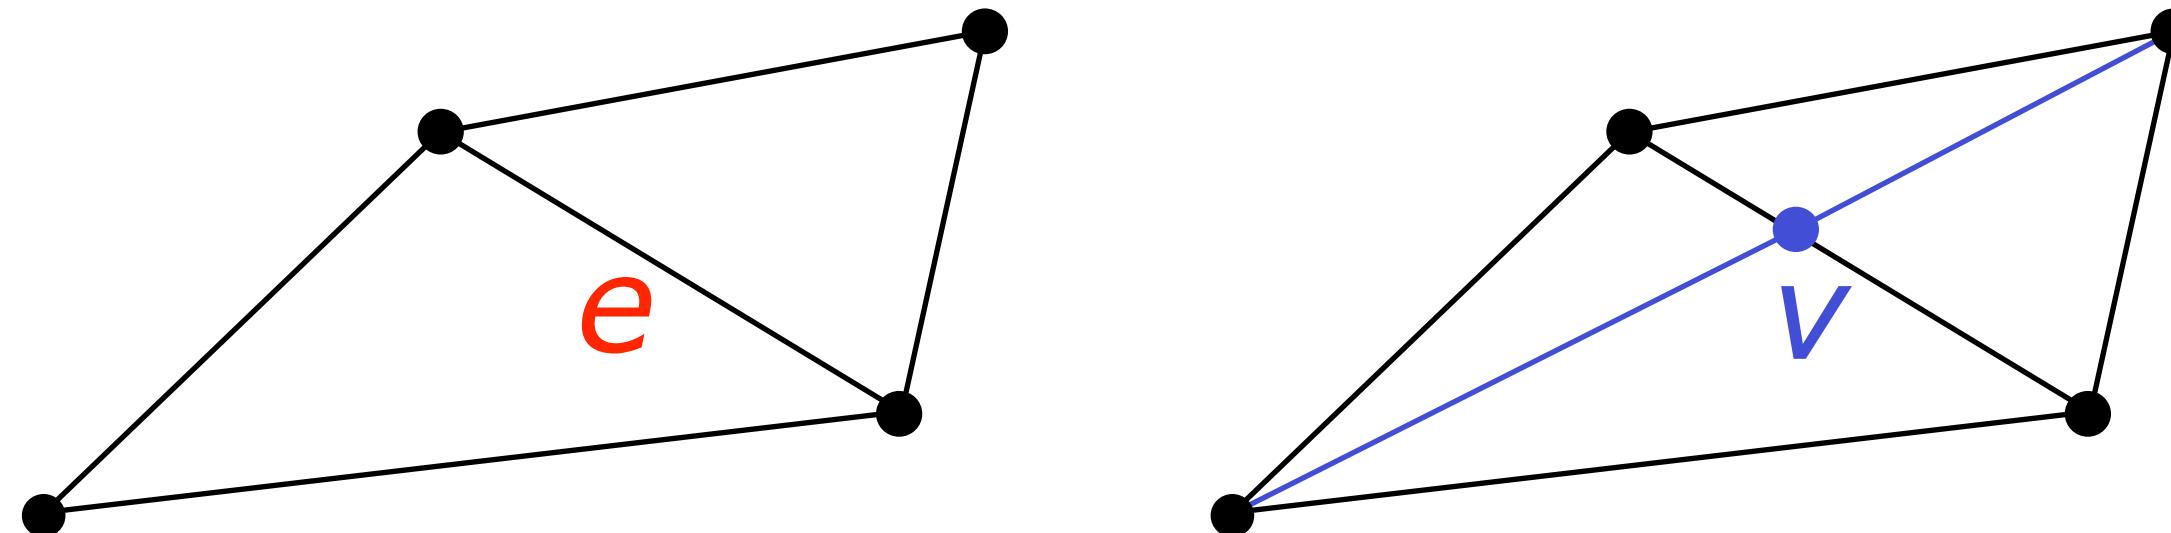
Refinement Operators

- Triangle split:
 - insert a new vertex v in a triangle t and connect v to the vertices of t by splitting it into three triangles



Refinement Operators

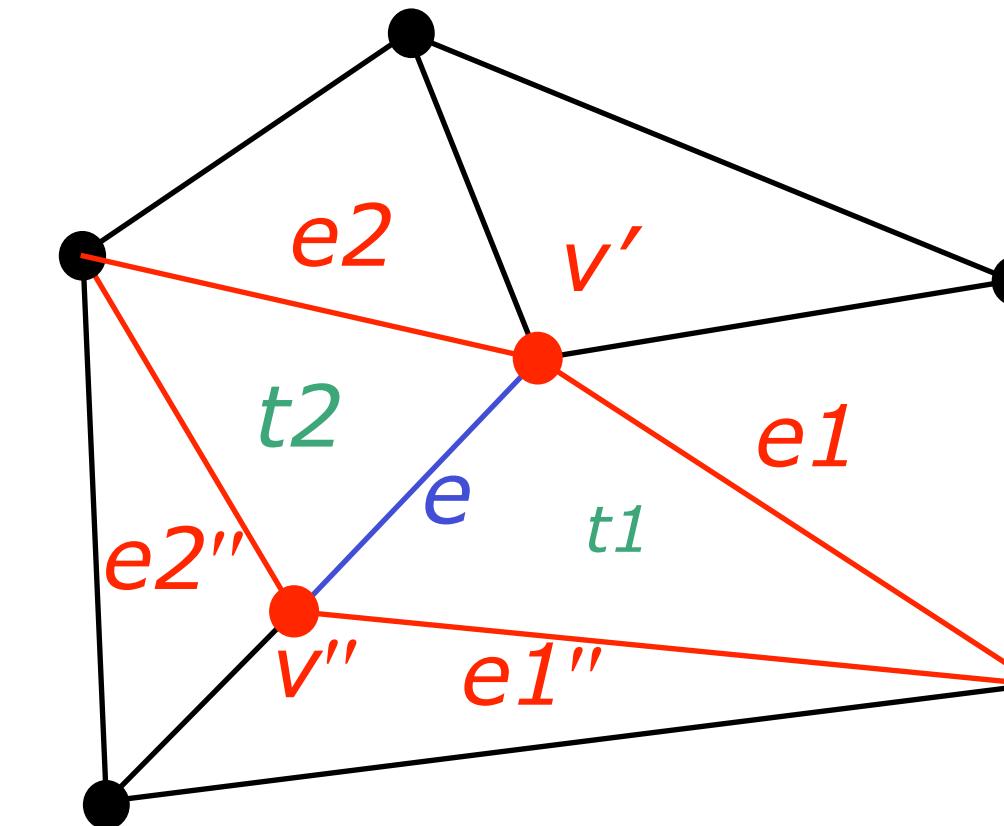
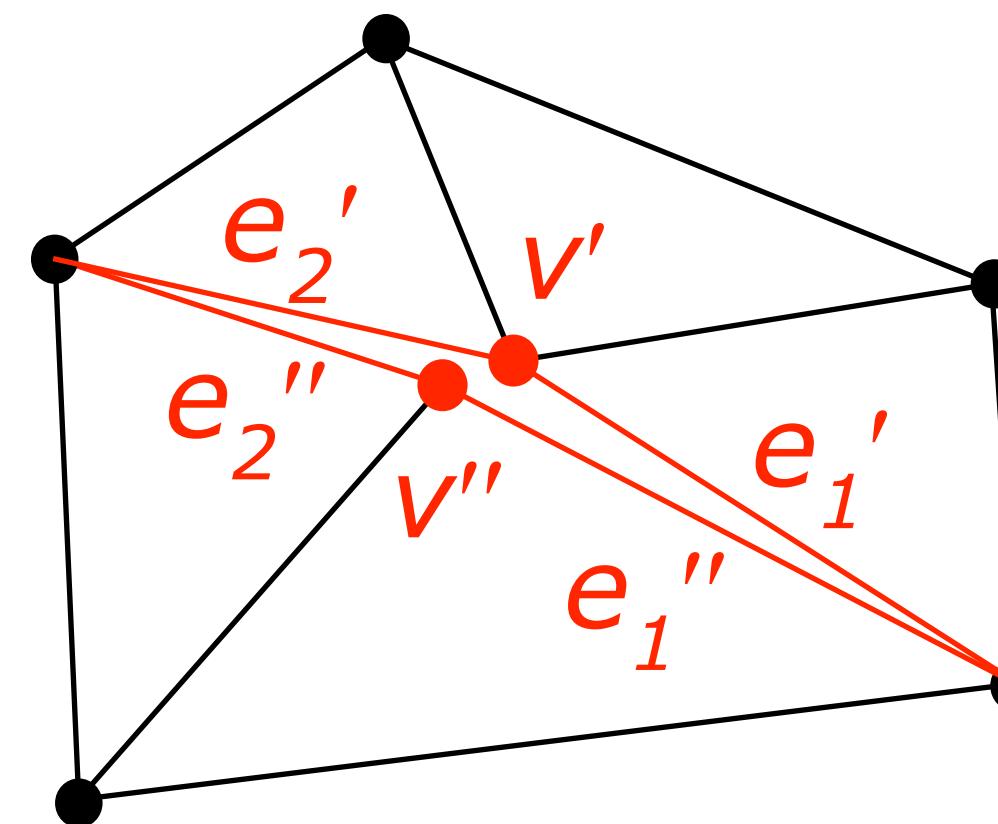
- Edge split:
 - insert a new vertex v on an edge e and connect v to the opposite vertices of triangles incident at e by splitting e as well as each such triangle into two



Refinement Operators

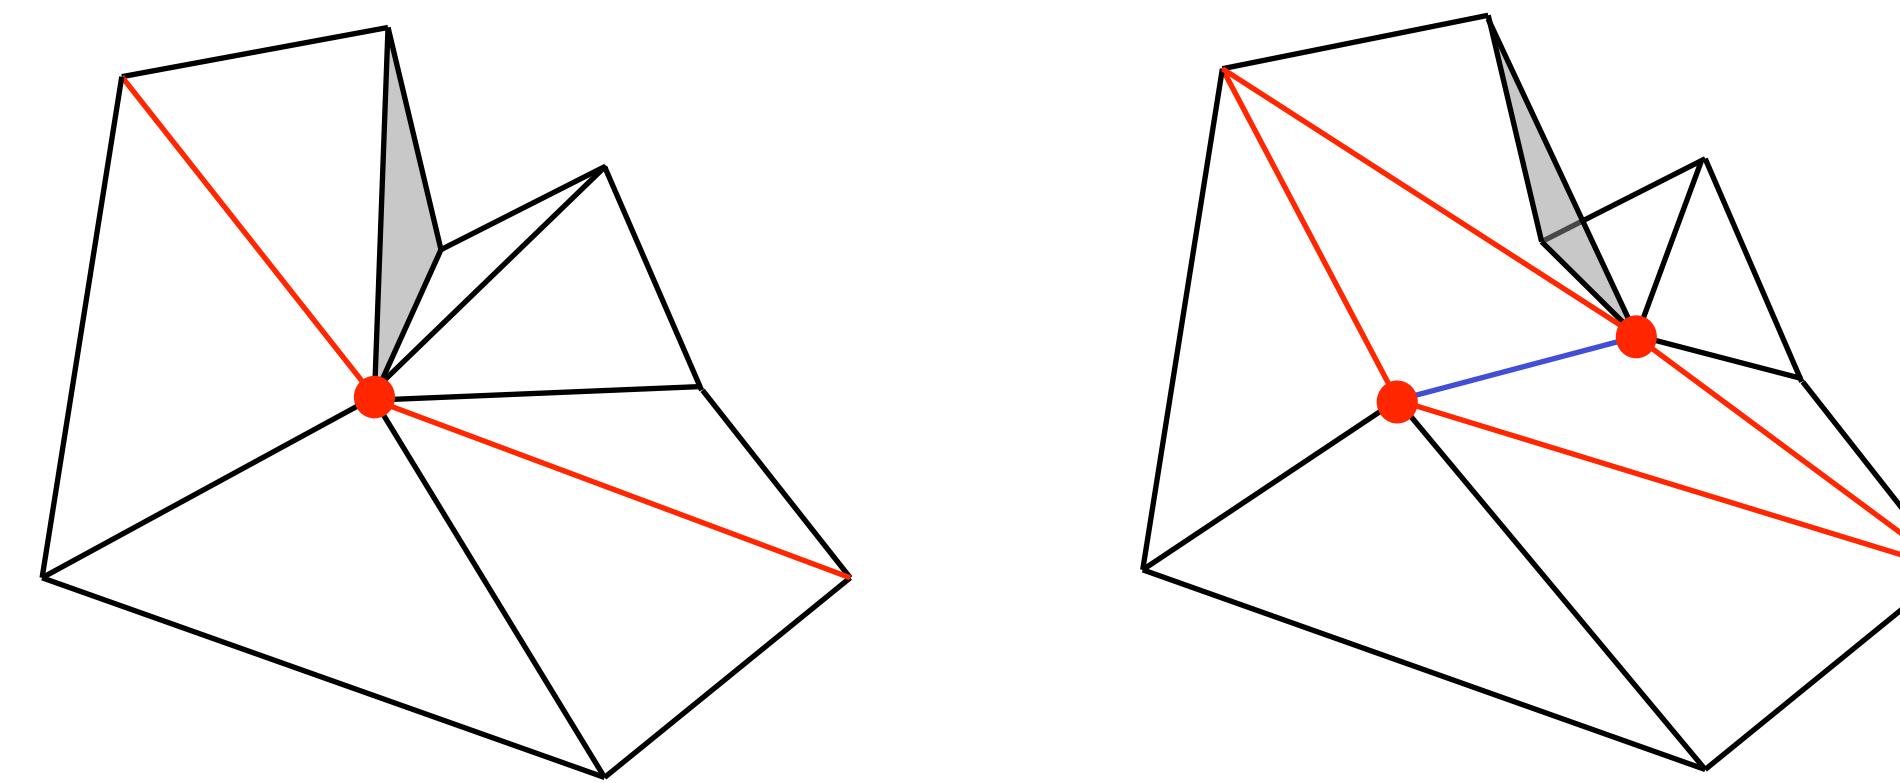
- Vertex split:

- cut open the mesh along two edges e_1 and e_2 incident at a common vertex v , by duplicating such edges as well as v
- fill the quadrangular hole with two new triangles and an edge joining the two copies of v



Refinement Operators

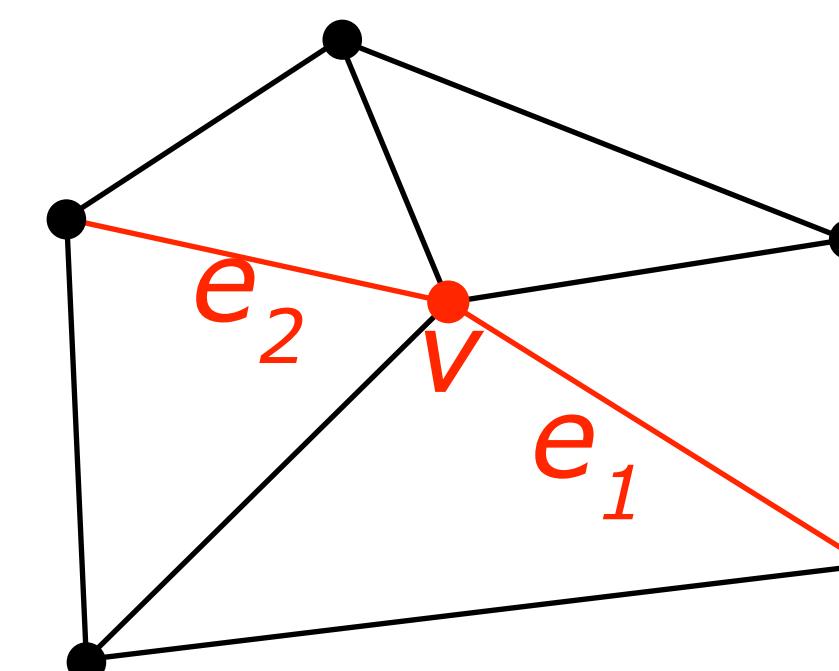
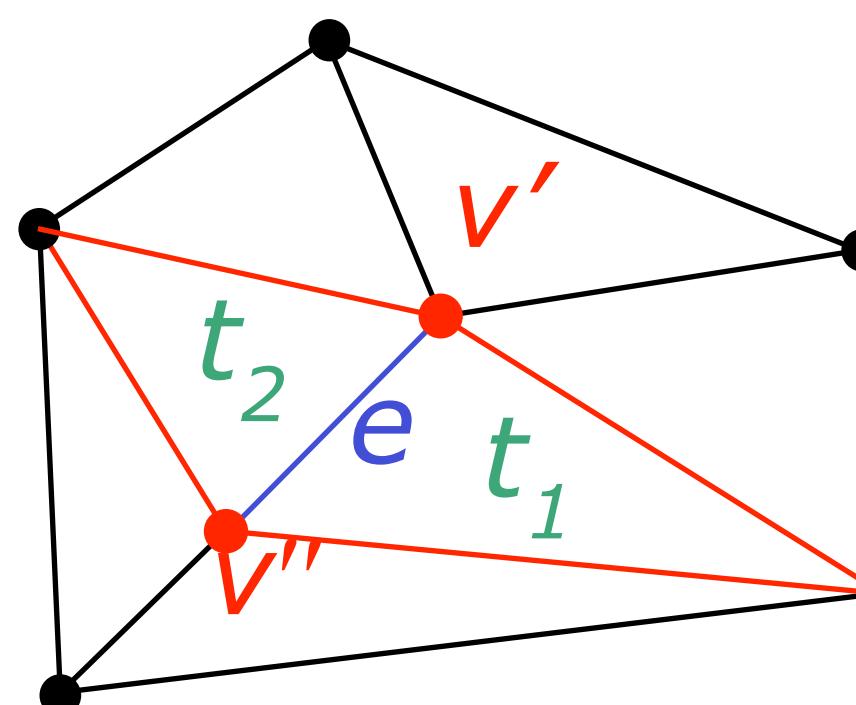
- Vertex split:
 - possible inconsistencies because of *triangle flip*



- flips can be detected by a local test on the orientation of faces: flips changes orientation from clockwise to counter-clockwise and vice-versa

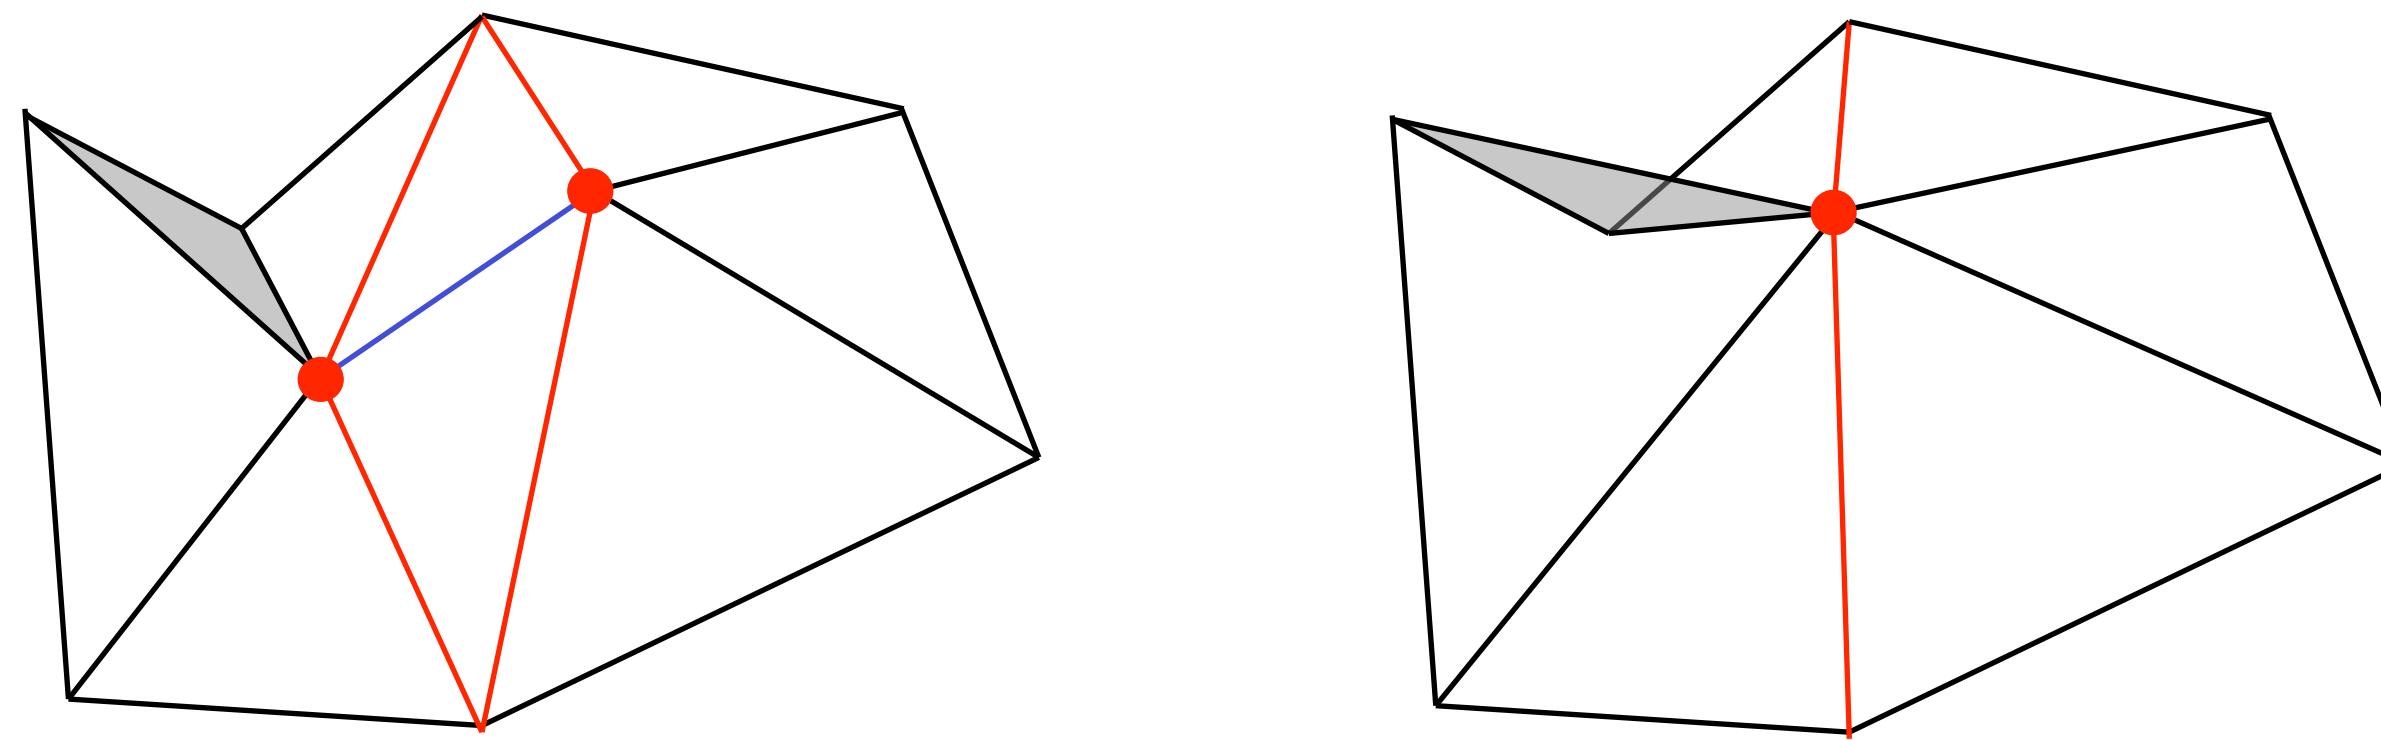
Simplification Operators

- Edge collapse (reverse of vertex split):
 - collapse an edge e to a single point
 - e is removed together with its two incident triangles
 - the endpoints of e are identified
 - the other edges bounding the deleted triangles are pairwise identified



Simplification Operators

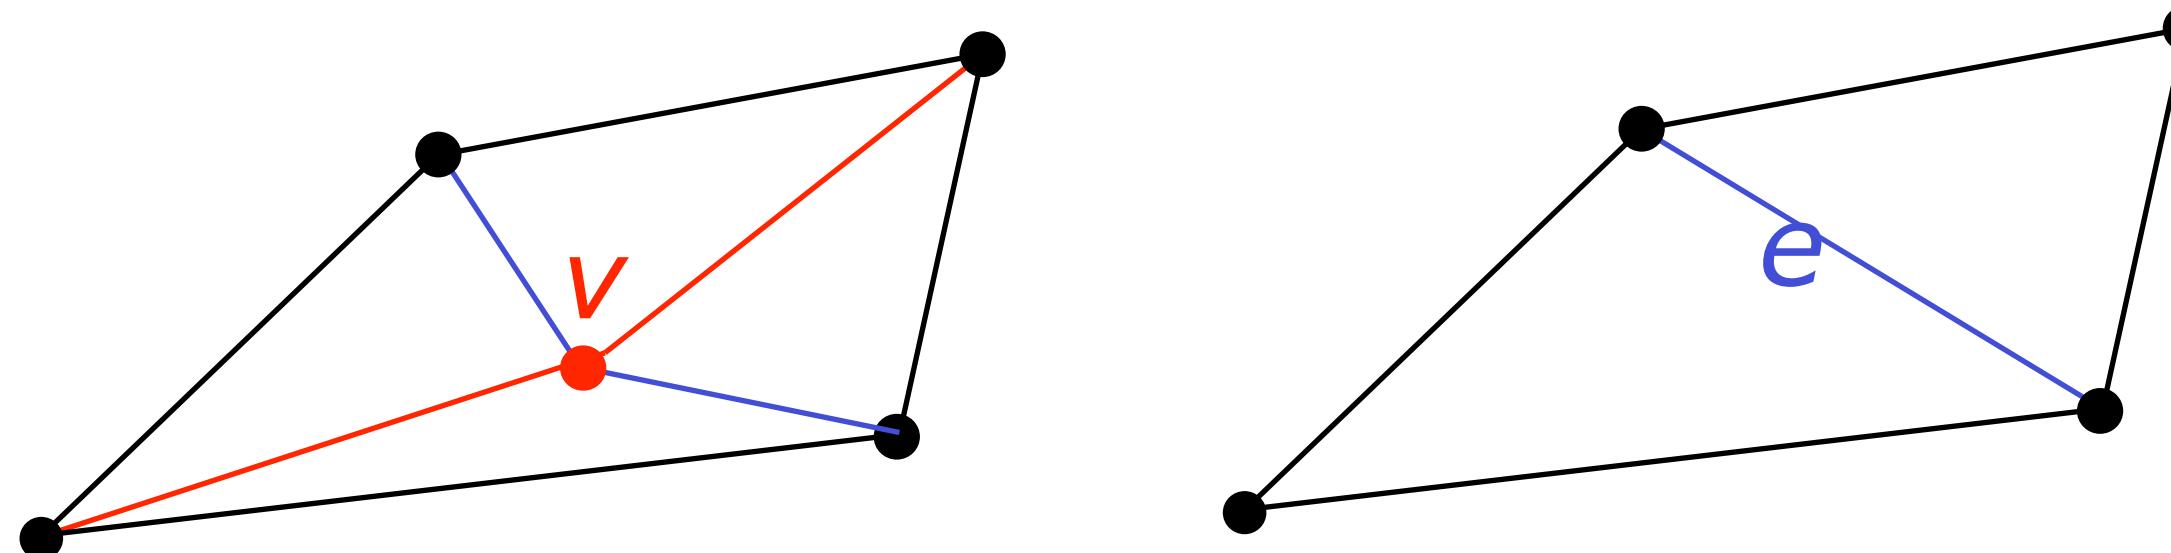
- Edge collapse:
 - possible inconsistencies because of *triangle flip*



- consistency check analogous to vertex split

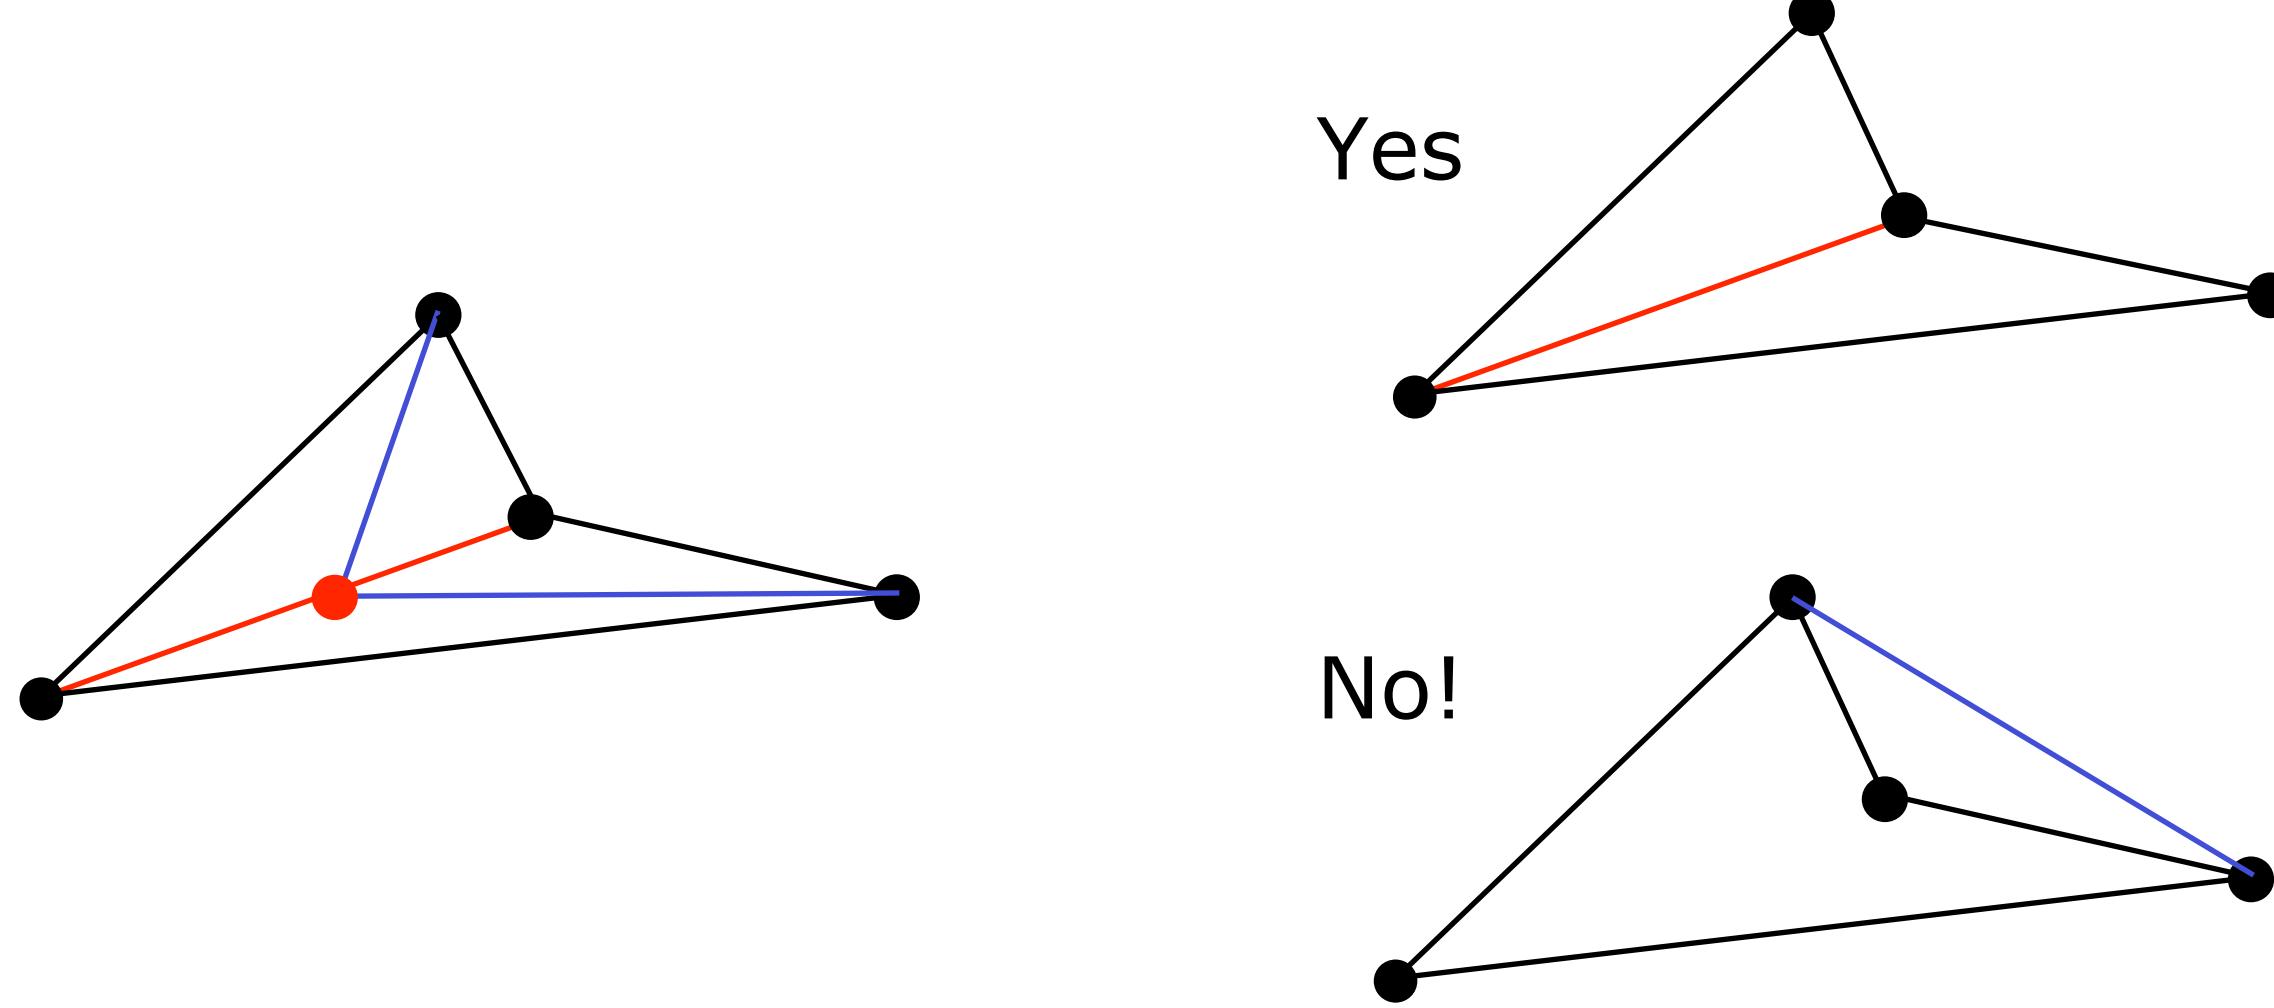
Simplification Operators

- Edge merge (reverse of edge split):
 - take an internal vertex v with valence 4
 - delete v together with its incident triangles and edges and fill the hole with two new triangles sharing a new edge e



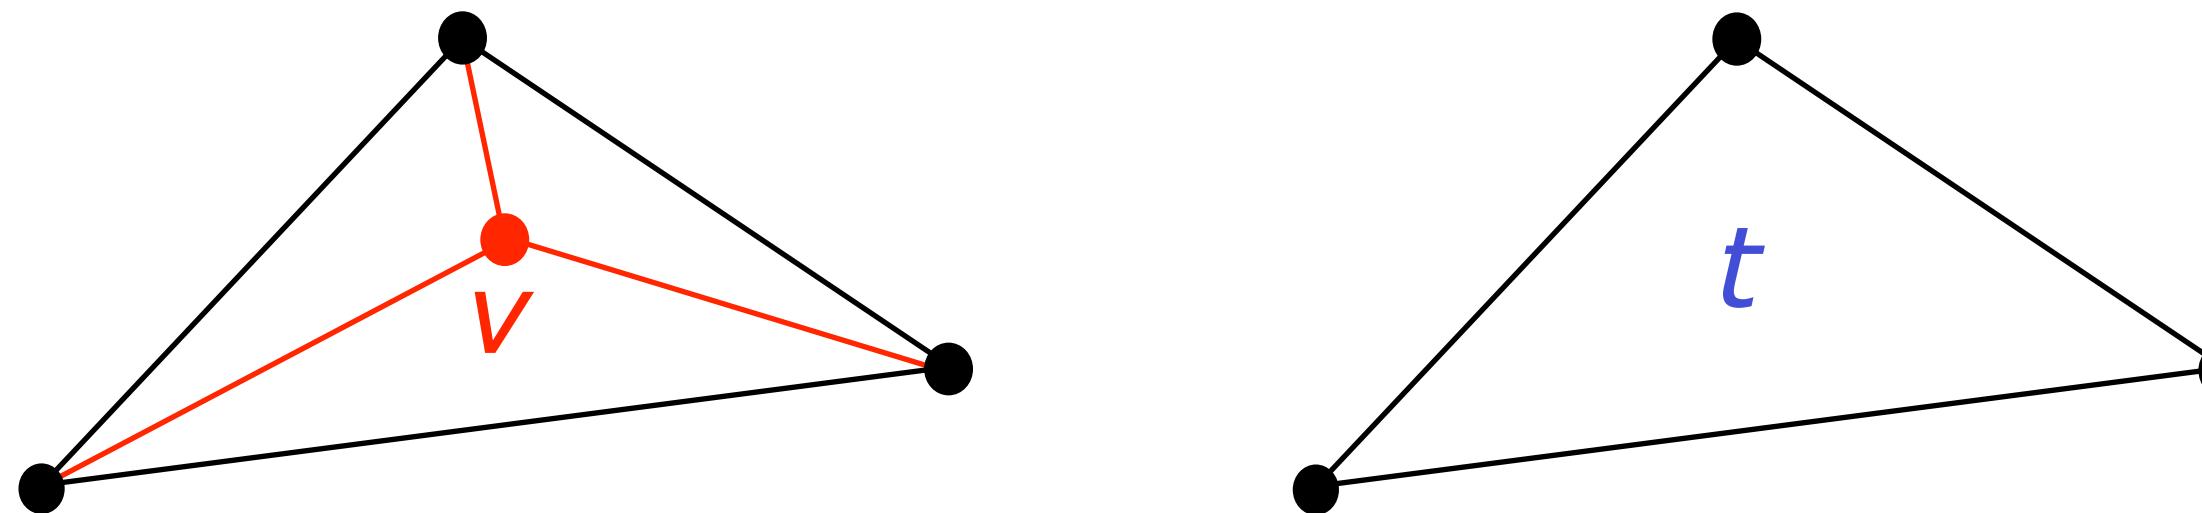
Simplification Operators

- Edge merge:
 - if the hole is not convex, only one diagonal edge can be inserted



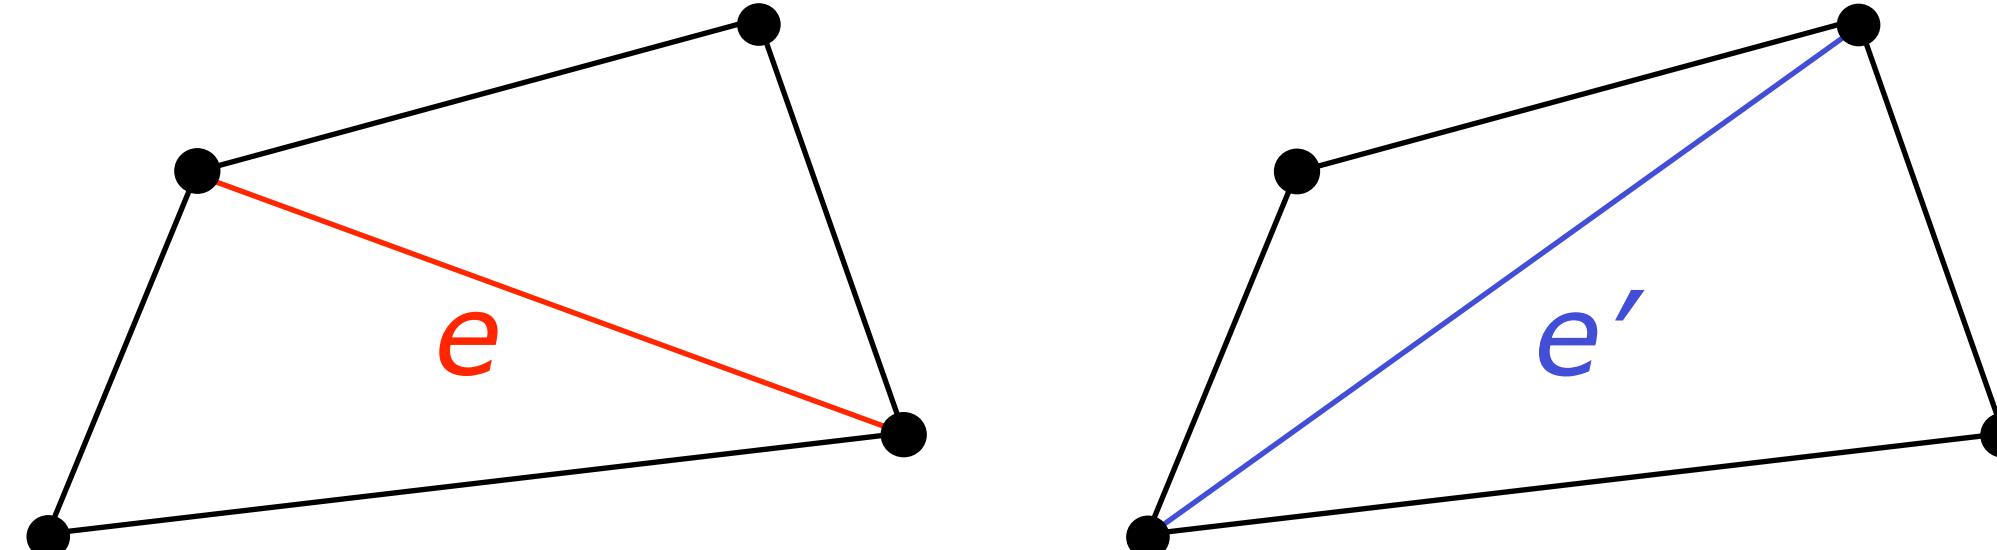
Simplification Operators

- Delete vertex (reverse of triangle split):
 - remove an internal vertex v of valence 3 together with its incident triangles and edges and fill the hole with a new triangle



Neutral Operator

- Edge swap:
 - consider an edge e such that its two incident triangles form a convex quadrilateral
 - replace e with the opposite diagonal of the quadrilateral, rearranging the two incident triangles accordingly



Mesh Data Structures

Data Structures

- Storing a mesh:
 - geometry: position of vertices
 - connectivity: edges, faces
 - topology: topological relations
- Compactness vs efficiency trade-off: data structures should be compact, while efficiently supporting time-critical operations
 - storage requirements
 - traversal operations
 - update operations

Polygon Soup

- a.k.a. Face set - STL files
- Just store a list of faces
- For each face, store positions of its vertices
 - For general polygonal mesh, the number of vertices of each face must also be stored
 - Number implicit for triangle meshes
- For a triangle mesh: 36 bytes/face \approx 72 bytes/vertex
- No connectivity! just a collection of polygons
- Streaming structure: no need to store it all in memory to render the mesh

Triangles		
$x_{11} \ y_{11} \ z_{11}$	$x_{12} \ y_{12} \ z_{12}$	$x_{13} \ y_{13} \ z_{13}$
$x_{21} \ y_{21} \ z_{21}$	$x_{22} \ y_{22} \ z_{22}$	$x_{23} \ y_{23} \ z_{23}$
...
$x_{F1} \ y_{F1} \ z_{F1}$	$x_{F2} \ y_{F2} \ z_{F2}$	$x_{F3} \ y_{F3} \ z_{F3}$

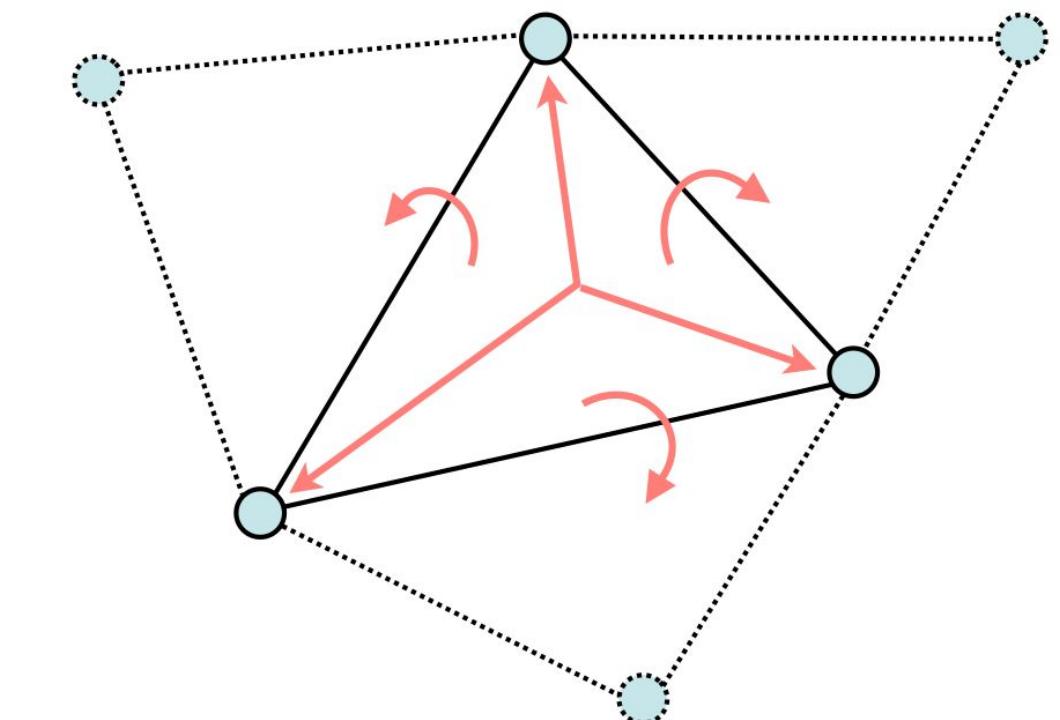
Indexed Structure

- Avoid replication of vertices - OBJ, OFF, PLY files
- Maintain a vector of vertices and a vector of faces
 - For each vertex: position
 - For each face: references to positions of its vertices (FV) in the vector
 - in general, the number of vertices of each face must also be stored
 - number implicit for triangle meshes
- For a triangle mesh: 12 bytes/vertex +12 bytes/face \approx 36 bytes/vertex
- Encodes connectivity through the FV relation
- Main memory structure: need to store in memory at least the whole list of vertices

Vertices	Triangles
$x_1 \ y_1 \ z_1$	$v_{11} \ v_{12} \ v_{13}$
...	...
$x_v \ y_v \ z_v$	$v_{F1} \ v_{F2} \ v_{F3}$
...	...
...	...
...	...

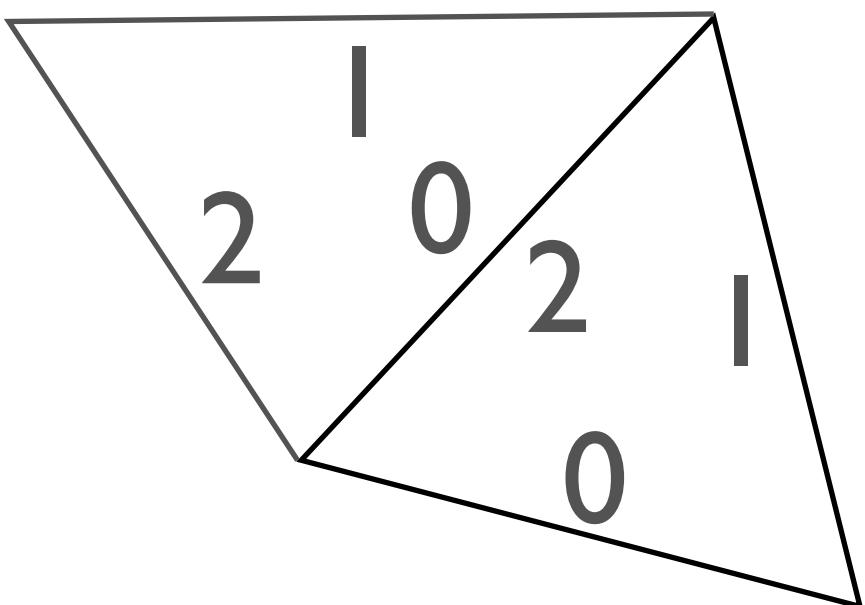
Indexed Structure with Adjacencies

- Just for triangle meshes
- Extends indexed structure with some topological relations:
 - For each vertex:
 - position
 - one reference to an incident triangle (VF^* relation)
 - For each face:
 - references to its three vertices (FV)
 - references to its three adjacent triangles (FF)



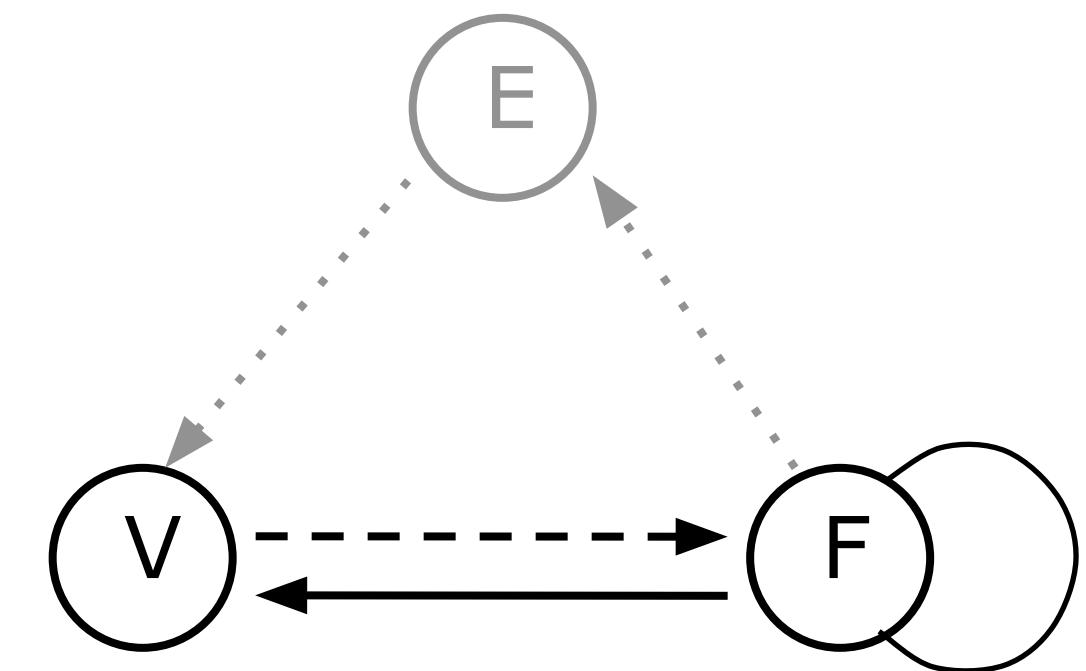
Indexed Structure with Adjacencies

- No explicit edges!
 - implicitly defined as pairs of vertices (unique)
 - or as pairs (f, i) with f triangle and i index in 0,1,2 (not unique!)
- Attributes for edges require care, especially when modifying the mesh with editing operators



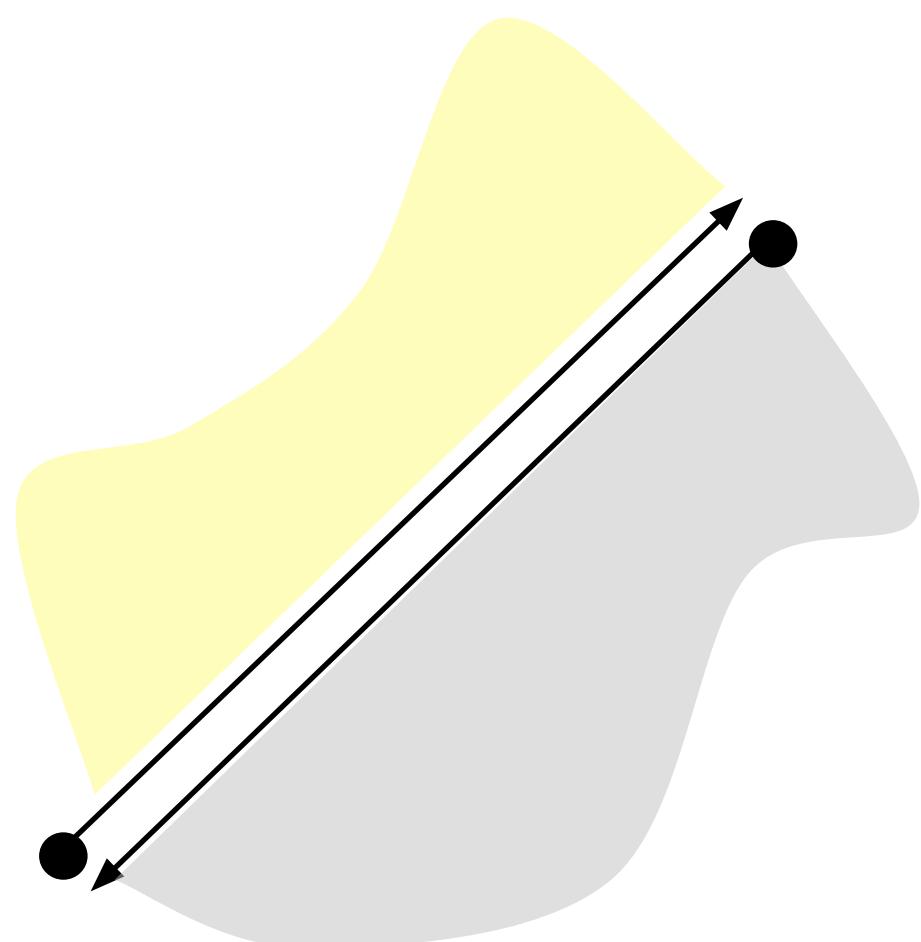
Indexed Structure with Adjacencies

- Evaluation of topological relations:
 - FV, FF: encoded - optimal
 - VF = VF* + FF + FV optimal
 - WV analogous to VF
- Relations involving edges are either implicit or they can be evaluated in optimal time by the (f,i) encoding



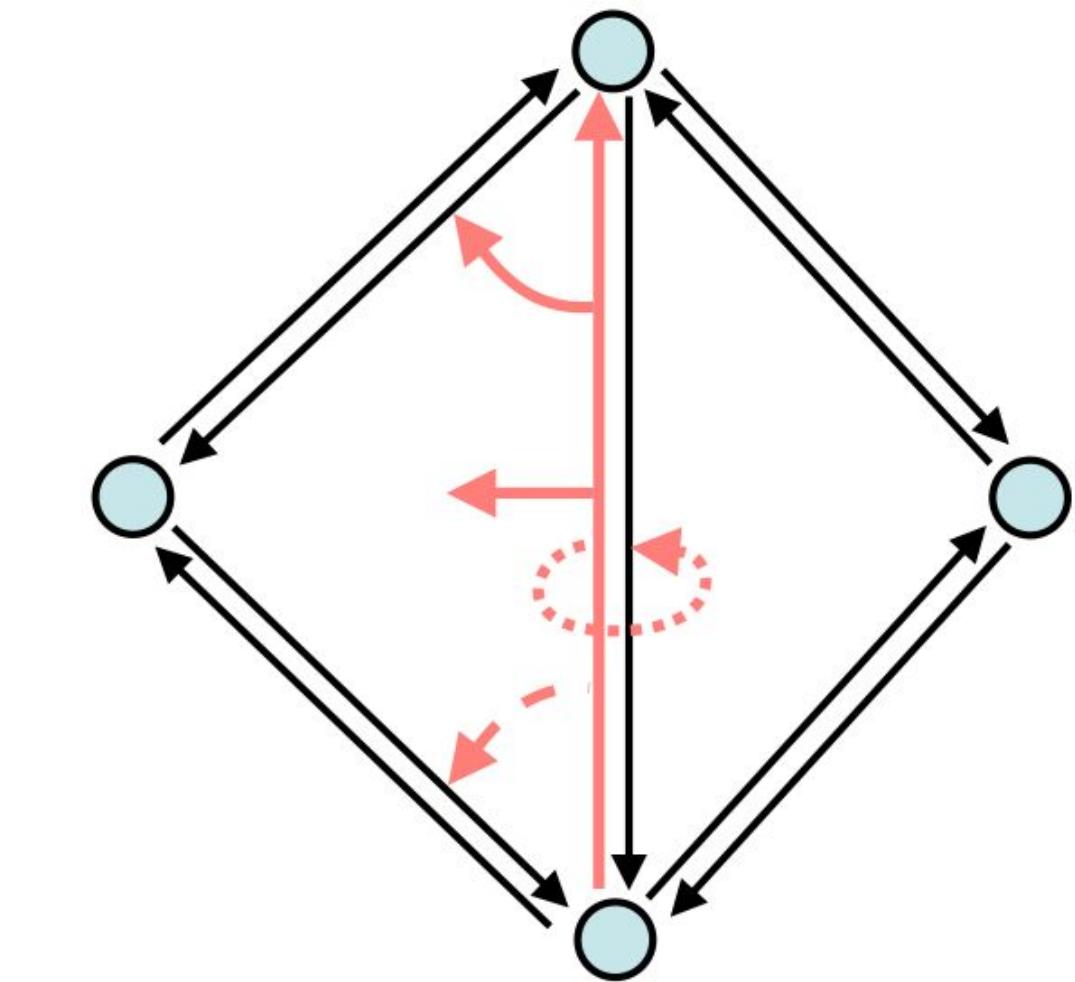
Half-Edge Data Structure

- Half-edge: each edge is duplicated by also considering its orientation
- An edge corresponds to a pair of sibling half-edges with opposite orientations
- Each half-edge stores half topological information concerning the edge



Half-Edge Data Structure

- For each vertex:
 - position
 - one reference to an incident half-edge
- For each half-edge:
 - reference to one its two vertices (origin)
 - references to one of its two incident faces (face on its left)
 - references to: next/previous half-edge on the same face, sibling half-edge
- For each face:
 - one reference to an incident half-edge (FE^*)

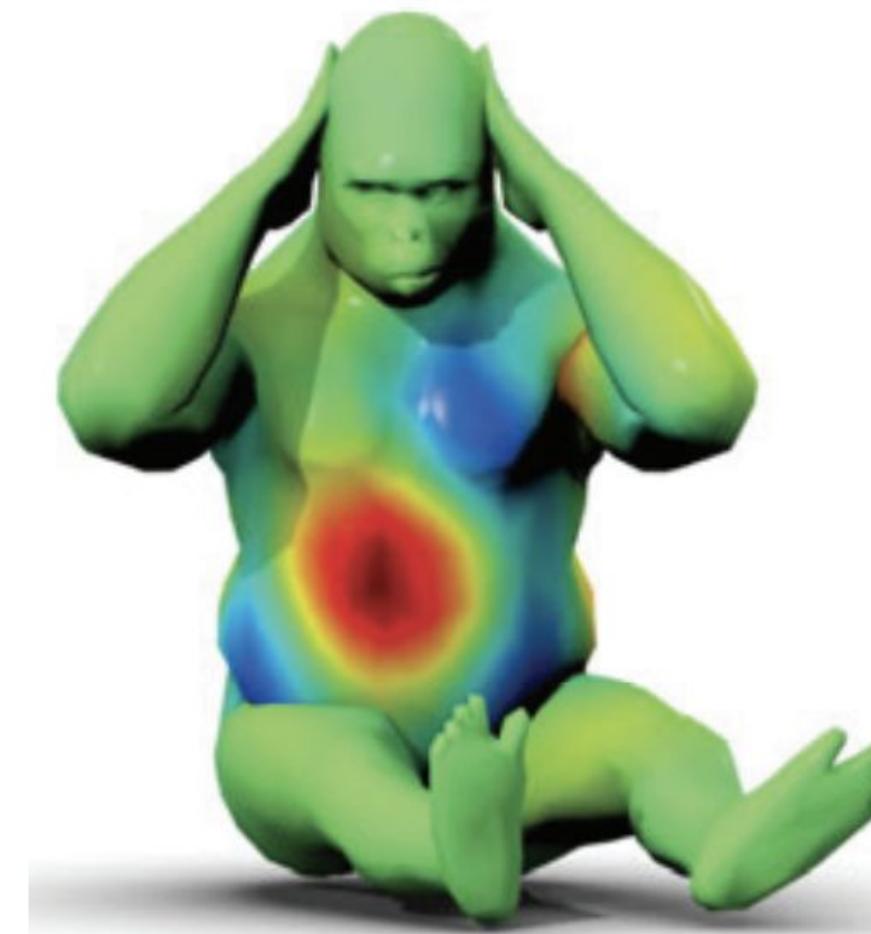


Half-Edge Data Structure

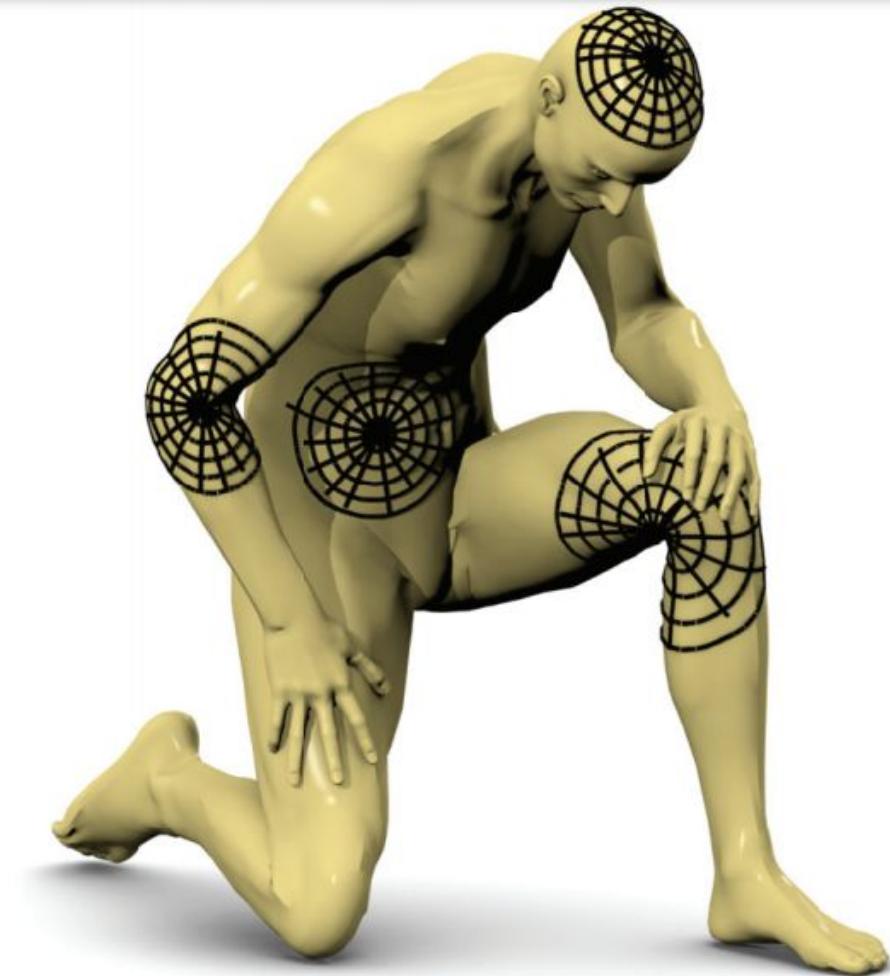
- 96-144 bytes/vertex depending on number of references to adjacent edges
 - reference to sibling half-edge can be avoided by storing siblings at consecutive entries of a vector
 - for triangle meshes, just one reference to either next or previous half-edge is sufficient
- Efficient traversal and update operations
- Attributes for edges must be stored separately

Mesh Neural Networks

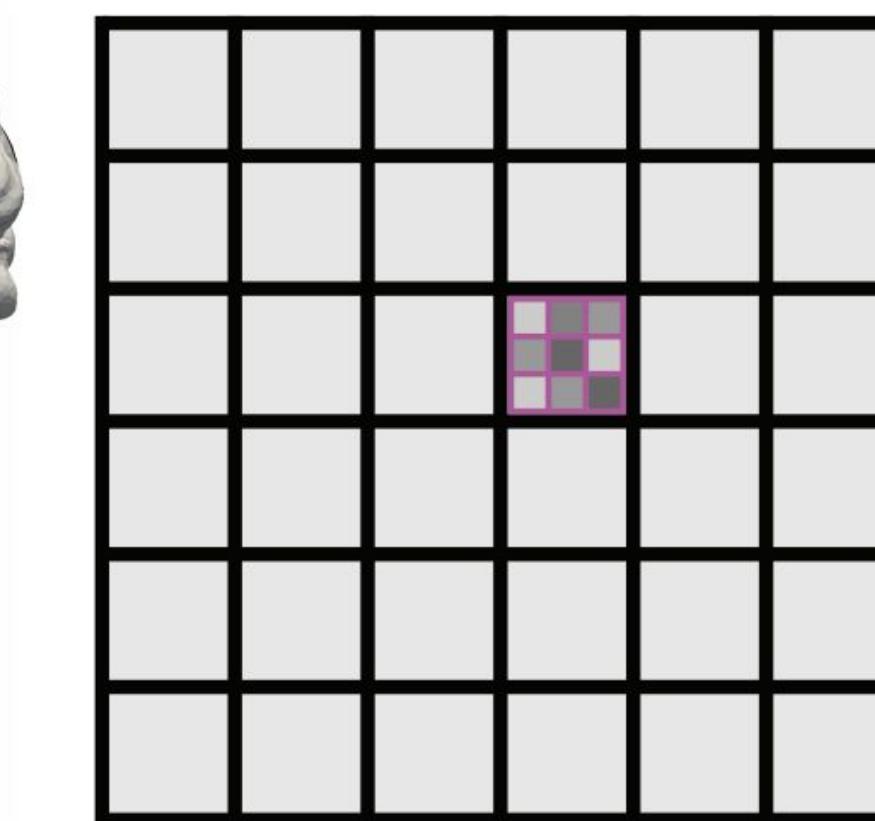
Formulations of non-Euclidean CNNs



Spectral domain

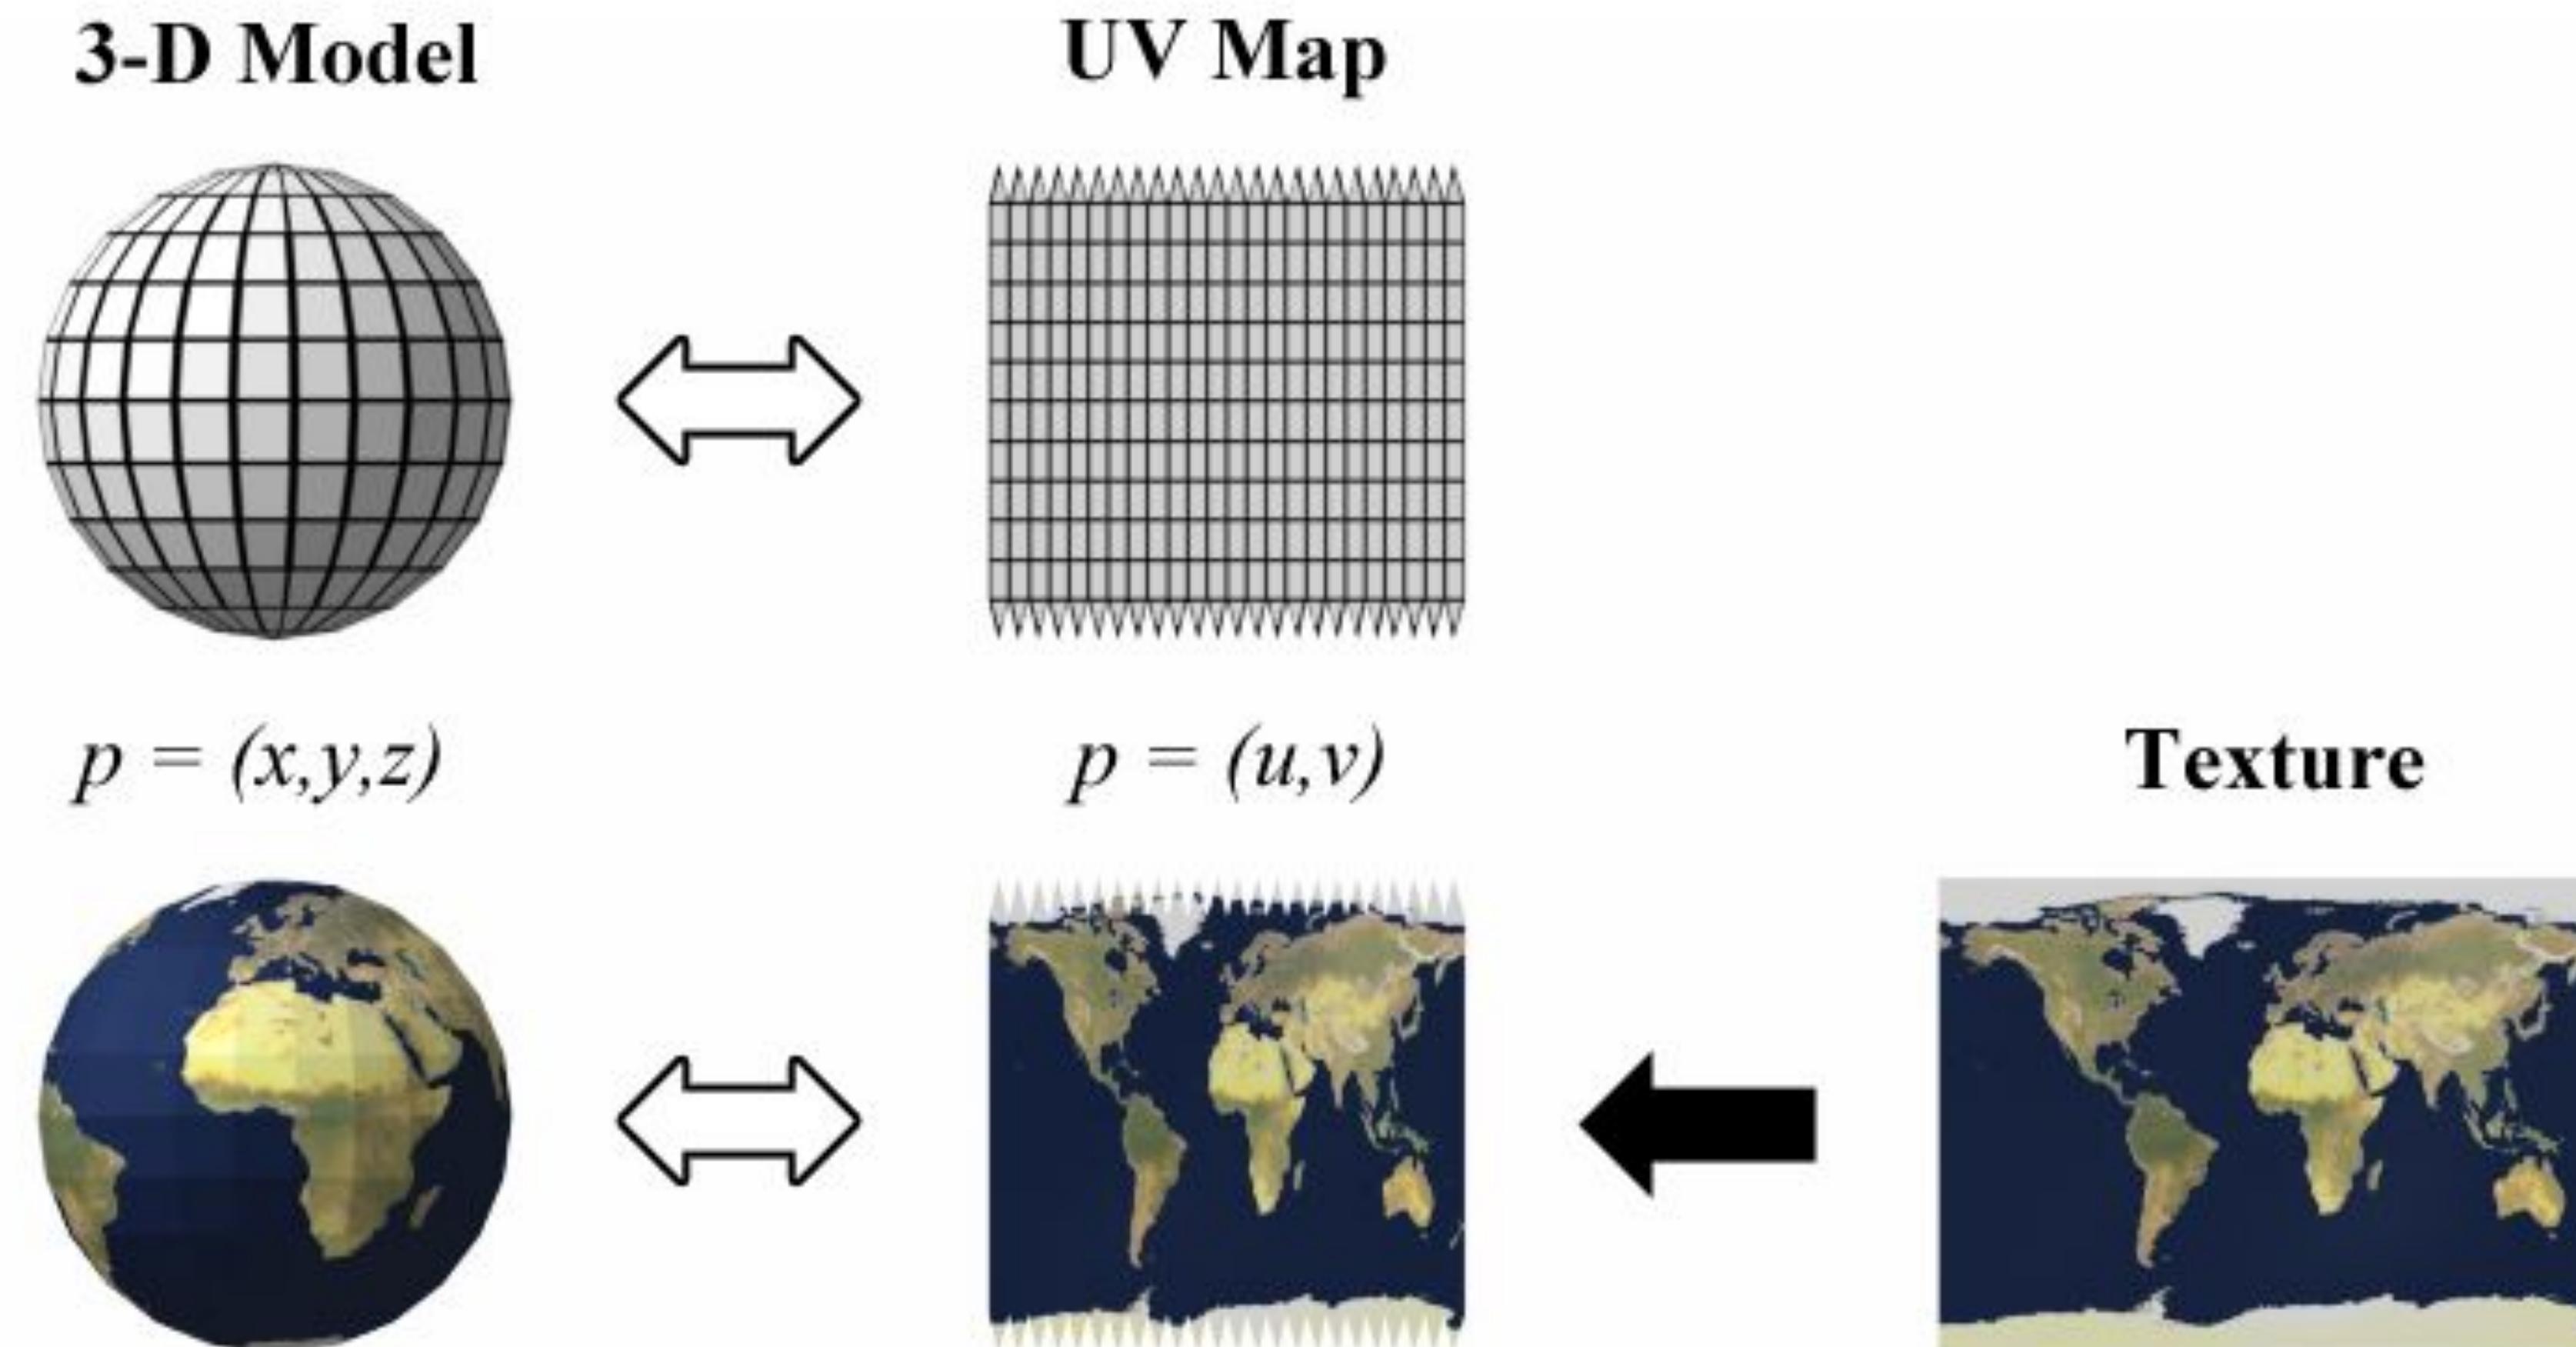


Spatial domain

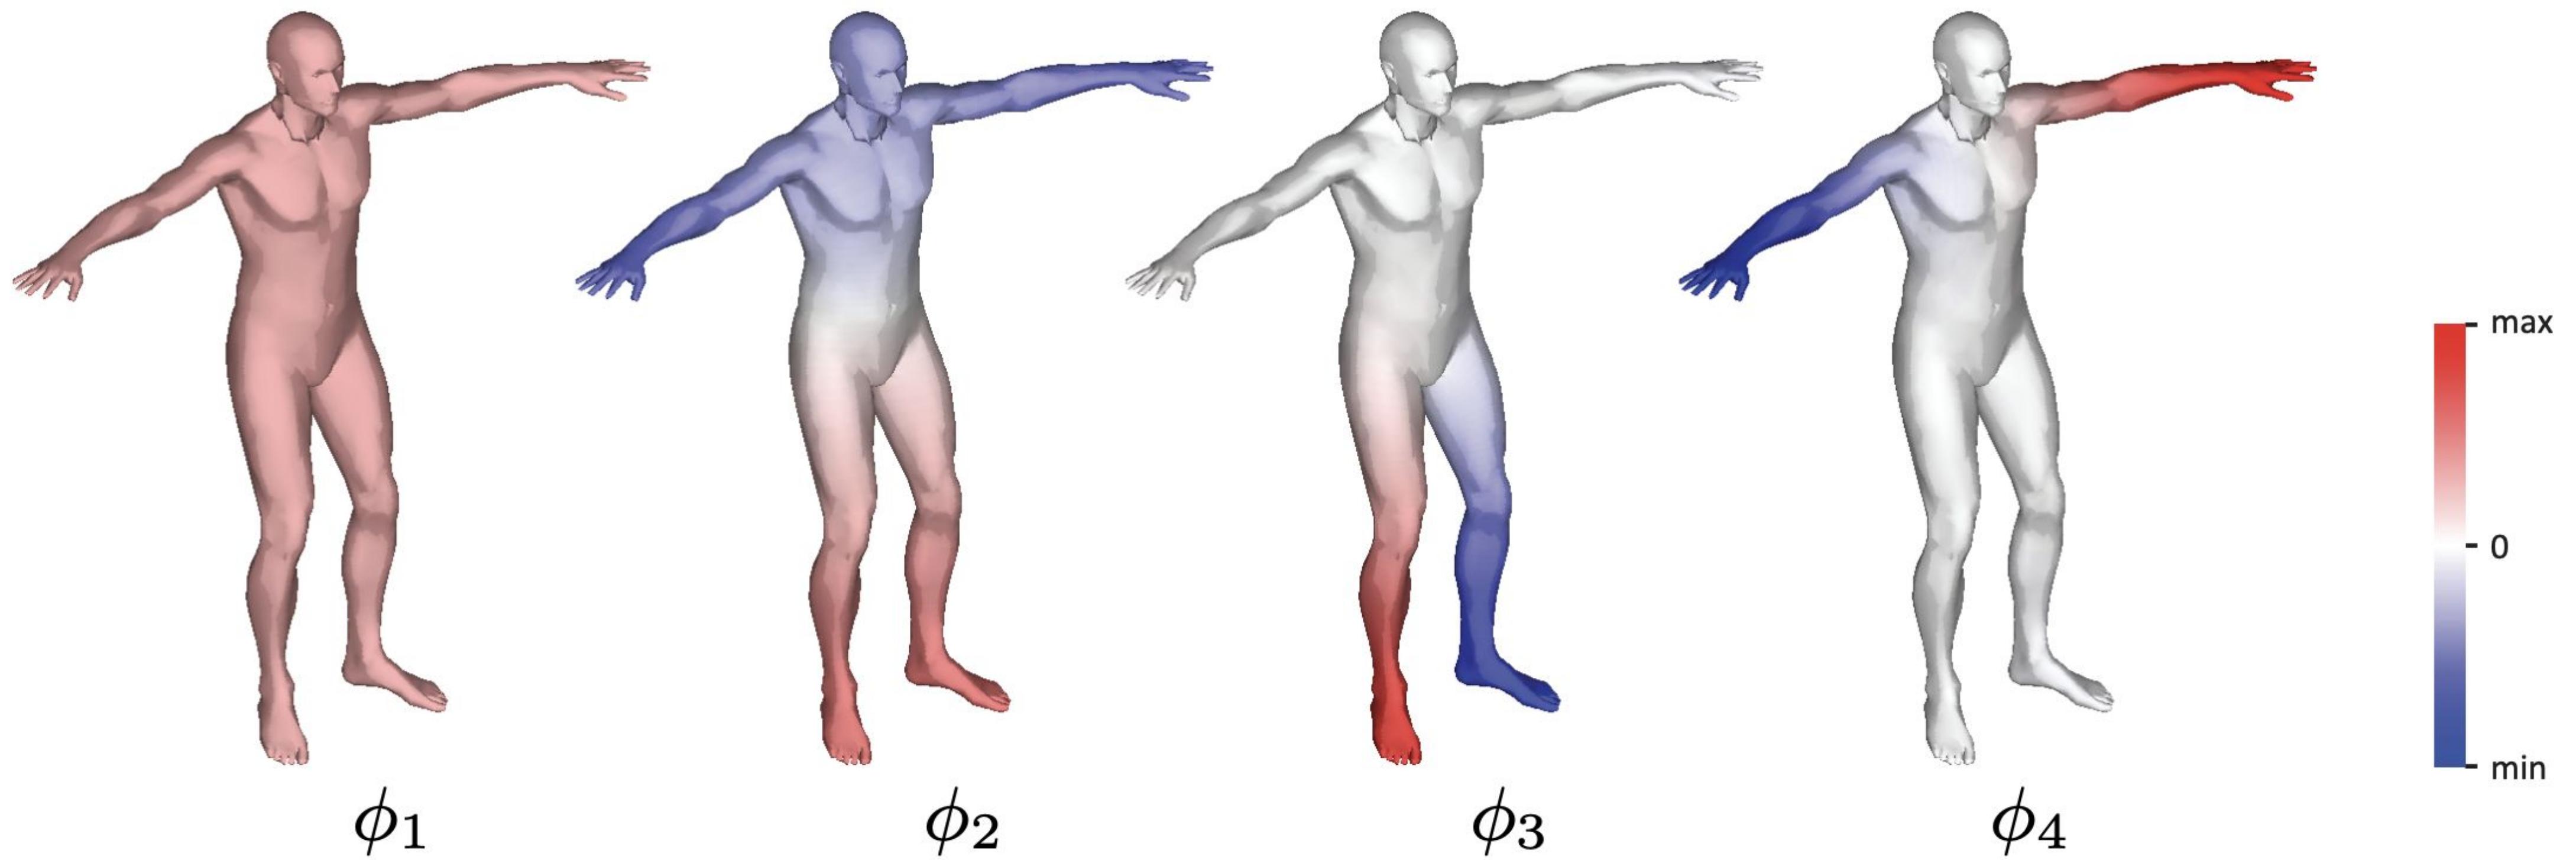


Parametric domain

Non-Euclidean CNNs via conventional CNNs



Laplacian eigenfunctions and eigenvalues



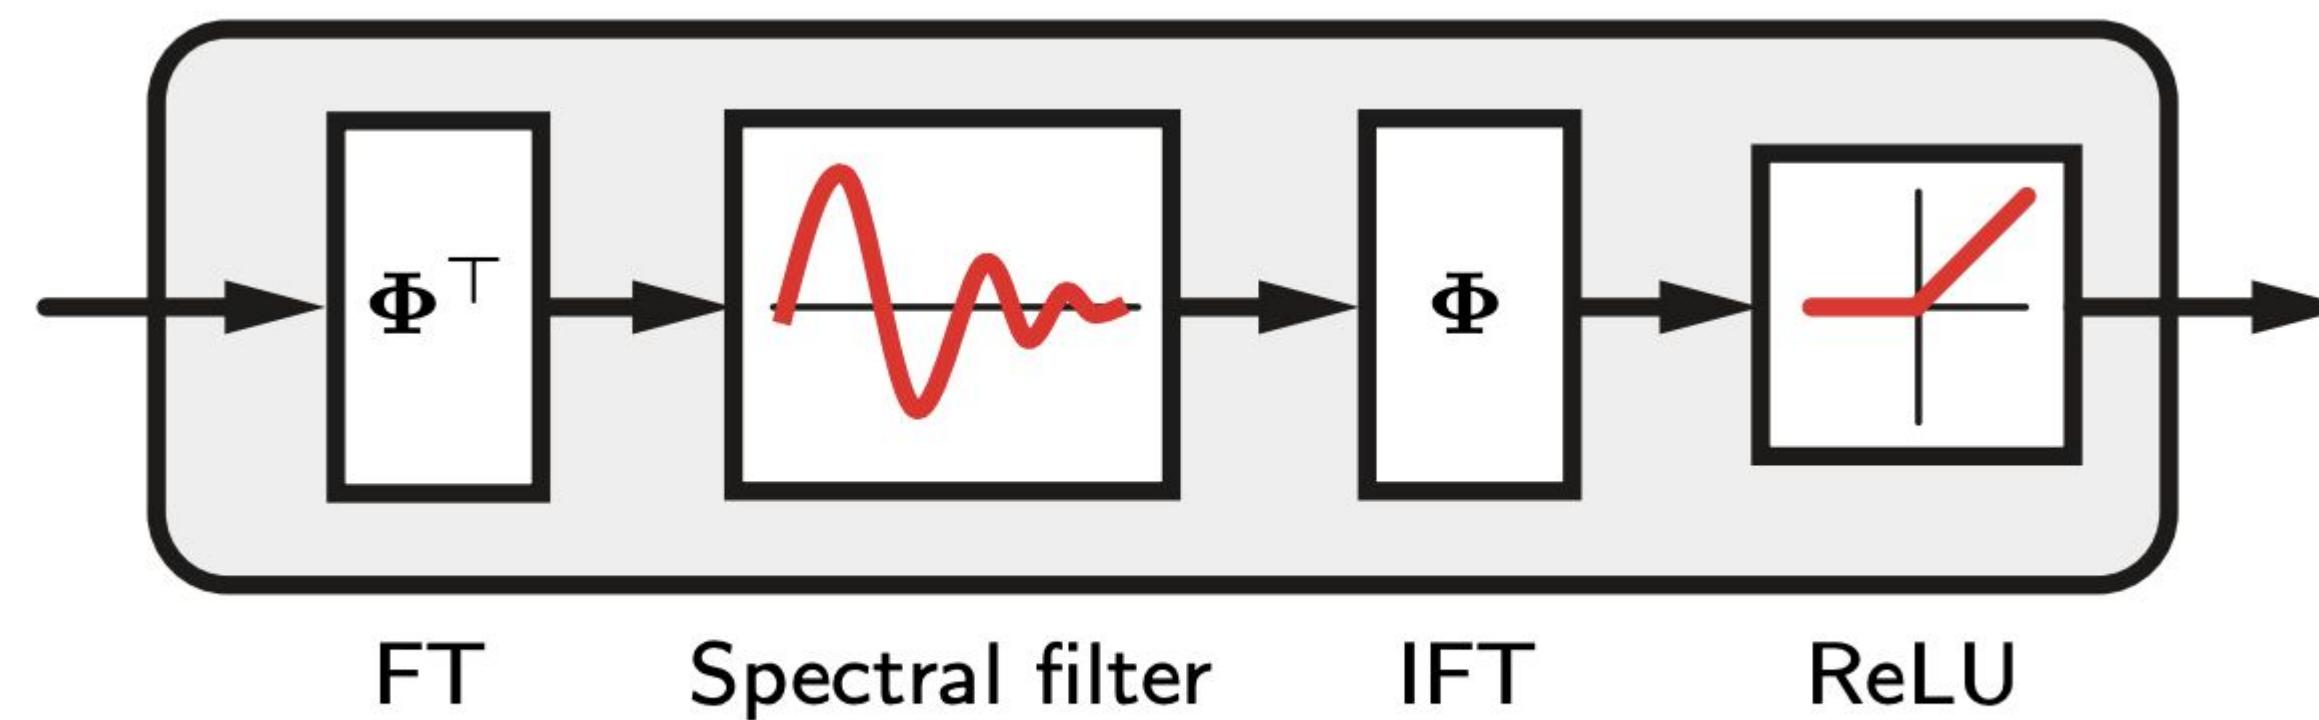
First eigenfunctions of a manifold Laplacian

Spectral CNN (2014)

Convolution expressed in the spectral domain:

$$\mathbf{g} = \Phi \mathbf{W} \Phi^\top \mathbf{f}$$

where \mathbf{W} is $n \times n$ diagonal matrix of learnable spectral filter coefficients



Convolutional filter of a Spectral CNN

Spectral CNN with polynomial filters (2016)

Represent spectral transfer function as a polynomial of order r

$$\hat{g}_{\boldsymbol{\theta}}(\Delta) = \sum_{\ell=0}^r \theta_\ell \Delta^\ell$$

where $\boldsymbol{\theta} = (\theta_0, \dots, \theta_r)^\top$ is the vector of filter parameters

- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ No explicit computation of $\Phi^\top, \Phi \Rightarrow \mathcal{O}(nr)$ complexity
- ☺ Filters have guaranteed r -ring support
- ☹ Mesh dependent

Spectral CNN with rational filters (2017)

Represent spectral transfer function as a Cayley polynomial of order r

$$\hat{g}_{\mathbf{c},h}(\lambda) = c_0 + 2\text{Re}\left\{\sum_{\ell=1}^r c_\ell (h\lambda - i)^\ell (h\lambda + i)^{-\ell}\right\}$$

where the filter parameters are the vector of real/complex coefficients $\mathbf{c} = (c_0, \dots, c_r)^\top$ and the spectral zoom h

Spectral CNN with rational filters (2017)

- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ Filters have guaranteed exponential spatial decay
- ☺ Spectral zoom property allowing to better localize in frequency
- ☺ Richer class of filters than polynomials for the same order
- ☺ Scale invariance
- ☺ $\mathcal{O}(n)$ computational complexity with **Jacobi approximate matrix inversion**

Spatial convolutions on meshes

Local system of coordinates u_{ij} around i (e.g. geodesic polar)

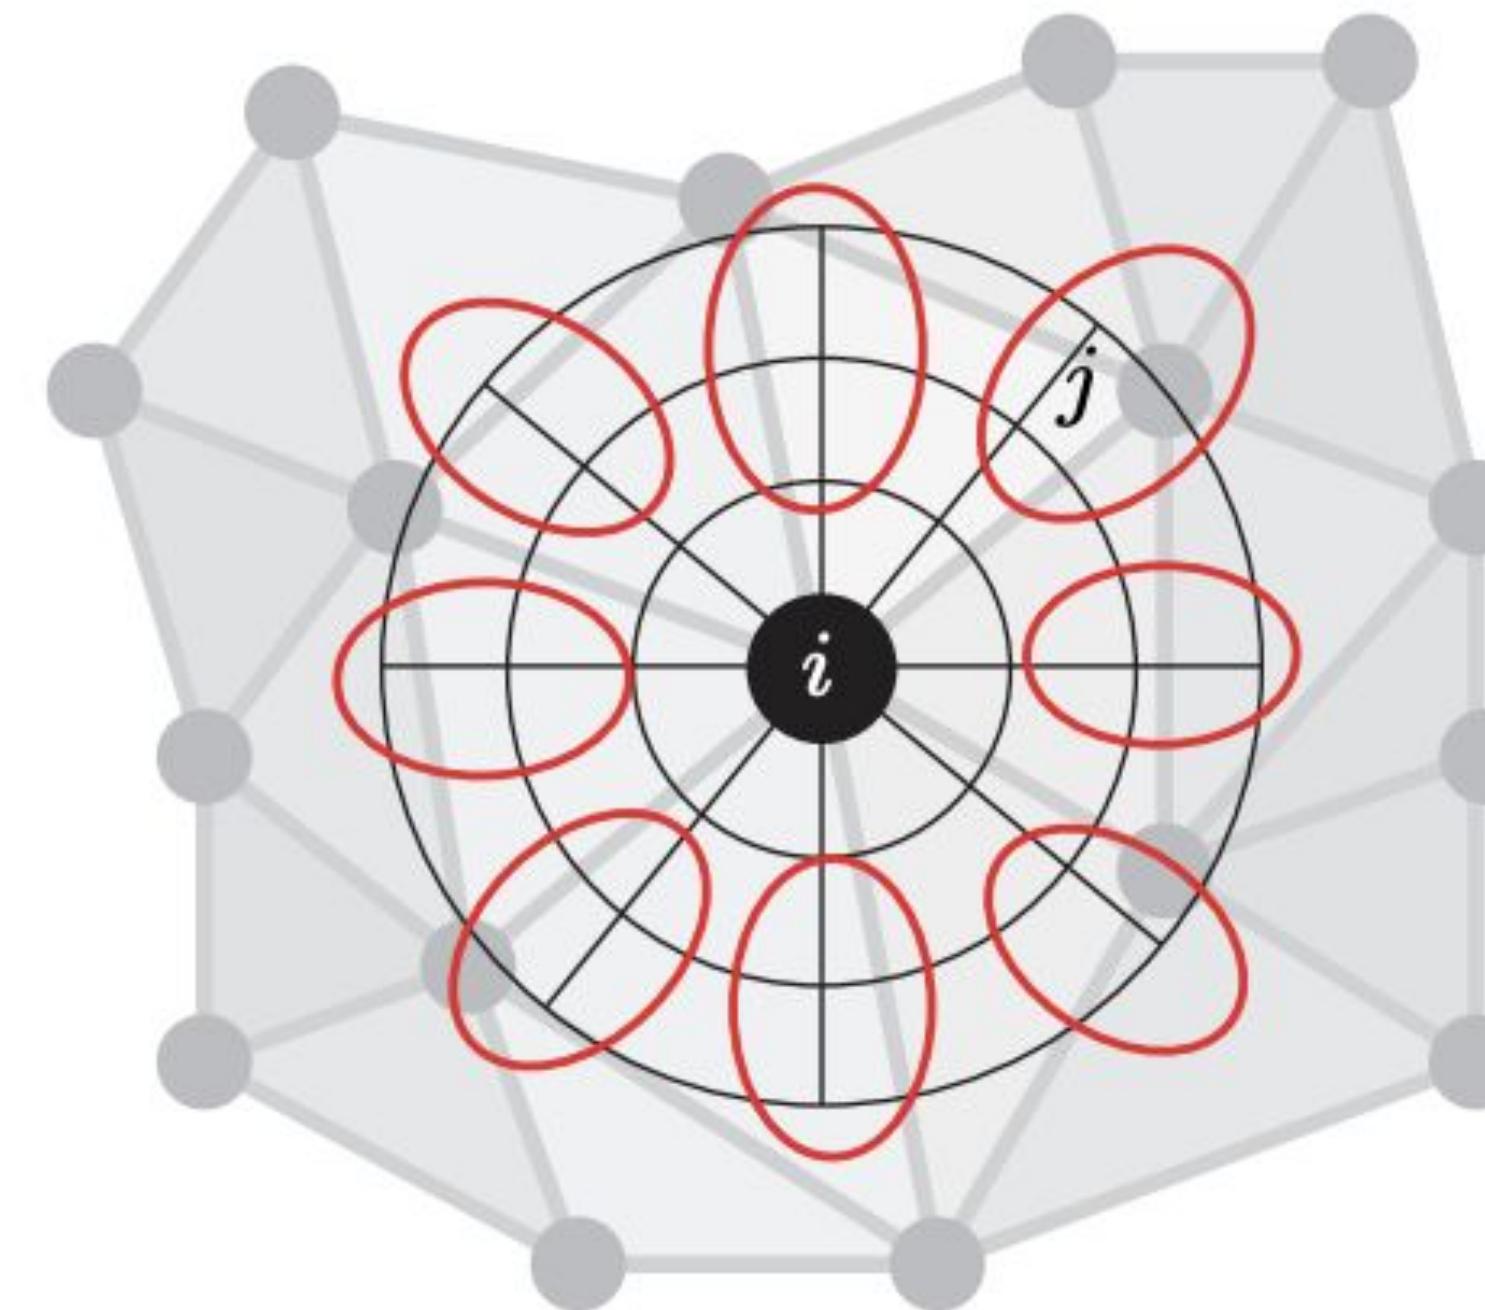
Local weights $w_1(\mathbf{u}), \dots, w_L(\mathbf{u})$ w.r.t. \mathbf{u} , e.g. Gaussians

$$w_\ell = \exp(-(\mathbf{u} - \boldsymbol{\mu}_\ell)^\top \boldsymbol{\Sigma}_\ell^{-1} (\mathbf{u} - \boldsymbol{\mu}_\ell))$$

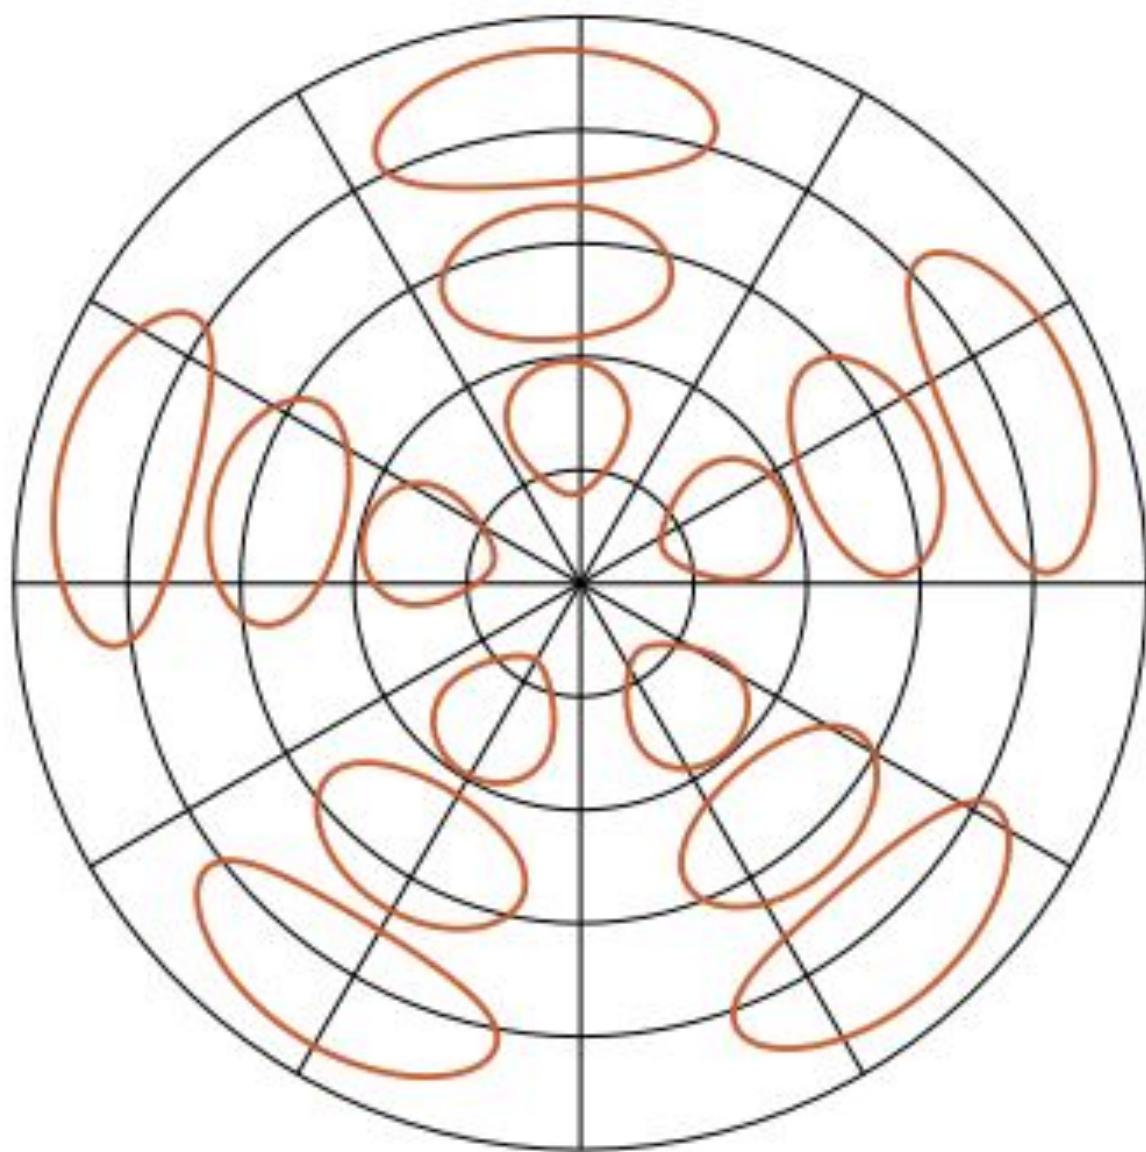
Spatial convolution with filter g

$$\mathbf{x}'_i \propto \sum_{\ell=1}^L g_\ell \underbrace{\sum_{j=1}^n w_\ell(\mathbf{u}_{ij}) \mathbf{x}_j}_{\text{patch operator}}$$

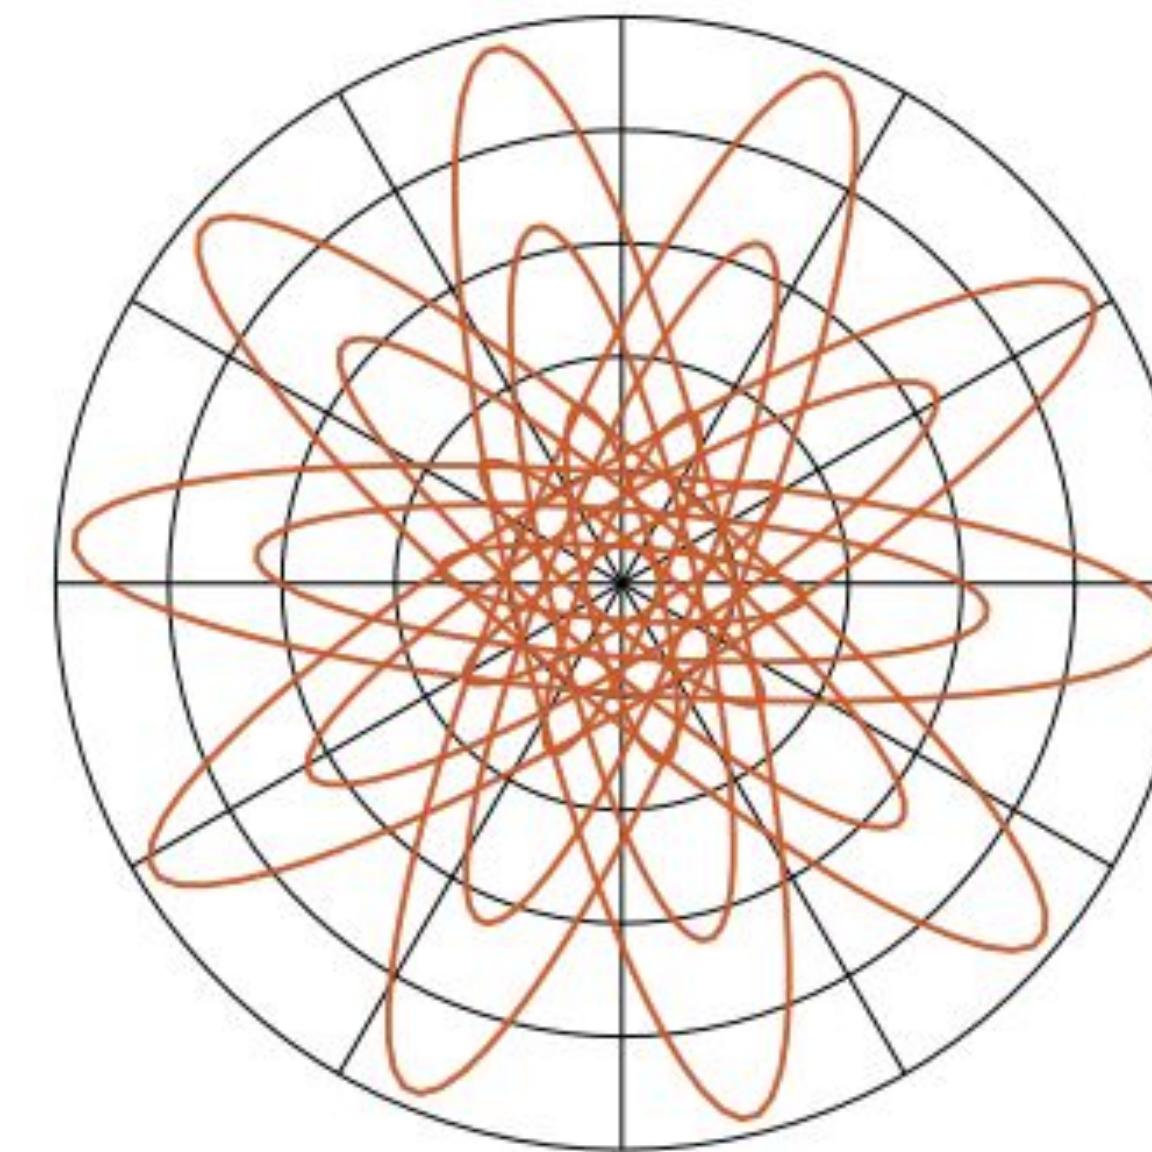
where $\mathbf{x}_i \in \mathbb{R}^d$ is feature at vertex i



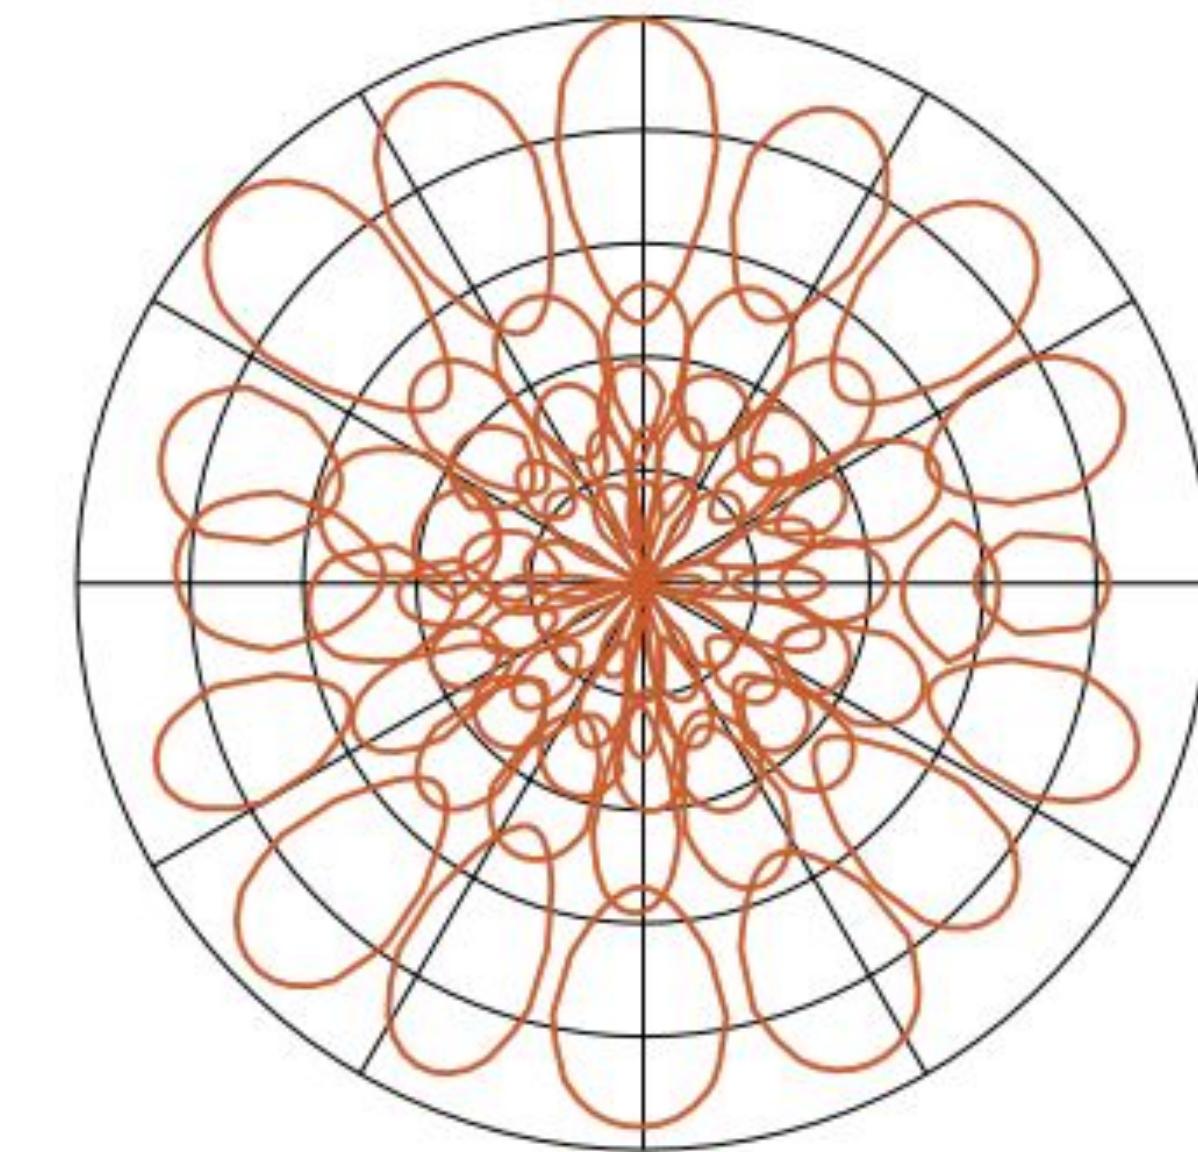
Patch operator weight functions



GCNN



ACNN



MoNet

MoNet (2017)

Geodesic polar coordinates $\mathbf{u}_{ij} = (\rho_{ij}, \theta_{ij})$

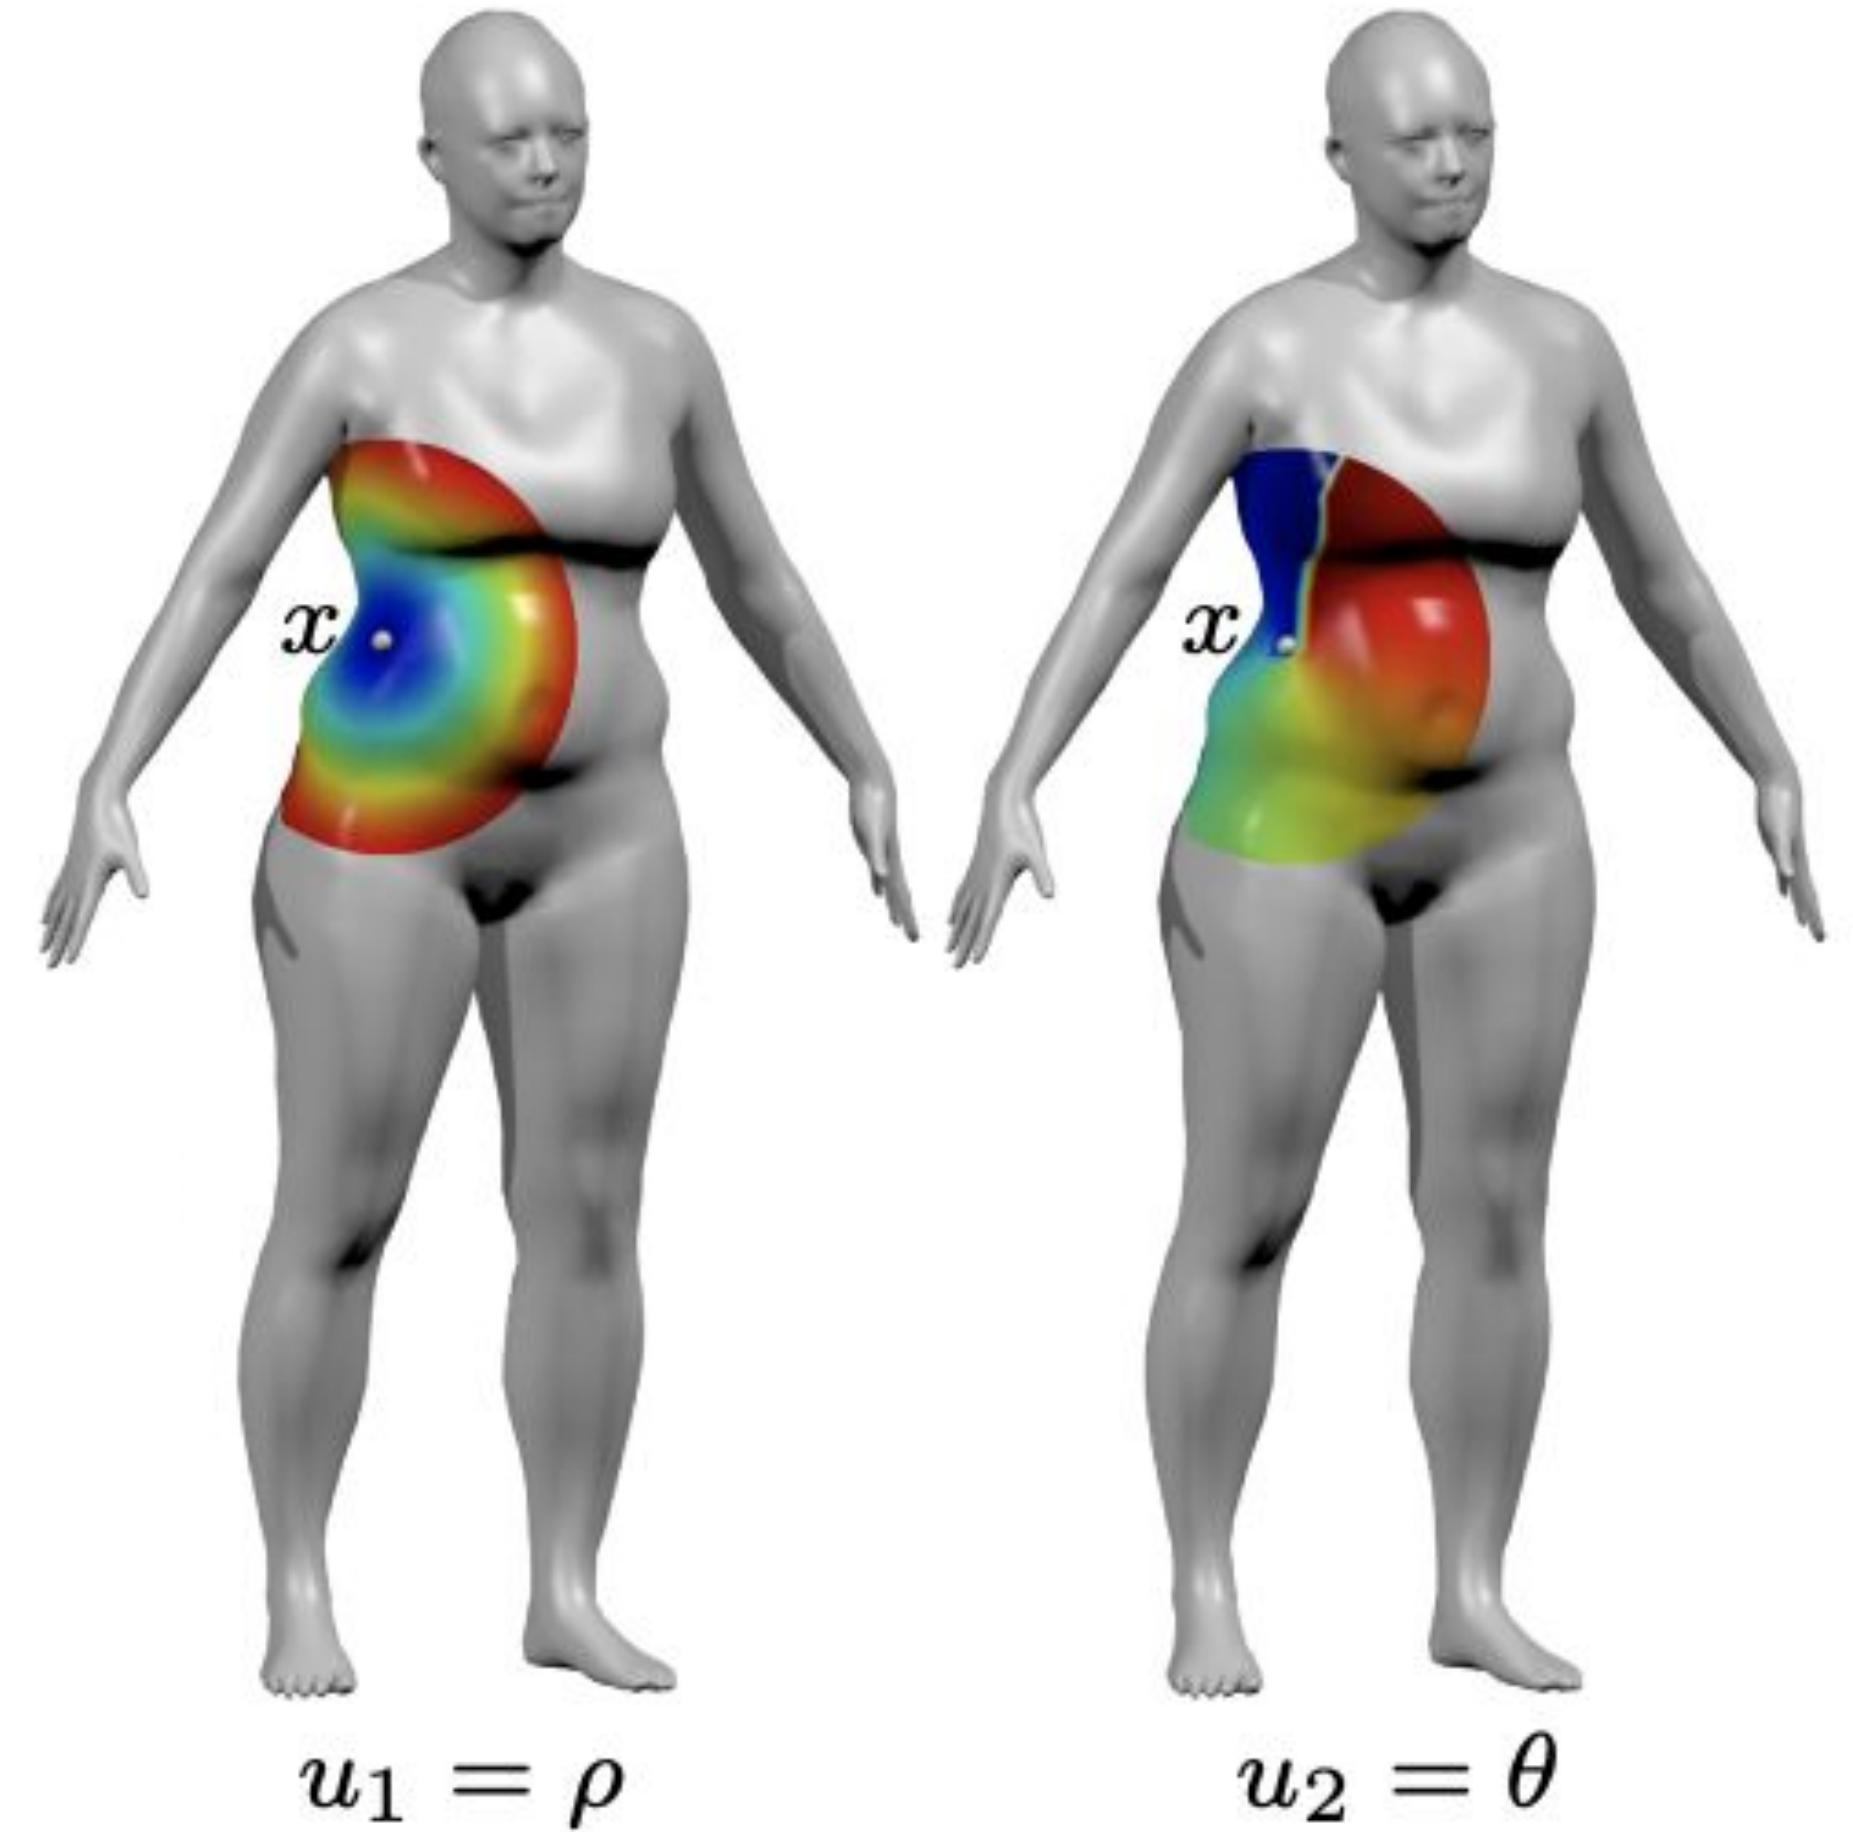
Gaussian weighting functions

$$w_{\mu, \Sigma}(\mathbf{u}) = \exp\left(-\frac{1}{2}(\mathbf{u} - \mu)^\top \Sigma^{-1}(\mathbf{u} - \mu)\right)$$

with learnable covariance Σ and mean μ

Spatial convolution

$$\mathbf{x}'_i \propto \underbrace{\sum_{j=1}^n \sum_{\ell=1}^L g_\ell w_{\mu_\ell, \Sigma_\ell}(\mathbf{u}_{ij}) \mathbf{x}_j}_{\text{Gaussian mixture}}$$



$$u_1 = \rho$$

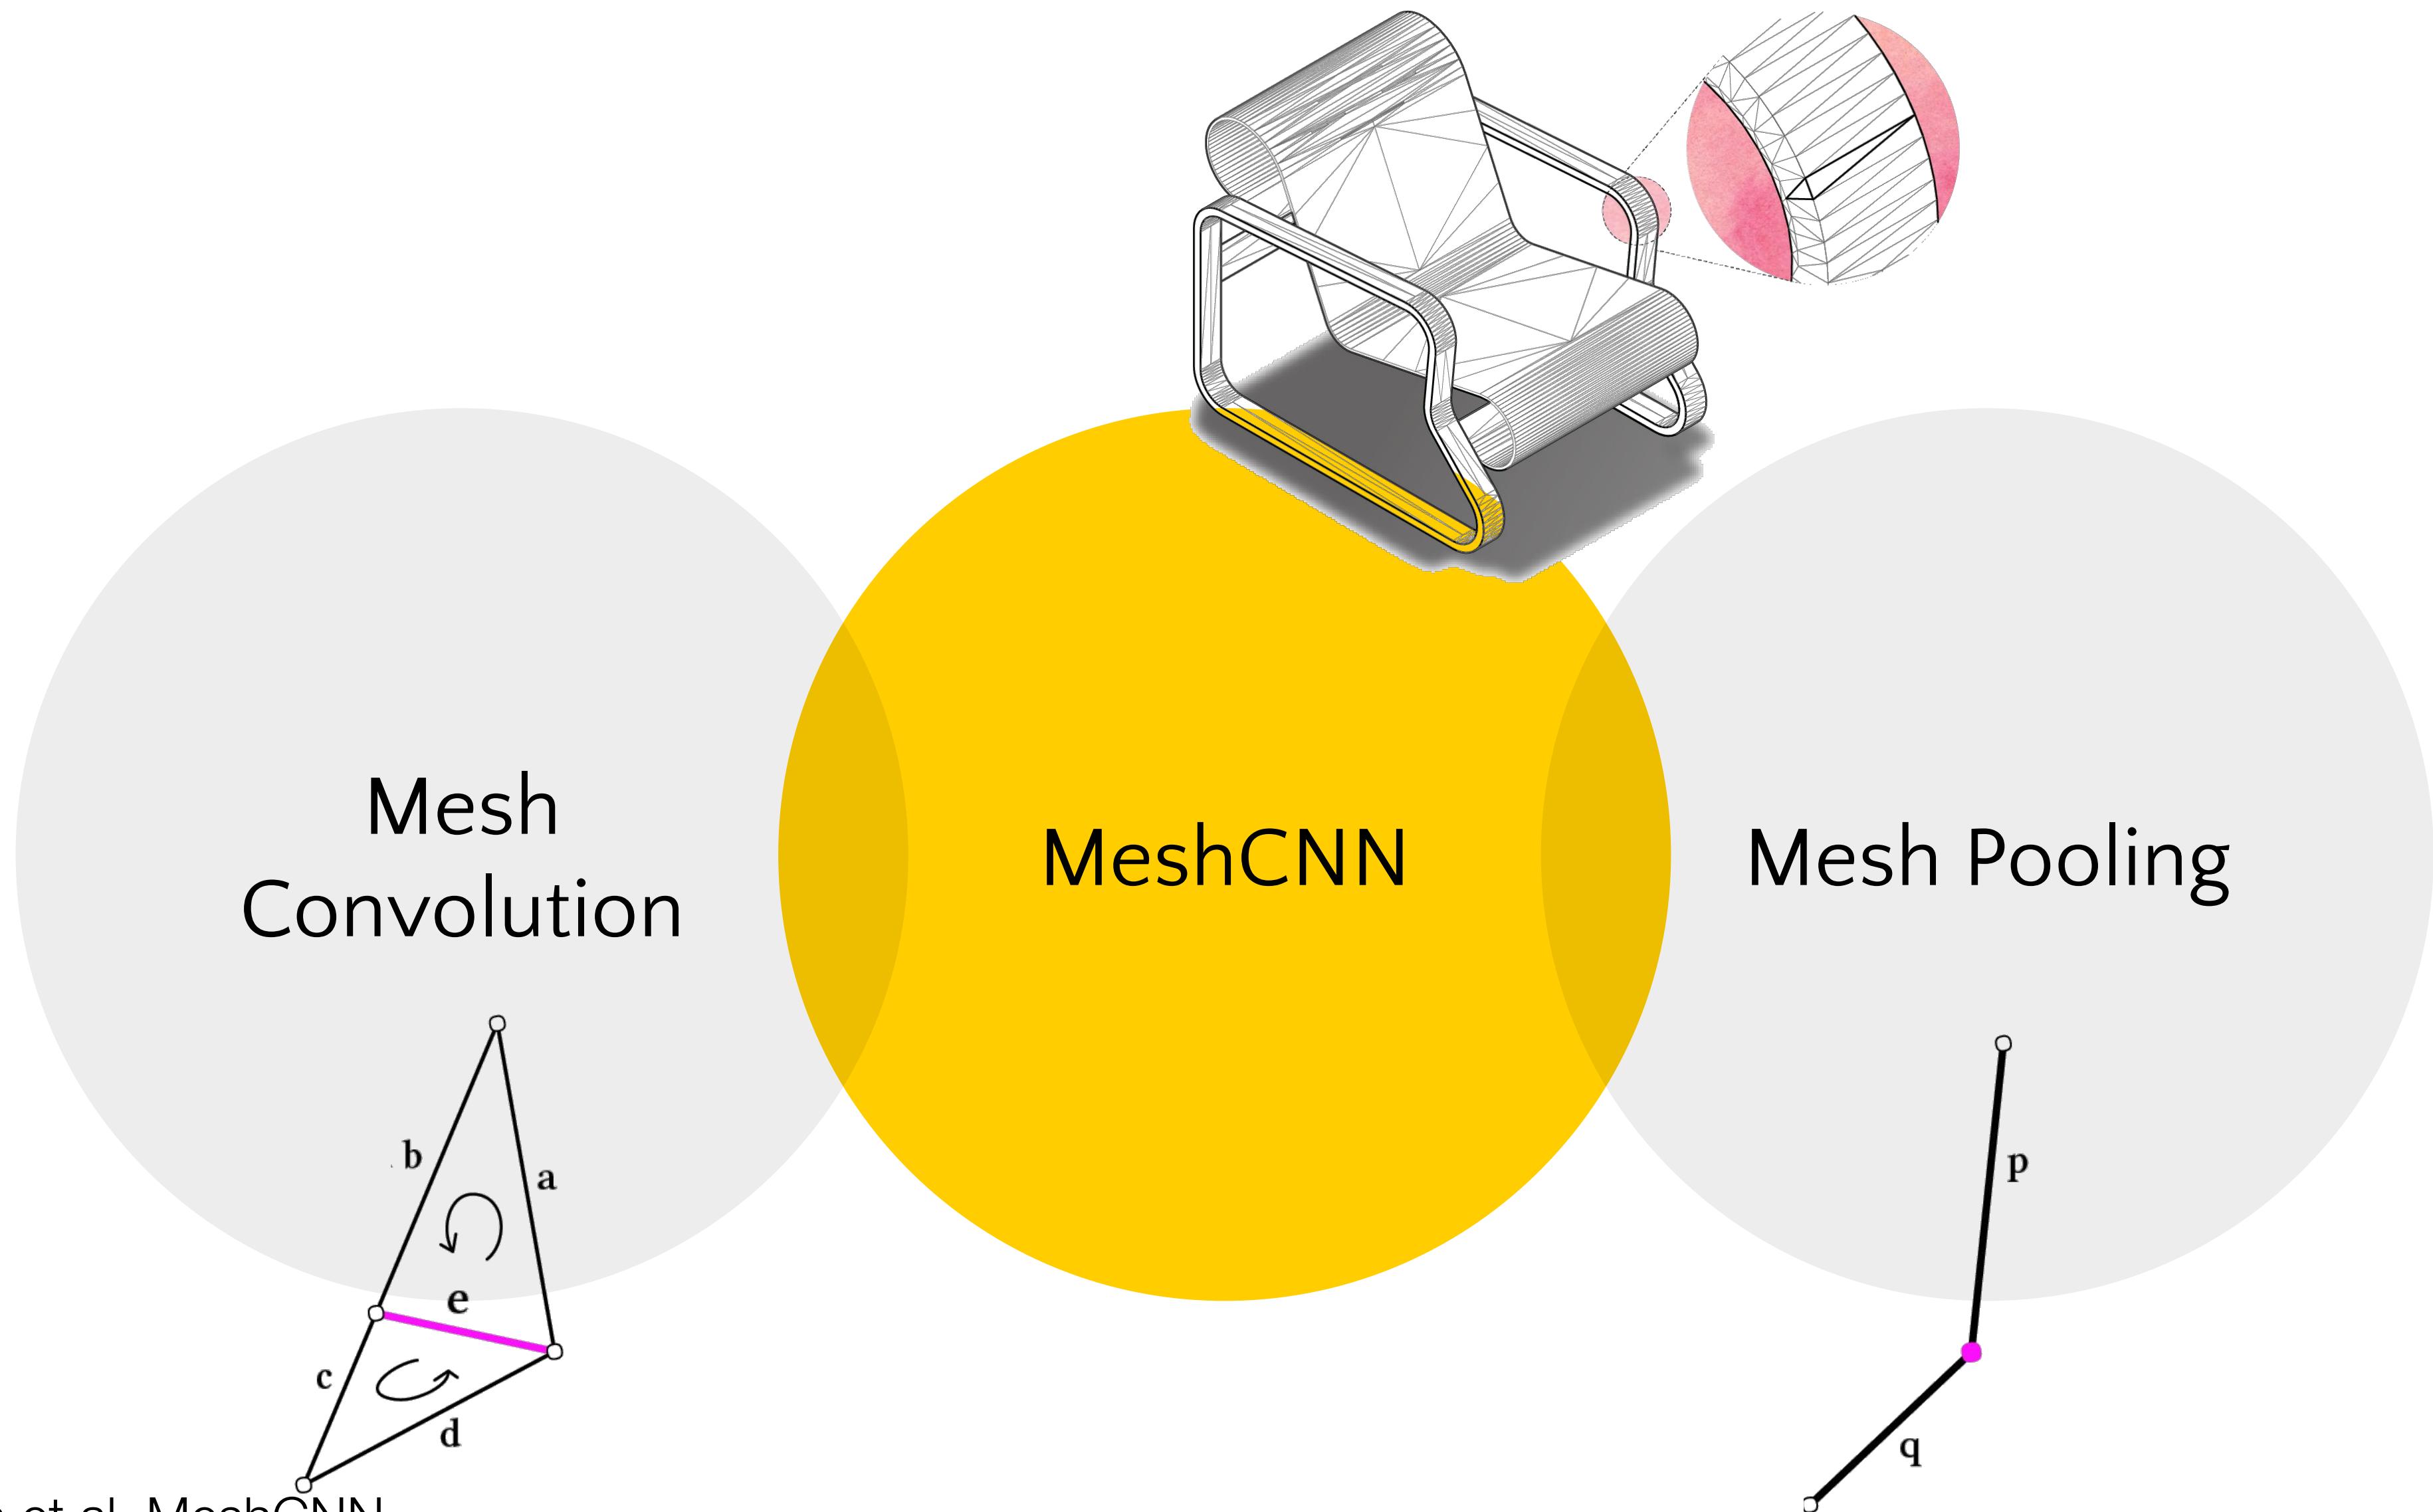
$$u_2 = \theta$$

MoNet as a generalization

Method	Coordinates \mathbf{u}_{ij}	Weight function $w_{\Theta}(\mathbf{u})$
CNN ¹	$\mathbf{u}_j - \mathbf{u}_i$	$\delta(\mathbf{u} - \mathbf{v})$
GCNN ²	ρ_{ij}, θ_{ij}	$\exp\left(-\frac{1}{2}(\mathbf{u} - \mathbf{v})^\top \begin{pmatrix} \sigma_\rho^2 & \\ & \sigma_\theta^2 \end{pmatrix}^{-1} (\mathbf{u} - \mathbf{v})\right)$
ACNN ³	ρ_{ij}, θ_{ij}	$\exp\left(-t\mathbf{u}^\top \mathbf{R}_\varphi \begin{pmatrix} \alpha & 1 \end{pmatrix} \mathbf{R}_\varphi^\top \mathbf{u}\right)$
MoNet ⁴	ρ_{ij}, θ_{ij}	$\exp\left(-\frac{1}{2}(\mathbf{u} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{u} - \boldsymbol{\mu})\right)$
		learnable parameters $\Theta = (\boldsymbol{\mu}, \boldsymbol{\Sigma})$

Some CNN models can be considered as particular settings of MoNet
with weighting functions of different form

MeshCNN (2019)



Mesh Convolution

Face normal

- Consistent ordering in each face

Two *valid* orderings

- (a, b, c, d) or (c, d, a, b)

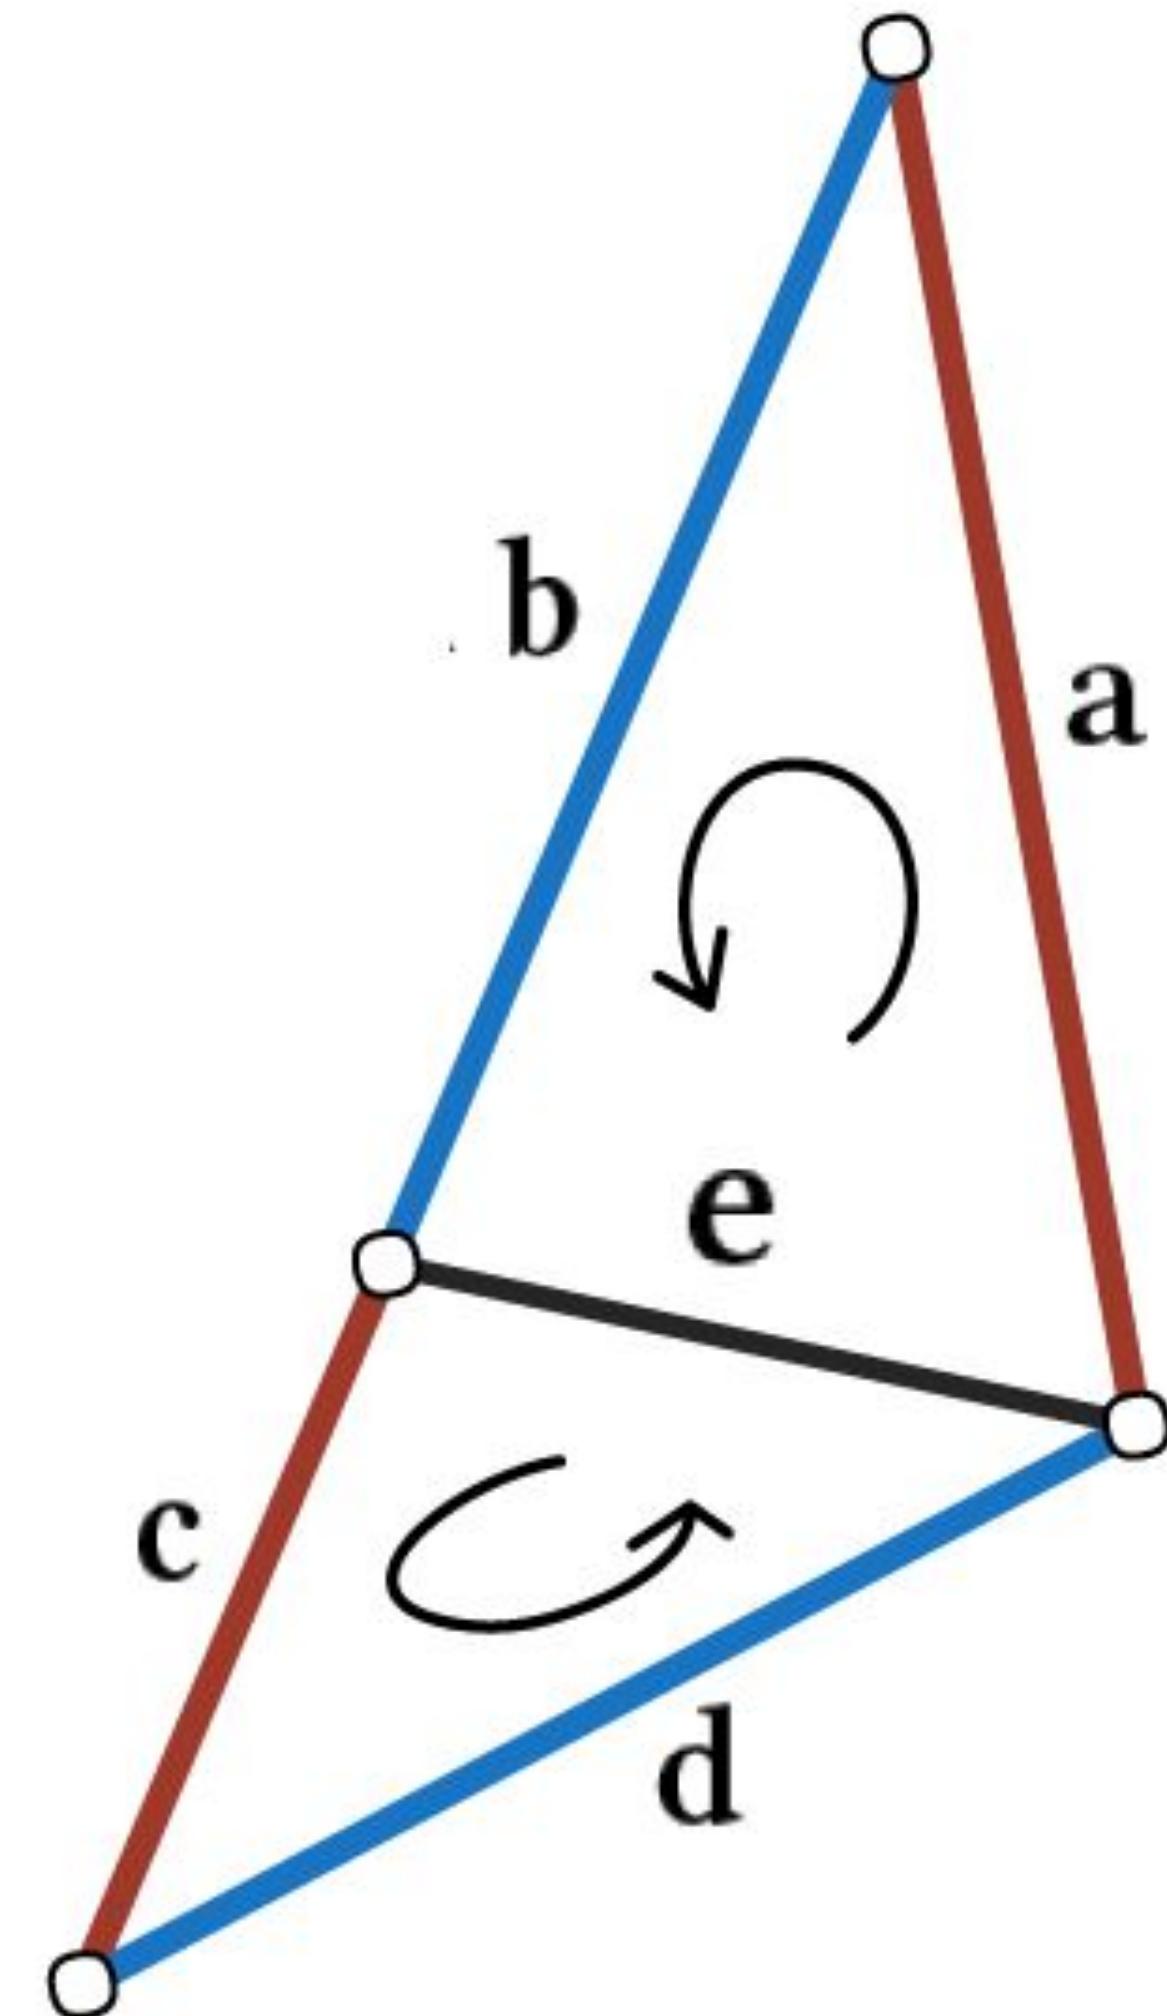
How to handle ambiguity?

Build symmetric features

- $e \rightarrow (a+c, |a-c|, b+d, |b-d|)$

Conv Filter

- $(k_0, k_1, k_2, k_3, k_4)$



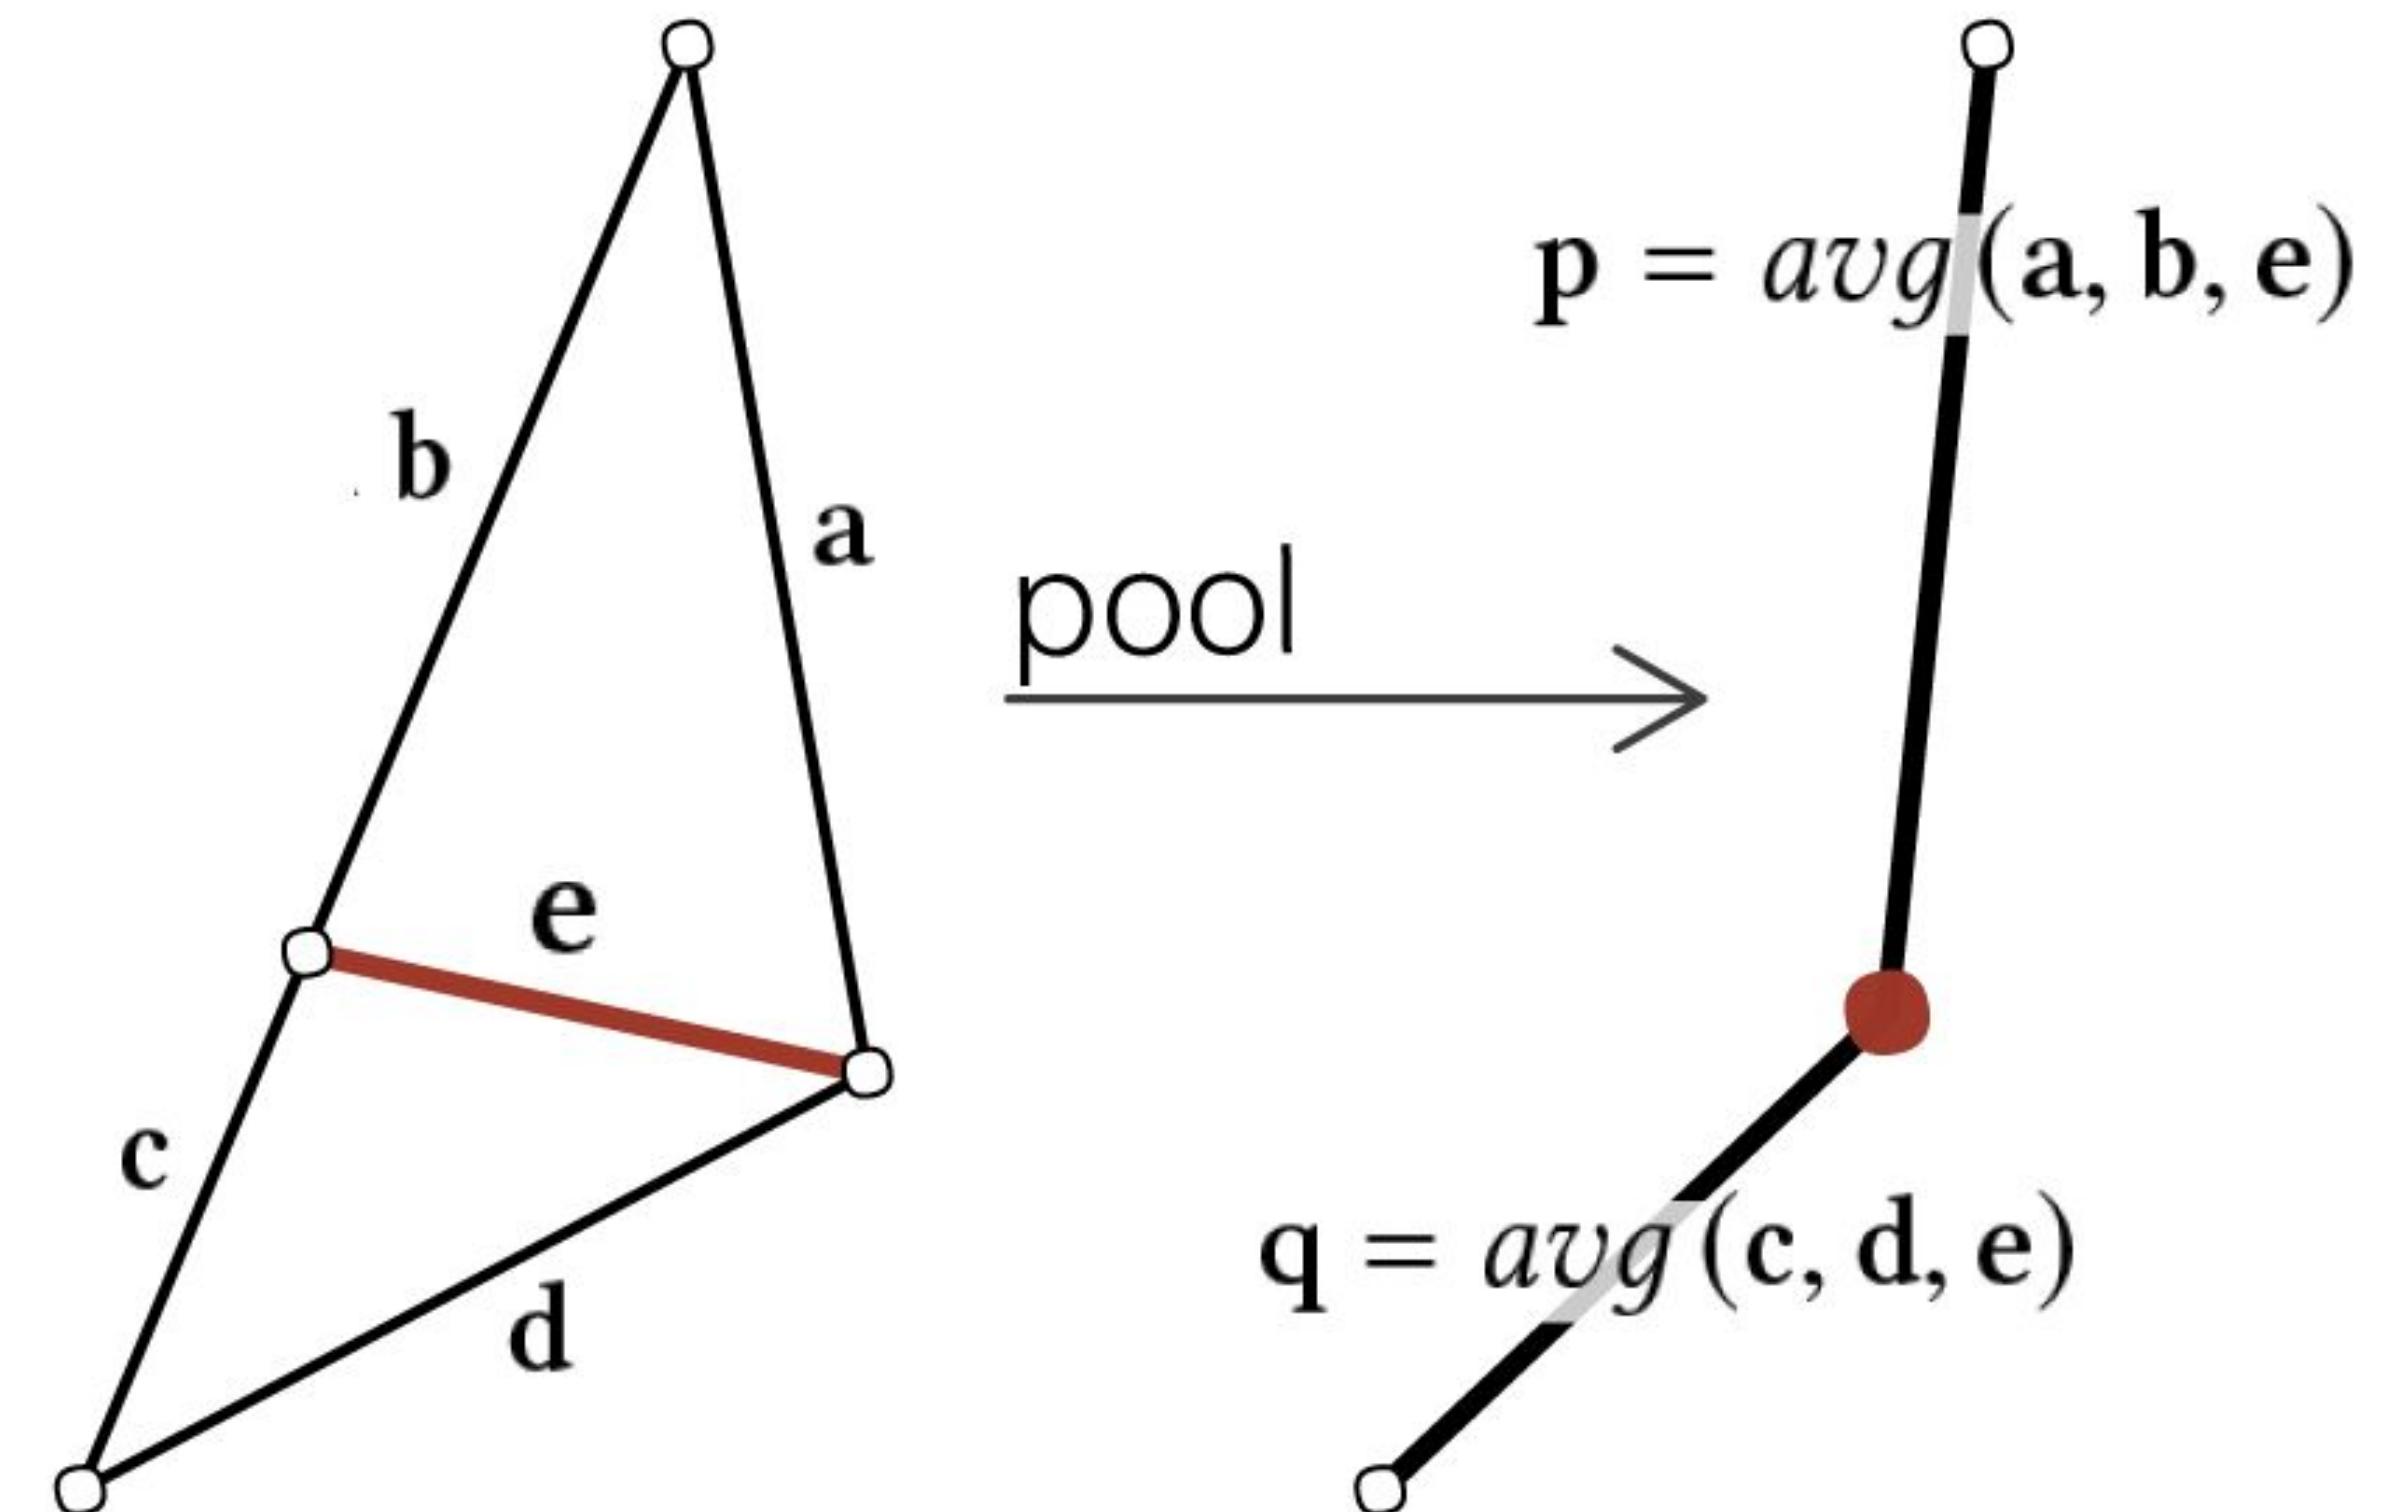
Mesh Pooling

Delete edge with smallest feature activations

- Aggregate features (average pooling)
- Update topology

Average pooling aggregation

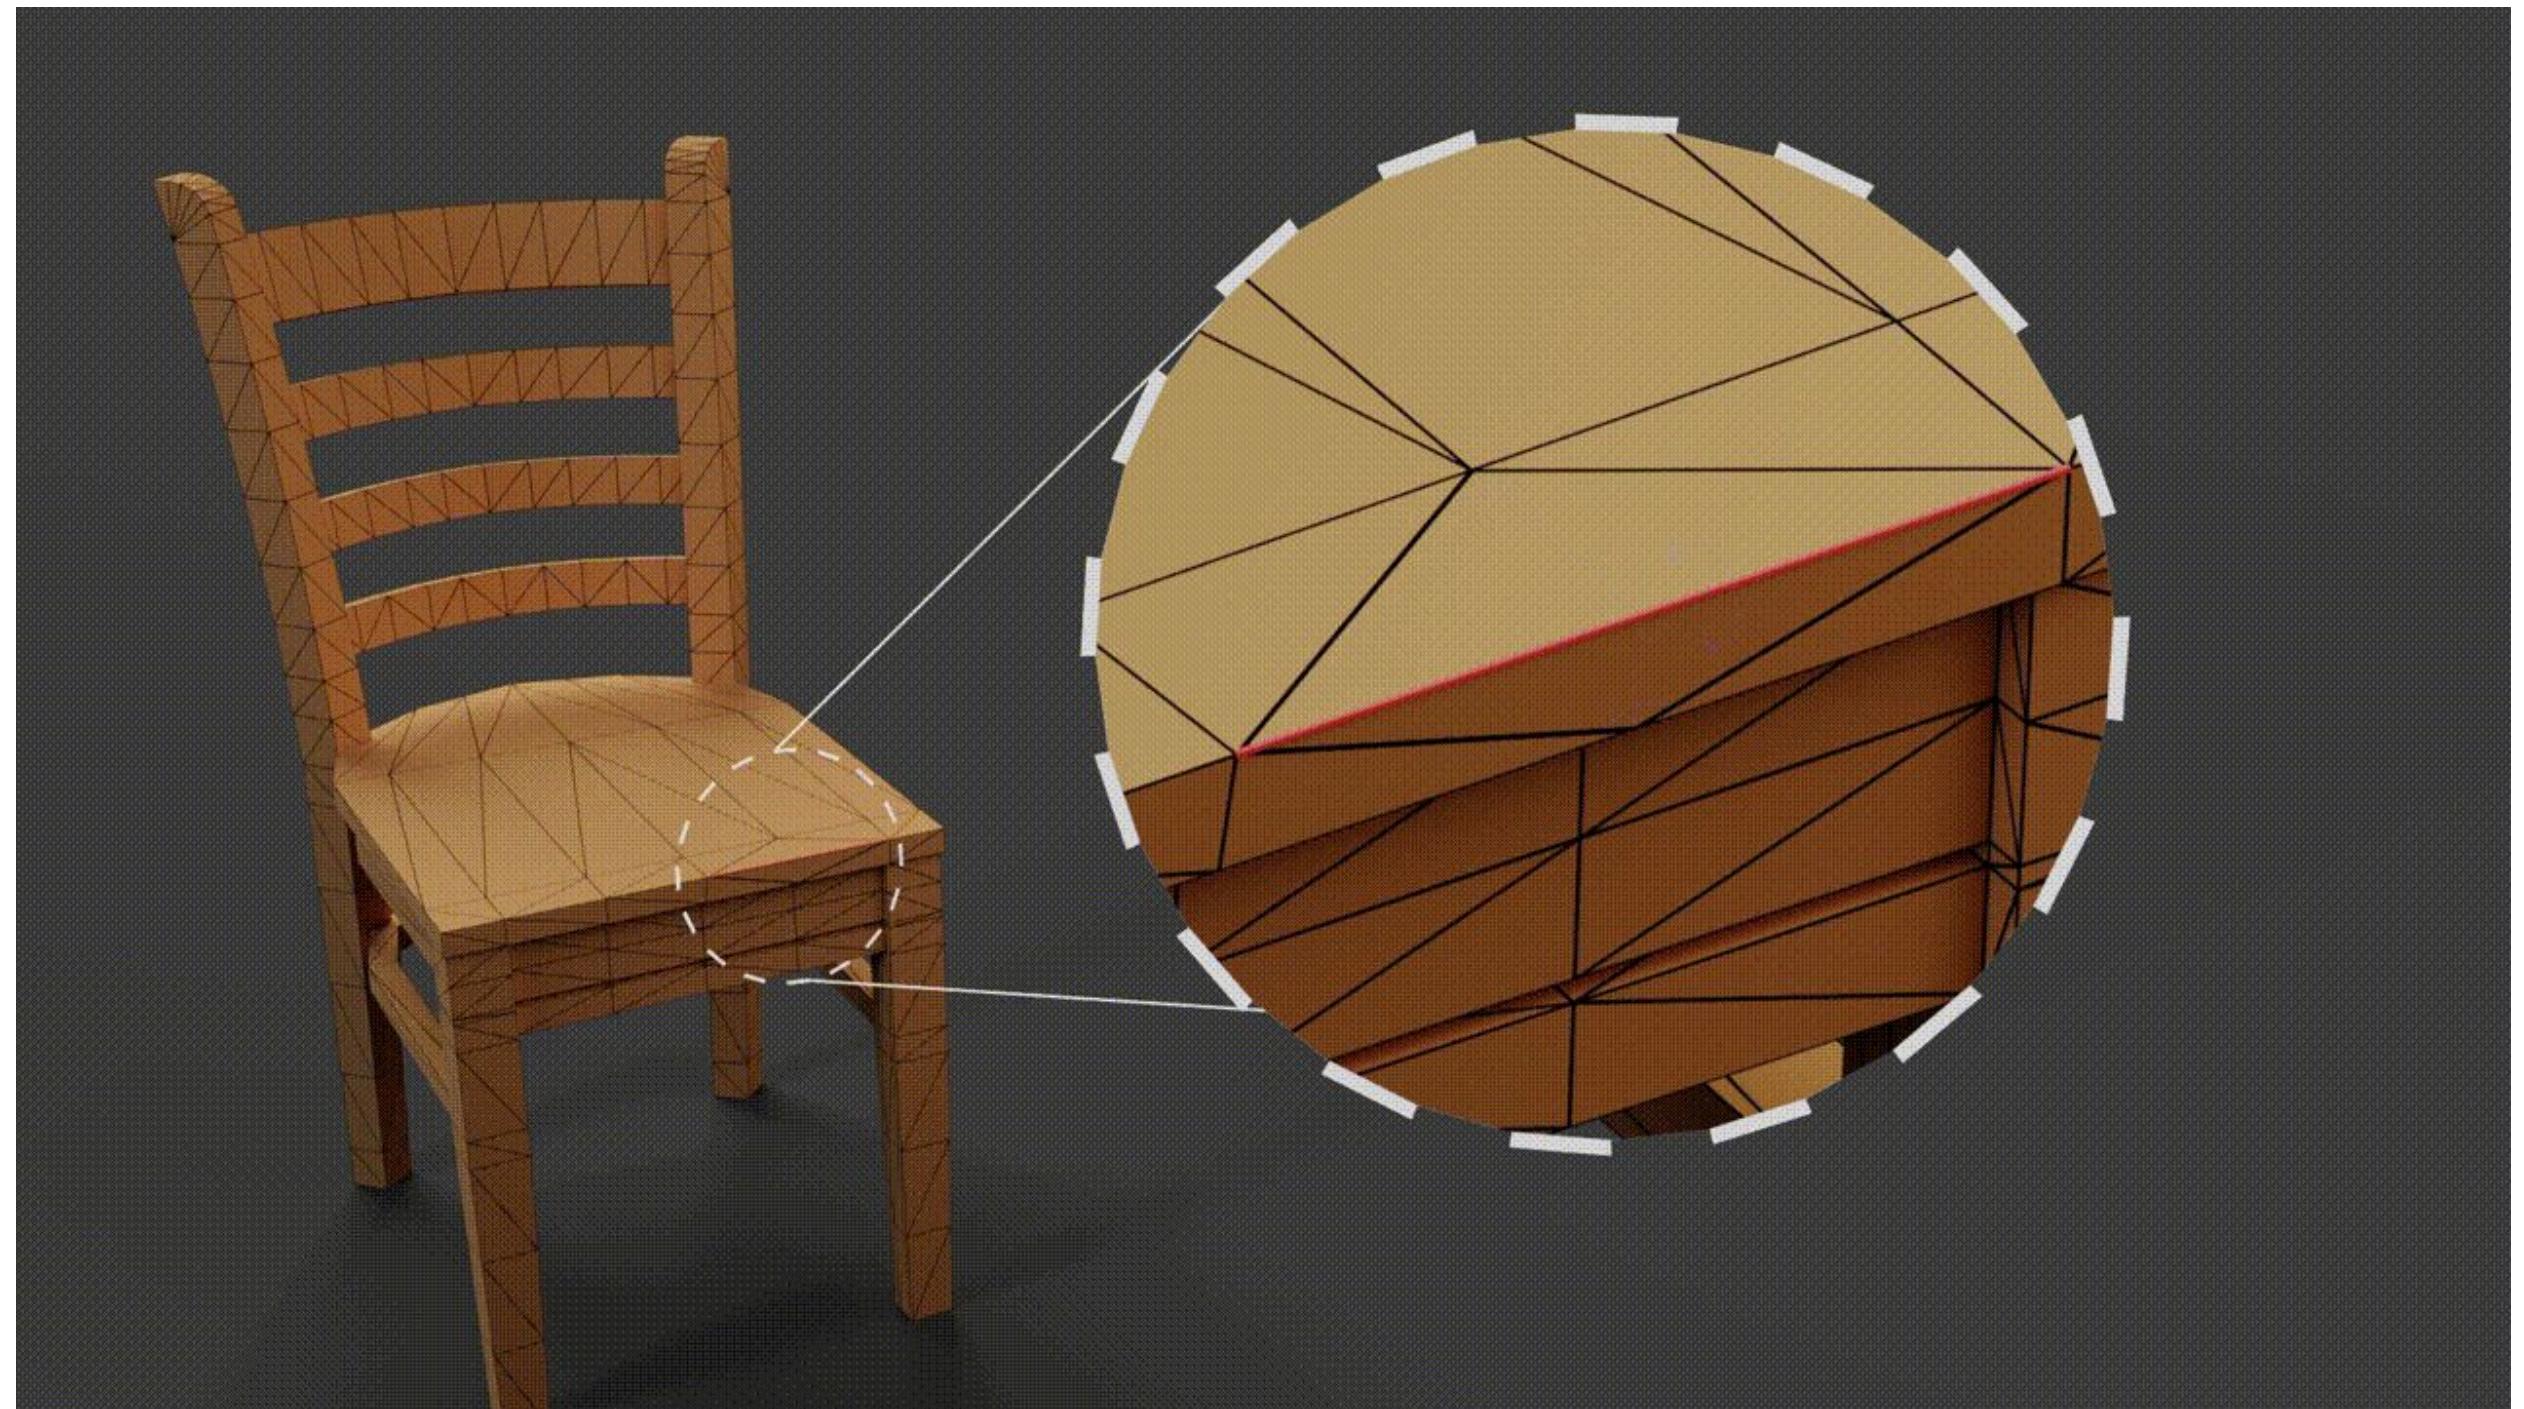
- 5 edges to 2 edges
- Average per-channel



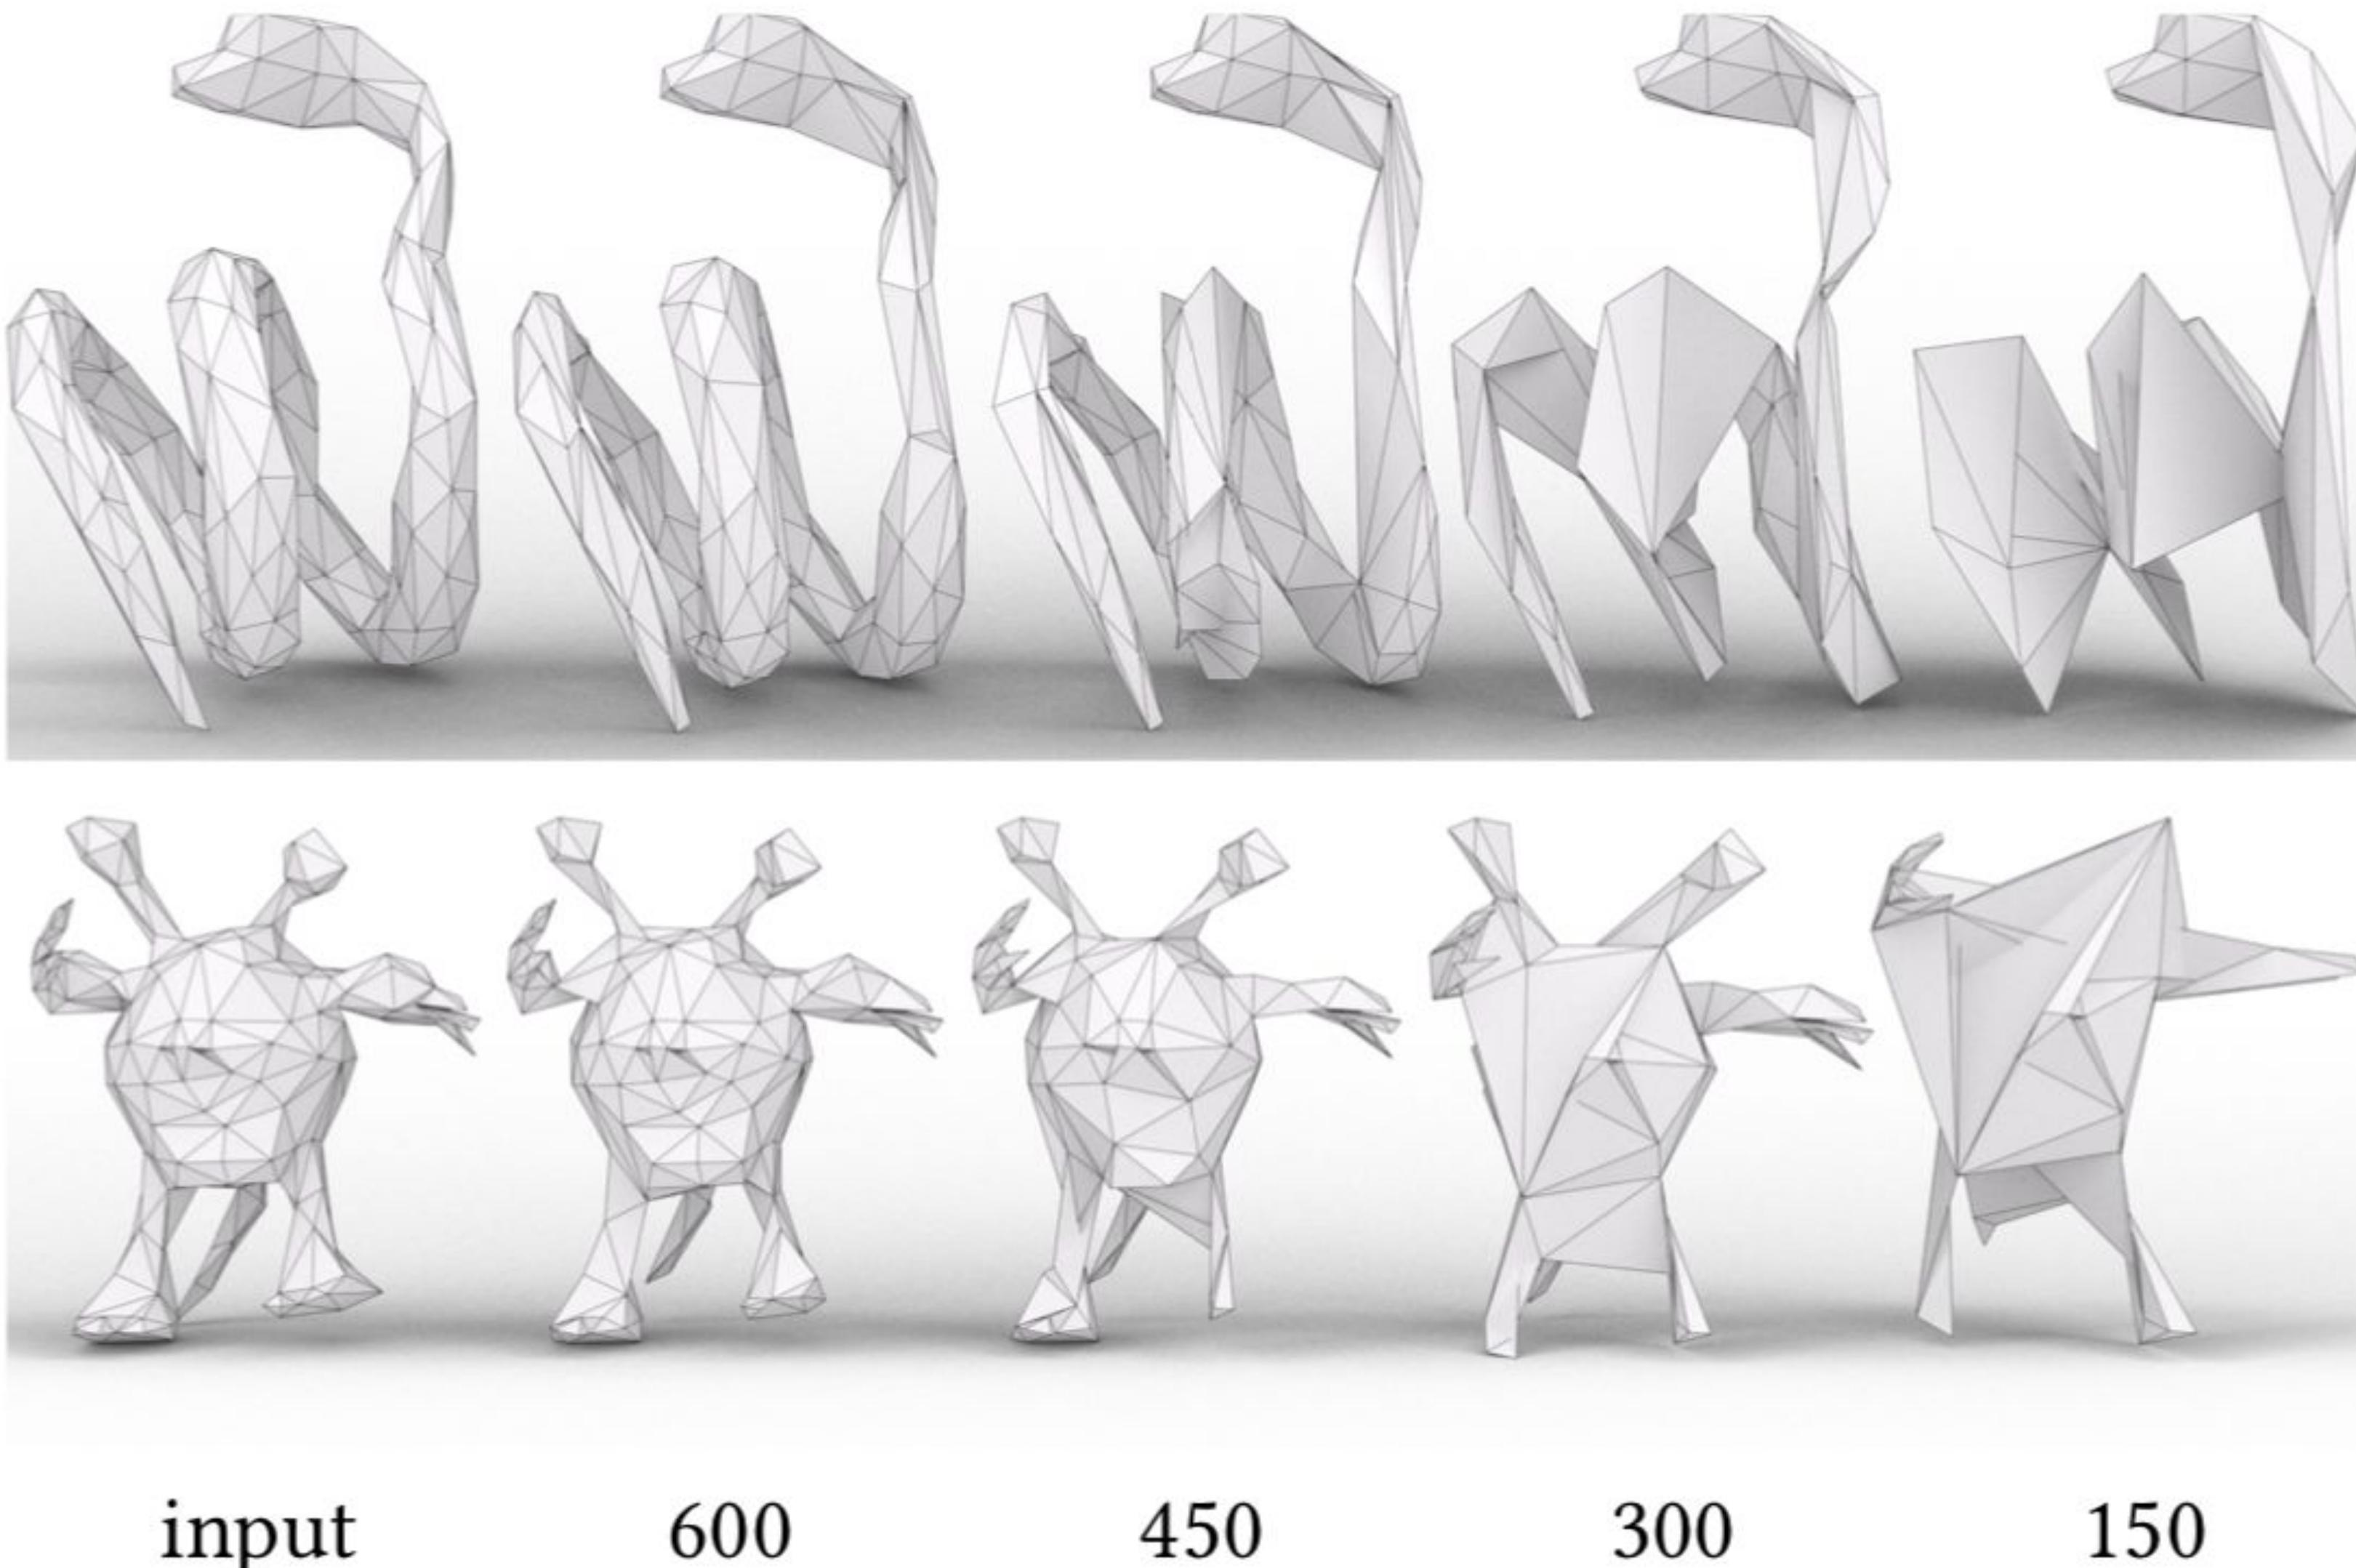
Update Mesh Topology

New mesh is generated

- Defines new neighbors in next convolution
- Update relevant data structures



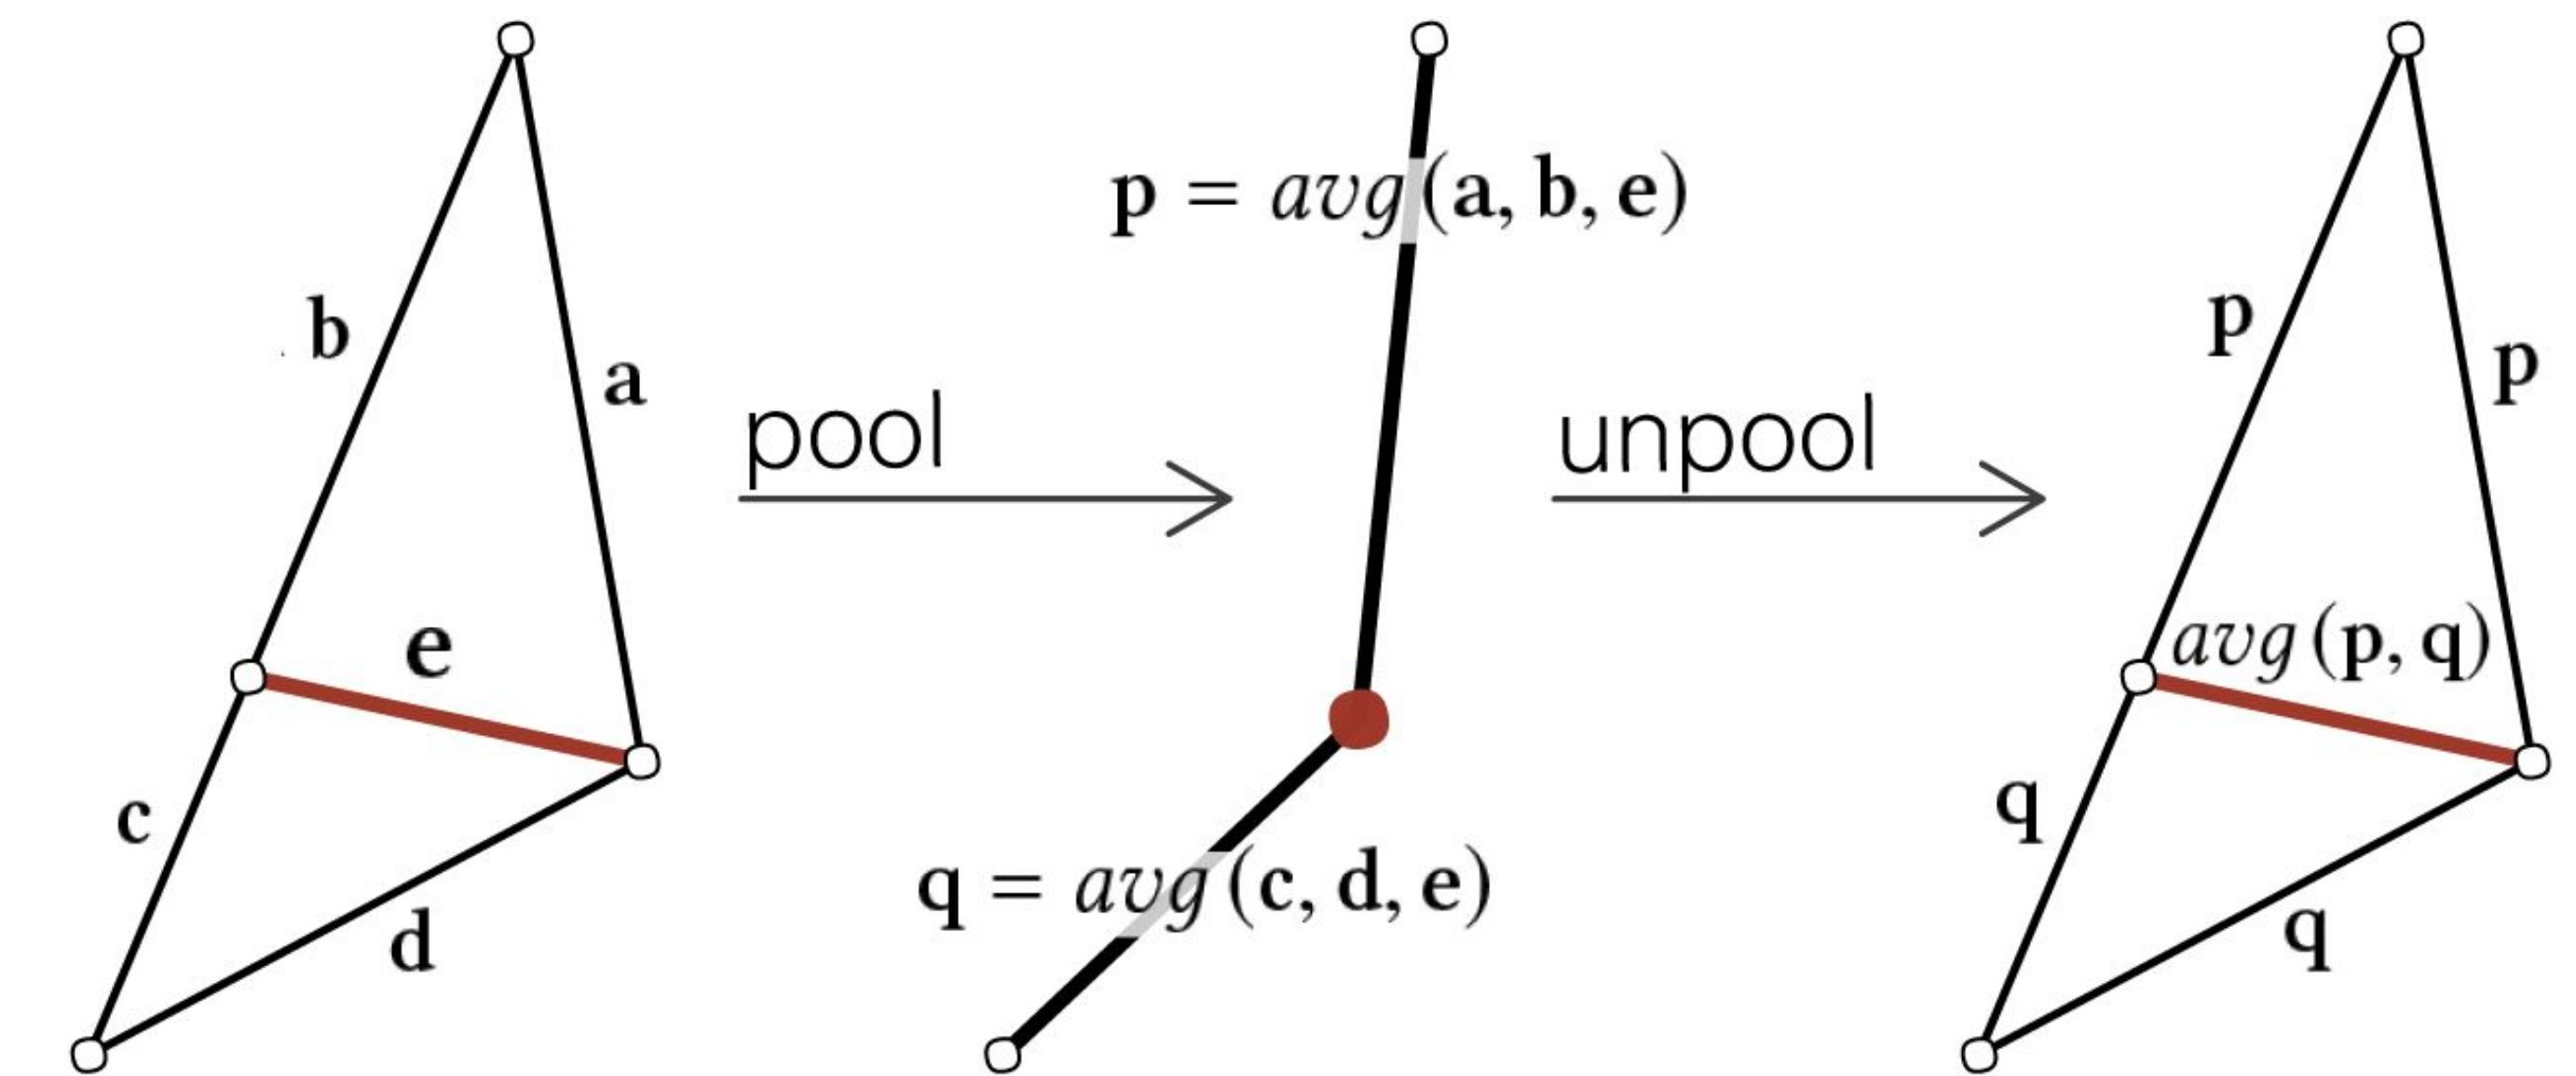
Common pooling results



Mesh Unpooling

Partial Inverse of Pooling

- Restores upsampled topology (reversible)
- Unpooled features weighted combination of pooled features



Key Attributes of Mesh Pooling

- Spatially adapts to the task
- Network decides collapse
- Flexible input mesh resolution
- Strengthens the learned feature representation
- Visual insights from the network



Next part: Spiral CNN

See the notebook!